

OpenImageIO 0.9

Programmer Documentation

(in progress)

Editor: Larry Gritz
lg@openimageio.org

Date: 19 Nov, 2010

*I kinda like “Oy-e-oh” with a bit of a groaning Yiddish accent, as in
“OIIO, did you really write yet another file I/O library?”*

Dan Wexler

Contents

1	Introduction	1
1.1	Overview	1
1.2	Simplifying Assumptions	2
I	The OpenImageIO Library	5
2	Image I/O API	7
2.1	Data Type Descriptions: TypeDesc	7
2.2	Image Specification: ImageSpec	9
3	ImageOutput: Writing Images	17
3.1	Image Output Made Simple	17
3.2	Advanced Image Output	18
3.3	ImageOutput Class Reference	34
4	Image I/O: Reading Images	39
4.1	Image Input Made Simple	39
4.2	Advanced Image Input	40
4.3	ImageInput Class Reference	49
5	Writing ImageIO Plugins	55
5.1	Plugin Introduction	55
5.2	Image Readers	55
5.3	Image Writers	62
5.4	Building ImageIO Plugins	72
6	Bundled ImageIO Plugins	73
6.1	BMP	74
6.2	Cineon	74
6.3	DDS	74
6.4	DPX	74
6.5	Field3D	74
6.6	FITS	74
6.7	HDR/RGBE	74

6.8	ICO	74
6.9	JPEG	74
6.10	JPEG-2000	74
6.11	OpenEXR	74
6.12	PNG	74
6.13	PNM	74
6.14	Ptex	74
6.15	SGI	74
6.16	Softimage PIC	74
6.17	Targa	74
6.18	TIFF	74
6.19	Zfile	74
7	Cached Images	75
7.1	Image Cache Introduction and Theory of Operation	75
7.2	ImageCache API	77
8	Texture Access: TextureSystem	85
8.1	Texture System Introduction and Theory of Operation	85
8.2	Helper Classes	85
8.3	TextureSystem API	92
9	Python Bindings	107
10	Image Buffer	109
II	Image Utilities	111
11	The <code>iv</code> Image Viewer	113
12	Getting Image information With <code>iinfo</code>	115
12.1	Using <code>iinfo</code>	115
12.2	<code>iinfo</code> command-line options	117
13	Converting Image Formats With <code>iconvert</code>	119
13.1	Overview	119
13.2	<code>iconvert</code> Recipes	119
13.3	<code>iconvert</code> command-line options	121
14	Searching Image Metadata With <code>igrep</code>	125
14.1	Using <code>igrep</code>	125
14.2	<code>igrep</code> command-line options	125

15 Comparing Images With <code>idiff</code>	127
15.1 Overview	127
15.2 Using <code>idiff</code>	127
15.3 <code>idiff</code> Reference	129
16 Making Tiled MIP-Map Texture Files With <code>maketx</code>	131
16.1 Overview	131
16.2 <code>maketx</code> command-line options	131
 III Appendices	 135
A Building OpenImageIO	137
B Metadata conventions	139
B.1 Description of the image	139
B.2 Display hints	140
B.3 Disk file format info/hints	140
B.4 Photographs or scanned images	141
B.5 Texture Information	142
B.6 Exif metadata	143
B.7 GPS Exif metadata	148
B.8 IPTC metadata	151
B.9 Extension conventions	152
 C Glossary	 153
 Index	 155

1 Introduction

Welcome to OpenImageIO!

1.1 Overview

OpenImageIO provides simple but powerful `ImageInput` and `ImageOutput` APIs that abstract the reading and writing of 2D image file formats. They don't support every possible way of encoding images in memory, but for a reasonable and common set of desired functionality, they provide an exceptionally easy way for an application using the APIs support a wide — and extensible — selection of image formats without knowing the details of any of these formats.

Concrete instances of these APIs, each of which implements the ability to read and/or write a different image file format, are stored as plugins (i.e., dynamic libraries, DLL's, or DSO's) that are loaded at runtime. The OpenImageIO distribution contains such plugins for several popular formats. Any user may create conforming plugins that implement reading and writing capabilities for other image formats, and any application that uses OpenImageIO would be able to use those plugins.

The library also implements the helper class `ImageBuf`, which is a handy way to store and manipulate images in memory. `ImageBuf` itself uses `ImageInput` and `ImageOutput` for its file I/O, and therefore is also agnostic as to image file formats.

The `ImageCache` class transparently manages a cache so that it can access truly vast amounts of image data (thousands of image files totaling tens of GB) very efficiently using only a tiny amount (tens of megabytes at most) of runtime memory. Additionally, a `TextureSystem` class provides filtered MIP-map texture lookups, atop the nice caching behavior of `ImageCache`.

Finally, the OpenImageIO distribution contains several utility programs that operate on images, each of which is built atop `ImageInput` and `ImageOutput`, and therefore may read or write any image file type for which an appropriate plugin is found at runtime. Paramount among these utilities is `iv`, a really fantastic and powerful image viewing application. Additionally, there are programs for converting images among different formats, comparing image data between two images, and examining image metadata.

All of this is released as “open source” software using the very permissive BSD license. So you should feel free to use any or all of OpenImageIO in your own software, whether it is private or public, open source or proprietary, free or commercial. You may also modify it on your own. You are also encouraged to contribute to the continued development of OpenImageIO and to share any improvements that you make on your own, though you are by no means required to do so.

1.2 Simplifying Assumptions

OpenImageIO is not the only image library in the world. Certainly there are many fine libraries that implement a single image format (including the excellent `libtiff`, `jpeg-6b`, and `OpenEXR` that OpenImageIO itself relies on). Many libraries attempt to present a uniform API for reading and writing multiple image file formats. Most of these support a fixed set of image formats, though a few of these also attempt to provide an extensible set by using the plugin approach.

But in our experience, these libraries are all flawed in one or more ways: (1) They either support only a few formats, or many formats but with the majority of them somehow incomplete or incorrect. (2) Their APIs are not sufficiently expressive as to handle all the image features we need (such as tiled images, which is critical for our texture library). (3) Their APIs are *too complete*, trying to handle every possible permutation of image format features, and as a result are horribly complicated.

The third sin is the most severe, and is almost always the main problem at the end of the day. Even among the many open source image libraries that rely on extensible plugins, we have not found one that is both sufficiently flexible and has APIs anywhere near as simple to understand and use as those of OpenImageIO.

Good design is usually a matter of deciding what *not* to do, and OpenImageIO is no exception. We achieve power and elegance only by making simplifying assumptions. Among them:

- OpenImageIO only deals with ordinary 2D images, and to a limited extent 3D volumes, and image files that contain multiple (but finite) independent images within them. OpenImageIO **does not deal with motion picture files**. At least, not currently.
- Pixel data are 8- 16- or 32-bit int (signed or unsigned), 16- 32- or 64-bit float. NOTHING ELSE. No < 8 bit images, or pixels boundaries that aren't byte boundaries. Files with < 8 bits will appear to the client as 8-bit unsigned grayscale images.
- Only fully elaborated, non-compressed data are accepted and returned by the API. Compression or special encodings are handled entirely within an OpenImageIO plugin.
- Color space is grayscale or RGB. Non-spectral data, such as XYZ, CMYK, or YUV, are converted to RGB upon reading.
- All color channels can be treated (by apps or readers/writers) as having the same data format (though there is a way to deal with per-channel formats for apps and readers/writers that truly need it).
- All image channels in a subimage are sampled at the same resolution. For file formats that allow some channels to be subsampled, they will be automatically up-sampled to the highest resolution channel in the subimage.
- Color information is always in the order R, G, B, and the alpha channel, if any, always follows RGB, and z channel (if any) always follows alpha. So if a file actually stores ABGR, the plugin is expected to rearrange it as RGBA.

It's important to remember that these restrictions apply to data passed through the APIs, not to the files themselves. It's perfectly fine to have an `OpenImageIO` plugin that supports YUV data, or 4 bits per channel, or any other exotic feature. You could even write a movie-reading `ImageInput` (despite `OpenImageIO`'s claims of not supporting movies) and make it look to the client like it's just a series of images within the file. It's just that all the nonconforming details are handled entirely within the `OpenImageIO` plugin and are not exposed through the main `OpenImageIO` APIs.

Historical Origins

`OpenImageIO` is the evolution of concepts and tools I've been working on for two decades.

In the 1980's, every program I wrote that output images would have a simple, custom format and viewer. I soon graduated to using a standard image file format (TIFF) with my own library implementation. Then I switched to Sam Leffler's stable and complete `libtiff`.

In the mid-to-late-1990's, I worked at Pixar as one of the main implementors of *PhotoRealistic RenderMan*, which had *display drivers* that consisted of an API for opening files and outputting pixels, and a set of DSO/DLL plugins that each implement image output for each of a dozen or so different file format. The plugins all responded to the same API, so the renderer itself did not need to know how to the details of the image file formats, and users could (in theory, but rarely in practice) extend the set of output image formats the renderer could use by writing their own plugins.

This was the seed of a good idea, but *PRMan*'s display driver plugin API was abstruse and hard to use. So when I started *Exluna* in 2000, Matt Pharr, Craig Kolb, and I designed a new API for image output for our own renderer, *Entropy*. This API, called "ExDisplay," was C++, and much simpler, clearer, and easier to use than *PRMan*'s display drivers.

NVIDIA's *Gelato* (circa 2002), whose early work was done by myself, Dan Wexler, Jonathan Rice, and Eric Enderton, had an API called "ImageIO." ImageIO was *much* more powerful and descriptive than ExDisplay, and had an API for *reading* as well as writing images. *Gelato* was not only "format agnostic" for its image output, but also for its image input (textures, image viewer, and other image utilities). We released the API specification and headers (though not the library implementation) using the BSD open source license, firmly repudiating any notion that the API should be specific to NVIDIA or *Gelato*.

For *Gelato* 3.0 (circa 2007), we refined ImageIO again (by this time, Philip Nemec was also a major influence, in addition to Dan, Eric, and myself¹). This revision was not a major overhaul but more of a fine tuning. Our ideas were clearly approaching stability. But, alas, the *Gelato* project was canceled before *Gelato* 3.0 was released, and despite our prodding, NVIDIA executives would not open source the full ImageIO code and related tools.

After I left NVIDIA, I was determined to recreate this work once again – and ONLY once more – and release it as open source from the start. Thus, `OpenImageIO` was born. I started with the existing *Gelato* ImageIO specification and headers (which were BSD licensed all along), and made some more refinements since I had to rewrite the entire implementation from scratch anyway. I think the additional changes are all improvements. This is the software you have in your hands today.

¹Gelato as a whole had many other contributors; those I've named here are the ones I recall contributing to the design or implementation of the ImageIO APIs

Acknowledgments

OpenImageIO incorporates or depends upon several other open source packages:

- libtiff © 1988-1997 Sam Leffler and 1991-1997 Silicon Graphics, Inc.
<http://www.remotesensing.org/libtiff>
- jpeg-6b © 1991-1998, Thomas G. Lane. <http://www.ijg.org>
- OpenEXR, Ilmbase, and Half © 2006, Industrial Light & Magic.
<http://www.openexr.com>
- zlib © 1995-2005 Jean-loup Gailly and Mark Adler. <http://www.zlib.net>
- libpng © 1998-2008 Glenn Randers-Pehrson, et al. <http://www.libpng.org>
- Thread Building Blocks (TBB) © 2005–2008 Intel Corporation.
<http://www.threadingbuildingblocks.org/>
- The SHA-1 implementation we use is public domain by Dominik Reichl
<http://www.dominik-reichl.de/>
- Boost © various authors. <http://www.boost.org>
- GLEW © 2002-2007 Milan Ikits, et al. <http://glew.sourceforge.net>
- Jasper © 2001-2006 Michael David Adams, et al.
<http://www.ece.uvic.ca/~mdadams/jasper/>
- Squish © 2006 Simon Brown. <http://sjbrown.co.uk/?code=squish>
- PugiXML © 2006-2009 by Arseny Kapoulkine (based on work ©2003 Kristen Wegner).
<http://code.google.com/p/pugixml/>
- DPX reader/writer © 2009 Patrick A. Palmer. <http://code.google.com/p/dpx>
- Ptex © 2009 Disney Enterprises, Inc. <http://ptex.us>

These other packages are all distributed under licenses that allow them to be used by and distributed with OpenImageIO.

Part I

The OpenImageIO Library

2 Image I/O API

2.1 Data Type Descriptions: TypeDesc

There are two kinds of data that are important to OpenImageIO:

- *Internal data* is in the memory of the computer, used by an application program.
- *Native file data* is what is stored in an image file itself (i.e., on the “other side” of the abstraction layer that OpenImageIO provides).

Both internal and file data is stored in a particular *data format* that describes the numerical encoding of the values. OpenImageIO understands several types of data encodings, and there is a special type, TypeDesc, that allows their enumeration. A TypeDesc describes a base data format type, aggregation into simple vector and matrix types, and an array length (if it’s an array).

TypeDesc supports the following base data format types, given by the enumerated type BASETYPE:

UINT8	8-bit integer values ranging from 0..255, corresponding to the C/C++ unsigned char.
INT8	8-bit integer values ranging from -128..127, corresponding to the C/C++ char.
UINT16	16-bit integer values ranging from 0..65535, corresponding to the C/C++ unsigned short.
INT16	16-bit integer values ranging from -32768..32767, corresponding to the C/C++ short.
UINT	32-bit integer values, corresponding to the C/C++ unsigned int.
INT	signed 32-bit integer values, corresponding to the C/C++ int.
UINT64	64-bit integer values, corresponding to the C/C++ unsigned long long (on most architectures).
INT64	signed 64-bit integer values, corresponding to the C/C++ long long (on most architectures).
FLOAT	32-bit IEEE floating point values, corresponding to the C/C++ float.
DOUBLE	64-bit IEEE floating point values, corresponding to the C/C++ double.
HALF	16-bit floating point values in the format supported by OpenEXR and OpenGL.

A TypeDesc can be constructed using just this information, either as a single scalar value, or an array of scalar values:

`TypeDesc (BASETYPE btype)`

`TypeDesc (BASETYPE btype, int arraylength)`

Construct a type description of a single scalar value of the given base type, or an array of such scalars if an array length is supplied. For example, `TypeDesc(UINT8)` describes an unsigned 8-bit integer, and `TypeDesc(FLOAT, 7)` describes an array of 7 32-bit float values. Note also that a non-array `TypeDesc` may be implicitly constructed from just the `BASETYPE`, so it's okay to pass a `BASETYPE` to any function parameter that takes a full `TypeDesc`.

In addition, `TypeDesc` supports certain aggregate types, described by the enumerated type `AGGREGATE`:

<code>SCALAR</code>	a single scalar value (such as a raw <code>int</code> or <code>float</code> in C). This is the default.
<code>VEC2</code>	two values representing a 2D vector.
<code>VEC3</code>	three values representing a 3D vector.
<code>VEC4</code>	four values representing a 4D vector.
<code>MATRIX44</code>	sixteen values representing a 4×4 matrix.

And optionally, several vector transformation semantics, described by the enumerated type `VECSEMANTICS`:

<code>NOXFORM</code>	indicates that the item is not a spatial quantity that undergoes any particular transformation.
<code>COLOR</code>	indicates that the item is a “color,” not a spatial quantity (and of course therefore does not undergo a transformation).
<code>POINT</code>	indicates that the item represents a spatial position and should be transformed by a 4×4 matrix as if it had a 4th component of 1.
<code>VECTOR</code>	indicates that the item represents a spatial direction and should be transformed by a 4×4 matrix as if it had a 4th component of 0.
<code>NORMAL</code>	indicates that the item represents a surface normal and should be transformed like a vector, but using the inverse-transpose of a 4×4 matrix.

These can be combined to fully describe a complex type:

`TypeDesc (BASETYPE btype, AGGREGATE agg, VECSEMANTICS xform=NOXFORM)`

`TypeDesc (BASETYPE btype, AGGREGATE agg, int arraylen)`

`TypeDesc (BASETYPE btype, AGGREGATE agg, VECSEMANTICS xform, int arraylen)`

Construct a type description of an aggregate (or array of aggregates), with optional vector transformation semantics. For example, `TypeDesc(HALF, COLOR)` describes an aggregate of 3 16-bit floats comprising a color, and `TypeDesc(FLOAT, VEC3, POINT)` describes an aggregate of 3 32-bit floats comprising a 3D position.

Note that aggregates and arrays are different. A `TypeDesc(FLOAT, 3)` is an array of three floats, a `TypeDesc(FLOAT, COLOR)` is a single 3-channel color comprised of floats, and `TypeDesc(FLOAT, 3, COLOR)` is an array of 3 color values, each of which is comprised of 3 floats.

Of these, the only ones commonly used to store pixel values in image files are scalars of `UINT8`, `UINT16`, `FLOAT`, and `HALF` (the last only used by `OpenEXR`, to the best of our knowledge).

Note that the `TypeDesc` (which is also used for applications other than images) can describe many types not used by `OpenImageIO`. Please ignore this extra complexity; only the above simple types are understood by `OpenImageIO` as pixel storage data types, though a few others, including `STRING` and `MATRIX44` aggregates, are occasionally used for *metadata* for certain image file formats (see Sections 3.2.5, 4.2.4, and the documentation of individual `ImageIO` plugins for details).

2.2 Image Specification: ImageSpec

An `ImageSpec` is a structure that describes the complete format specification of a single image. It contains:

- The image resolution (number of pixels).
- The origin, if its upper left corner is not located beginning at pixel (0,0).
- The full size and offset of an abstract “full” or “display” image, useful for describing cropping or overscan.
- Whether the image is organized into *tiles*, and if so, the tile size.
- The *native data format* of the pixel values (e.g., float, 8-bit integer, etc.).
- The number of color channels in the image (e.g., 3 for RGB images), names of the channels, and whether any particular channels represent *alpha* and *depth*.
- Any presumed gamma correction or hints about color space of the pixel values.
- Quantization parameters describing how floating point values should be converted to integers (in cases where users pass real values but integer values are stored in the file). This is used only when writing images, not when reading them.
- A user-extensible (and format-extensible) list of any other arbitrarily-named and -typed data that may help describe the image or its disk representation.

2.2.1 ImageSpec Data Members

The `ImageSpec` contains data fields for the values that are required to describe nearly any image, and an extensible list of arbitrary attributes that can hold metadata that may be user-defined or specific to individual file formats. Here are the hard-coded data fields:

```
int width, height, depth
int x, y, z
```

`width`, `height`, `depth` are the size of the data of this image, i.e., the number of pixels in each dimension. A `depth` greater than 1 indicates a 3D “volumetric” image.

`x`, `y`, `z` indicate the *origin* of the pixel data of the image. These default to (0,0,0), but setting them differently may indicate that this image is offset from the usual origin.

Therefore the pixel data are defined over pixel coordinates [`x` ... `x+width-1`] horizontally, [`y` ... `y+height-1`] vertically, and [`z` ... `z+depth-1`] in depth.

```
int full_width, full_height, full_depth
int full_x, full_y, full_z
```

These fields define a “full” or “display” image window over the region [`full_x` ... `full_x+full_width-1`] horizontally, [`full_y` ... `full_y+full_height-1`] vertically, and [`full_z` ... `full_z+full_depth-1`] in depth.

Having the full display window different from the pixel data window can be helpful in cases where you want to indicate that your image is a *crop window* of a larger image (if the pixel data window is a subset of the full display window), or that the pixels include *overscan* (if the pixel data is a superset of the full display window), or may simply indicate how different non-overlapping images piece together.

```
int tile_width, tile_height, tile_depth
```

If nonzero, indicates that the image is stored on disk organized into rectangular *tiles* of the given dimension. The default of (0,0,0) indicates that the image is stored in scanline order, rather than as tiles.

```
int nchannels
```

The number of *channels* (color values) present in each pixel of the image. For example, an RGB image has 3 channels.

TypeDesc format

`std::vector<TypeDesc> channelformats`

Describes the native format of the pixel data values themselves, as a `TypeDesc` (see 2.1). Typical values would be `TypeDesc : :UINT8` for 8-bit unsigned values, `TypeDesc : :FLOAT` for 32-bit floating-point values, etc.

If all channels of the image have the same data format, that will be described by `format` and `channelformats` will be empty (zero length).

If there are different data formats for each channel, they will be described in the `channelformats` vector, and the `format` field will indicate a single default data format for applications that don't wish to support per-channel formats (usually this will be the format of the channel that has the most precision).

`std::vector<std::string> channelnames`

The names of each channel, in order. Typically this will be "R", "G", "B", "A" (alpha), "Z" (depth), or other arbitrary names.

`int alpha_channel`

The index of the channel that represents *alpha* (pixel coverage and/or transparency). It defaults to -1 if no alpha channel is present, or if it is not known which channel represents alpha.

`int z_channel`

The index of the channel that represents *z* or *depth* (from the camera). It defaults to -1 if no depth channel is present, or if it is not known which channel represents depth.

`LinearitySpec linearity`

Describes the mapping of pixel values to real-world units. `LinearitySpec` is an enumerated type that may take on the following values:

- `Linear` (the default) indicates that pixel values map linearly.
- `GammaCorrected` indicates that the color pixel values have already been gamma corrected, using the exponent given by the `gamma` field. (It is still assumed that non-color values, such as alpha and depth, are linear.)
- `sRGB` indicates that color values are encoded using the sRGB mapping. (It is still assumed that non-color values are linear.)
- `AdobeRGB` indicates that the values are encoded in the Adobe RGB color space. (It is still assumed that non-color values are linear.)
- `Rec709` indicates that color values are encoded using the Rec709 mapping. (It is still assumed that non-color values are linear.)
- `KodakLog` indicates that color values are encoded using the Kodak logarithmic mapping. (It is still assumed that non-color values are linear.)

`float gamma`

The gamma exponent, if the pixel values in the image have already been gamma corrected (indicated by `linearity` having a value of `GammaCorrected`). The default of 1.0 indicates that no gamma correction has been applied.

`int quant_black, quant_white, quant_min, quant_max;`

`float quant_dither`

Describes the *quantization*, or mapping between real (floating-point) values and the stored integer values. Please refer to Section 3.2.6 for a more complete explanation of each of these parameters.

`ParamValueList extra_attrs`

A list of arbitrarily-named and arbitrarily-typed additional attributes of the image, for any metadata not described by the hard-coded fields described above. This list may be manipulated with the `attribute()` and `find_attribute()` methods.

2.2.2 ImageSpec member functions

`ImageSpec` contains the following methods that manipulate format specs or compute useful information about images given their format spec:

`ImageSpec (int xres, int yres, int nchans, TypeDesc fmt = UINT8)`

Constructs an `ImageSpec` with the given *x* and *y* resolution, number of channels, and pixel data format.

All other fields are set to the obvious defaults – the image is an ordinary 2D image (not a volume), the image is not offset or a crop of a bigger image, the image is scanline-oriented (not tiled), channel names are “R”, “G”, “B,” and “A” (up to and including 4 channels, beyond that they are named “channel *n*”), the fourth channel (if it exists) is assumed to be alpha, values are assumed to be linear, and quantization (if *fmt* describes an integer type) is done in such a way that the maximum positive integer range maps to (0.0, 1.0).

`void set_format (TypeDesc fmt)`

Sets the format as described, and also sets all quantization parameters to the default for that data type (as explained in Section 3.2.6).

`void default_channel_names ()`

Sets the `channelnames` to reasonable defaults for the number of channels. Specifically, channel names are set to “R”, “G”, “B,” and “A” (up to and including 4 channels, beyond that they are named “channel*n*”).

```
static TypeDesc
format_from_quantize (int quant_black, int quant_white,
                     int quant_min, int quant_max)
```

Utility function that, given quantization parameters, returns a data type that may be used without unacceptable loss of significant bits.

```
size_t channel_bytes () const
```

Returns the number of bytes comprising each channel of each pixel (i.e., the size of a single value of the type described by the format field).

```
size_t pixel_bytes (bool native=false) const
```

Returns the number of bytes comprising each pixel (i.e. the number of channels multiplied by the channel size).

If `native` is true, this will be the sum of all the per-channel formats in `channelformats`. If `native` is false (the default), or if all channels use the same format, this will simply be the number of channels multiplied by the width (in bytes) of the format.

```
imagesize_t scanline_bytes (bool native=false) const
```

Returns the number of bytes comprising each scanline, i.e.,

```
pixel_bytes(native) * width
```

This will return `std::numeric_limits<imagesize_t>::max()` in the event of an overflow where it's not representable in an `imagesize_t`.

```
imagesize_t tile_pixels () const
```

Returns the number of tiles comprising an image tile (if it's a tiled image). This will return `std::numeric_limits<imagesize_t>::max()` in the event of an overflow where it's not representable in an `imagesize_t`.

```
imagesize_t tile_bytes (bool native=false) const
```

Returns the number of bytes comprising an image tile (if it's a tiled image), i.e.,

```
pixel_bytes(native) * tile_width * tile_height * tile_depth
```

This will return `std::numeric_limits<imagesize_t>::max()` in the event of an overflow where it's not representable in an `imagesize_t`.

```
imagesize_t image_pixels () const
```

Returns the number of pixels comprising an entire image image of these dimensions. This will return `std::numeric_limits<imagesize_t>::max()` in the event of an overflow where it's not representable in an `imagesize_t`.

```
imagesize_t image_bytes (bool native=false) const
```

Returns the number of bytes comprising an entire image of these dimensions, i.e., `pixel_bytes(native) * width * height * depth`. This will return `std::numeric_limits<imagesize_t>::max()` in the event of an overflow where it's not representable in an `imagesize_t`.

```
bool size_t_safe () const
```

Return true if an image described by this spec can the sizes (in pixels or bytes) of its scanlines, tiles, and the entire image can be represented by a `size_t` on that platform. If this returns false, the client application should be very careful allocating storage!

```
void attribute (const std::string &name, TypeDesc type,
               const void *value)
```

Add a metadata attribute to `extra_attribs`, with the given name and data type. The value pointer specifies the address of the data to be copied.

```
void attribute (const std::string &name, unsigned int value)
```

```
void attribute (const std::string &name, int value)
```

```
void attribute (const std::string &name, float value)
```

```
void attribute (const std::string &name, const char *value)
```

```
void attribute (const std::string &name, const std::string &value)
```

Shortcuts for passing attributes comprised of a single integer, floating-point value, or string.

```
ImageIOParameter * find_attribute (const std::string &name,
                                  TypeDesc searchtype=UNKNOWN,
                                  bool casesensitive=false)
```

```
const ImageIOParameter * find_attribute (const std::string &name,
                                         TypeDesc searchtype=UNKNOWN,
                                         bool casesensitive=false) const
```

Searches `extra_attribs` for an attribute matching name, returning a pointer to the attribute record, or NULL if there was no match. If `searchtype` is `TypeDesc::UNKNOWN`, the search will be made regardless of the data type, whereas other values of `searchtype` will reject a matching name if the data type does not also match. The name comparison will be exact if `casesensitive` is true, otherwise in a case-insensitive manner if `casesensitive` is false.

```
int get_int_attribute (const std::string &name, int defaultval=0) const
```

Gets an integer metadata attribute (silently converting to int even if the data is really int8, uint8, int16, uint16, or uint32), and simply substituting the supplied default value if no such metadata exists. This is a convenience function for when you know you are just looking for a simple integer value.

```
float get_float_attribute (const std::string &name,  
                           float defaultval=0) const
```

Gets a float metadata attribute (silently converting to float even if the data is really half or double), simply substituting the supplied default value if no such metadata exists. This is a convenience function for when you know you are just looking for a simple float value.

```
std::string get_string_attribute (const std::string &name,  
                                 const std::string &defaultval=std::string()) const
```

Gets a string metadata attribute, simply substituting the supplied default value if no such metadata exists. This is a convenience function for when you know you are just looking for a simple string value.

```
std::string metadata_val (const ImageIOParameter &p, bool human=true) const
```

For a given parameter (in this ImageSpec's `extra_attribs` field), format the value nicely as a string. If `human` is true, use especially human-readable explanations (units, or decoding of values) for certain known metadata.

3 ImageOutput: Writing Images

3.1 Image Output Made Simple

Here is the simplest sequence required to write the pixels of a 2D image to a file:

```
#include "imageio.h"
using namespace OpenImageIO;
...

const char *filename = "foo.jpg";
const int xres = 640, yres = 480;
const int channels = 3; // RGB
unsigned char pixels[xres*yres*channels];

ImageOutput *out = ImageOutput::create (filename);
if (! out)
    return;
ImageSpec spec (xres, yres, channels, TypeDesc::UINT8);
out->open (filename, spec);
out->write_image (TypeDesc::UINT8, pixels);
out->close ();
delete out;
```

This little bit of code does a surprising amount of useful work:

- Search for an ImageIO plugin that is capable of writing the file ("foo.jpg"), deducing the format from the file extension. When it finds such a plugin, it creates a subclass instance of ImageOutput that writes the right kind of file format.

```
ImageOutput *out = ImageOutput::create (filename);
```

- Open the file, write the correct headers, and in all other important ways prepare a file with the given dimensions (640×480), number of color channels (3), and data format (unsigned 8-bit integer).

```
ImageSpec spec (xres, yres, channels, TypeDesc::UINT8);
out->open (filename, spec);
```

- Write the entire image, hiding all details of the encoding of image data in the file, whether the file is scanline- or tile-based, or what is the native format of data in the file (in this case, our in-memory data is unsigned 8-bit and we've requested the same format for disk storage, but if they had been different, `write_image()` would do all the conversions for us).

```
out->write_image (TypeDesc::UINT8, &pixels);
```

- Close the file, destroy and free the `ImageOutput` we had created, and perform all other cleanup and release of any resources needed by the plugin.

```
out->close ();  
delete out;
```

3.2 Advanced Image Output

Let's walk through many of the most common things you might want to do, but that are more complex than the simple example above.

3.2.1 Writing individual scanlines, tiles, and rectangles

The simple example of Section 3.1 wrote an entire image with one call. But sometimes you are generating output a little at a time and do not wish to retain the entire image in memory until it is time to write the file. `OpenImageIO` allows you to write images one scanline at a time, one tile at a time, or by individual rectangles.

Writing individual scanlines

Individual scanlines may be written using the `write_scanline()` API call:

```
...  
unsigned char scanline[xres*channels];  
out->open (filename, spec);  
int z = 0;    // Always zero for 2D images  
for (int y = 0; y < yres; ++y) {  
    ... generate data in scanline[0..xres*channels-1] ...  
    out->write_scanline (y, z, TypeDesc::UINT8, scanline);  
}  
out->close ();  
...
```

The first two arguments to `write_scanline()` specify which scanline is being written by its vertical (*y*) scanline number (beginning with 0) and, for volume images, its slice (*z*) number (the slice number should be 0 for 2D non-volume images). This is followed by a `TypeDesc` describing the data you are supplying, and a pointer to the pixel data itself. Additional optional

arguments describe the data stride, which can be ignored for contiguous data (use of strides is explained in Section 3.2.3).

All `ImageOutput` implementations will accept scanlines in strict order (starting with scanline 0, then 1, up to `yres-1`, without skipping any). See Section 3.2.7 for details on out-of-order or repeated scanlines.

The full description of the `write_scanline()` function may be found in Section 3.3.

Writing individual tiles

Not all image formats (and therefore not all `ImageOutput` implementations) support tiled images. If the format does not support tiles, then `write_tile()` will fail. An application using `OpenImageIO` should gracefully handle the case that tiled output is not available for the chosen format.

Once you `create()` an `ImageOutput`, you can ask if it is capable of writing a tiled image by using the `supports("tiles")` query:

```
...
ImageOutput *out = ImageOutput::create (filename);
if (! out->supports ("tiles")) {
    // Tiles are not supported
}
```

Assuming that the `ImageOutput` supports tiled images, you need to specifically request a tiled image when you `open()` the file. This is done by setting the tile size in the `ImageSpec` passed to `open()`. If the tile dimensions are not set, they will default to zero, which indicates that scanline output should be used rather than tiled output.

```
int tileSize = 64;
ImageSpec spec (xres, yres, channels, TypeDesc::UINT8);
spec.tile_width = tileSize;
spec.tile_height = tileSize;
out->open (filename, spec);
...
```

In this example, we have used square tiles (the same number of pixels horizontally and vertically), but this is not a requirement of `OpenImageIO`. However, it is possible that some image formats may only support square tiles, or only certain tile sizes (such as restricting tile sizes to powers of two). Such restrictions should be documented by each individual plugin.

```
unsigned char tile[tileSize*tileSize*channels];
int z = 0;    // Always zero for 2D images
for (int y = 0; y < yres; y += tileSize) {
    for (int x = 0; x < xres; x += tileSize) {
        ... generate data in tile[] ..
        out->write_tile (x, y, z, TypeDesc::UINT8, tile);
    }
}
out->close ();
...
```

The first three arguments to `write_tile()` specify which tile is being written by the pixel coordinates of any pixel contained in the tile: x (column), y (scanline), and z (slice, which should always be 0 for 2D non-volume images). This is followed by a `TypeDesc` describing the data you are supplying, and a pointer to the tile's pixel data itself, which should be ordered by increasing slice, increasing scanline within each slice, and increasing column within each scanline. Additional optional arguments describe the data stride, which can be ignored for contiguous data (use of strides is explained in Section 3.2.3).

All `ImageOutput` implementations that support tiles will accept tiles in strict order of increasing y rows, and within each row, increasing x column, without missing any tiles. See Section 3.2.7 for details on out-of-order or repeated tiles.

The full description of the `write_tile()` function may be found in Section 3.3.

Writing arbitrary rectangles

Some `ImageOutput` implementations — such as those implementing an interactive image display, but probably not any that are outputting directly to a file — may allow you to send arbitrary rectangular pixel regions. Once you create() an `ImageOutput`, you can ask if it is capable of accepting arbitrary rectangles by using the `supports("rectangles")` query:

```
...
ImageOutput *out = ImageOutput::create (filename);
if (! out->supports ("rectangles")) {
    // Rectangles are not supported
}
```

If rectangular regions are supported, they may be sent using the `write_rectangle()` API call:

```
unsigned int rect[...];
... generate data in rect[] ..
out->write_rectangle (xmin, xmax, ymin, ymax, zmin, zmax, TypeDesc::UINT8, rect);
...
```

The first six arguments to `write_rectangle()` specify the region of pixels that is being transmitted by supplying the minimum and maximum pixel indices in x (column), y (scanline), and z (slice, always 0 for 2D non-volume images). The total number of pixels being transmitted is therefore:

$$(xmax-xmin+1) * (ymax-ymin+1) * (zmax-zmin+1)$$

This is followed by a `TypeDesc` describing the data you are supplying, and a pointer to the rectangle's pixel data itself, which should be ordered by increasing slice, increasing scanline within each slice, and increasing column within each scanline. Additional optional arguments describe the data stride, which can be ignored for contiguous data (use of strides is explained in Section 3.2.3).

3.2.2 Converting data formats

The code examples of the previous sections all assumed that your internal pixel data is stored as unsigned 8-bit integers (i.e., 0-255 range). But `OpenImageIO` is significantly more flexible.

You may request that the output image be stored in any of several formats. This is done by setting the `format` field of the `ImageSpec` prior to calling `open`. You can do this upon construction of the `ImageSpec`, as in the following example that requests a spec that stores data as 16-bit unsigned integers:

```
ImageSpec spec (xres, yres, channels, TypeDesc::UINT16);
```

Or, for an `ImageSpec` that has already been constructed, you may reset its format using the `set_format()` method (which also resets the various quantization fields of the spec to the defaults for the data format you have specified).

```
ImageSpec spec (...);  
spec.set_format (TypeDesc::UINT16);
```

Note that resetting the format must be done *before* passing the spec to `open()`, or it will have no effect on the file.

Individual file formats, and therefore `ImageOutput` implementations, may only support a subset of the formats understood by the `OpenImageIO` library. Each `ImageOutput` plugin implementation should document which data formats it supports. An individual `ImageOutput` implementation may choose to simply fail to `open()`, though the recommended behavior is for `open()` to succeed but in fact choose a data format supported by the file format that best preserves the precision and range of the originally-requested data format.

It is not required that the pixel data passed to `write_image()`, `write_scanline()`, `write_tile()`, or `write_rectangle()` actually be in the same data format as that requested as the native format of the file. You can fully mix and match data you pass to the various `write` routines and `OpenImageIO` will automatically convert from the internal format to the native file format. For example, the following code will open a TIFF file that stores pixel data as 16-bit unsigned integers (values ranging from 0 to 65535), compute internal pixel values as floating-point values, with `write_image()` performing the conversion automatically:

```
ImageOutput *out = ImageOutput::create ("myfile.tif");  
ImageSpec spec (xres, yres, channels, TypeDesc::UINT16);  
out->open (filename, spec);  
...  
float pixels [xres*yres*channels];  
...  
out->write_image (TypeDesc::FLOAT, pixels);
```

Note that `write_scanline()`, `write_tile()`, and `write_rectangle` have a parameter that works in a corresponding manner.

Please refer to Section 3.2.6 for more information on how values are translated among the supported data formats by default, and how to change the formulas by specifying quantization in the `ImageSpec`.

3.2.3 Data Strides

In the preceeding examples, we have assumed that the block of data being passed to the `write` functions are *contiguous*, that is:

- each pixel in memory consists of a number of data values equal to the declared number of channels that are being written to the file;
- successive column pixels within a row directly follow each other in memory, with the first channel of pixel x immediately following last channel of pixel $x - 1$ of the same row;
- for whole images, tiles or rectangles, the data for each row immediately follows the previous one in memory (the first pixel of row y immediately follows the last column of row $y - 1$);
- for 3D volumetric images, the first pixel of slice z immediately follows the last pixel of slice $z - 1$.

Please note that this implies that data passed to `write_tile()` be contiguous in the shape of a single tile (not just an offset into a whole image worth of pixels), and that data passed to `write_rectangle()` be contiguous in the dimensions of the rectangle.

The `write_scanline()` function takes an optional `xstride` argument, and the `write_image()`, `write_tile()`, and `write_rectangle` functions take optional `xstride`, `ystride`, and `zstride` values that describe the distance, in *bytes*, between successive pixel columns, rows, and slices, respectively, of the data you are passing. For any of these values that are not supplied, or are given as the special constant `AutoStride`, contiguity will be assumed.

By passing different stride values, you can achieve some surprisingly flexible functionality. A few representative examples follow:

- Flip an image vertically upon writing, by using *negative y* stride:

```
unsigned char pixels[xres*yres*channels];
int scanlinesize = xres * channels * sizeof(pixels[0]);
...
out->write_image (TypeDesc::UINT8,
                  (char *)pixels+(yres-1)*scanlinesize, // offset to last
                  AutoStride,                          // default x stride
                  -scanlinesize,                        // special y stride
                  AutoStride);                          // default z stride
```

- Write a tile that is embedded within a whole image of pixel data, rather than having a one-tile-only memory layout:

```
unsigned char pixels[xres*yres*channels];
int pixelsize = channels * sizeof(pixels[0]);
int scanlinesize = xres * pixelsize;
...
out->write_tile (x, y, 0, TypeDesc::UINT8,
```

```
(char *)pixels + y*scanlinesize + x*pixelsize,
pixelsize,
scanlinesize);
```

- Write only a subset of channels to disk. In this example, our internal data layout consists of 4 channels, but we write just channel 3 to disk as a one-channel image:

```
// In-memory representation is 4 channel
const int xres = 640, yres = 480;
const int channels = 4; // RGBA
const int channelsize = sizeof(unsigned char);
unsigned char pixels[xres*yres*channelsize];

// File representation is 1 channel
ImageOutput *out = ImageOutput::create (filename);
ImageSpec spec (xres, yres, 1, TypeDesc::UINT8);
out->open (filename, spec);

// Use strides to write out a one-channel "slice" of the image
out->write_image (TypeDesc::UINT8,
                 (char *)pixels+3*channelsize, // offset to chan 3
                 channelsize,                // 4 channel x stride
                 AutoStride,                  // default y stride
                 AutoStride);                 // default z stride
...
```

Please consult Section 3.3 for detailed descriptions of the stride parameters to each write function.

3.2.4 Writing a crop window or overscan region

The ImageSpec fields `width`, `height`, and `depth` describe the dimensions of the actual pixel data.

At times, it may be useful to also describe an abstract *full* or *display* image window, whose position and size may not correspond exactly to the data pixels. For example, a pixel data window that is a subset of the full display window might indicate a *crop window*; a pixel data window that is a superset of the full display window might indicate *overscan* regions (pixels defined outside the eventual viewport).

The ImageSpec fields `full_width`, `full_height`, and `full_depth` describe the dimensions of the full display window, and `full_x`, `full_y`, `full_z` describe its origin (upper left corner). The fields `x`, `y`, `z` describe the origin (upper left corner) of the pixel data.

These fields collectively describe an abstract full display image ranging from `[full_x ... full_x+full_width-1]` horizontally, `[full_y ... full_y+full_height-1]` vertically, and `[full_z ... full_z+full_depth-1]` in depth (if it is a 3D volume), and actual pixel data over the pixel coordinate range `[x ... x+width-1]` horizontally, `[y ... y+height-1]` vertically, and `[z ... z+depth-1]` in depth (if it is a volume).

Not all image file formats have a way to describe display windows. An ImageOutput implementation that cannot express display windows will always write out the $\text{width} \times \text{height}$ pixel data, may upon writing lose information about offsets or crop windows.

Here is a code example that opens an image file that will contain a 32×32 pixel crop window within an abstract 640×480 full size image. Notice that the pixel indices (column, scanline, slice) passed to the write functions are the coordinates relative to the full image, not relative to the crop widow, but the data pointer passed to the write functions should point to the beginning of the actual pixel data being passed (not the the hypothetical start of the full data, if it was all present).

```
int fullwidth = 640, fulllength = 480; // Full display image size
int cropwidth = 16, croplength = 16;  // Crop window size
int xorigin = 32, yorigin = 128;      // Crop window position
unsigned char pixels [cropwidth * croplength * channels]; // Crop size!
...
ImageOutput *out = ImageOutput::create (filename);
ImageSpec spec (cropwidth, croplength, channels, TypeDesc::UINT8);
spec.full_x = 0;
spec.full_y = 0;
spec.full_width = fullwidth;
spec.full_length = fulllength;
spec.x = xorigin;
spec.y = yorigin;
out->open (filename, spec);
...
int z = 0; // Always zero for 2D images
for (int y = yorigin; y < yorigin+croplength; ++y) {
    out->write_scanline (y, z, TypeDesc::UINT8,
                        (y-yorigin)*cropwidth*channels);
}
out->close ();
```

3.2.5 Writing metadata

The ImageSpec passed to open() can specify all the common required properties that describe an image: data format, dimensions, number of channels, tiling. However, there may be a variety of additional *metadata*¹ that should be carried along with the image or saved in the file.

The remainder of this section explains how to store additional metadata in the ImageSpec. It is up to the ImageOutput to store these in the file, if indeed the file format is able to accept the data. Individual ImageOutput implementations should document which metadata they respect.

Channel names

In addition to specifying the number of color channels, it is also possible to name those channels. Only a few ImageOutput implementations have a way of saving this in the file, but some do, so you may as well do it if you have information about what the channels represent.

¹Metadata refers to data about data, in this case, data about the image that goes beyond the pixel values and description thereof.

By convention, channel names for red, green, blue, and alpha (or a main image) should be named "R", "G", "B", and "A", respectively. Beyond this guideline, however, you can use any names you want.

The ImageSpec has a vector of strings called `channelnames`. Upon construction, it starts out with reasonable default values. If you use it at all, you should make sure that it contains the same number of strings as the number of color channels in your image. Here is an example:

```
int channels = 4;
ImageSpec spec (width, length, channels, TypeDesc::UINT8);
spec.channelnames.clear ();
spec.channelnames.push_back ("R");
spec.channelnames.push_back ("G");
spec.channelnames.push_back ("B");
spec.channelnames.push_back ("A");
```

Here is another example in which custom channel names are used to label the channels in an 8-channel image containing beauty pass RGB, per-channel opacity, and texture *s*, *t* coordinates for each pixel.

```
int channels = 8;
ImageSpec spec (width, length, channels, TypeDesc::UINT8);
spec.channelnames.clear ();
spec.channelnames.push_back ("R");
spec.channelnames.push_back ("G");
spec.channelnames.push_back ("B");
spec.channelnames.push_back ("opacityR");
spec.channelnames.push_back ("opacityG");
spec.channelnames.push_back ("opacityB");
spec.channelnames.push_back ("texture_s");
spec.channelnames.push_back ("texture_t");
```

The main advantage to naming color channels is that if you are saving to a file format that supports channel names, then any application that uses `OpenImageIO` to read the image back has the option to retain those names and use them for helpful purposes. For example, the `iv` image viewer will display the channel names when viewing individual channels or displaying numeric pixel values in “pixel view” mode.

Specially-designated channels

The ImageSpec contains two fields, `alpha_channel` and `z_channel`, which can be used to designate which channel indices are used for alpha and *z* depth, if any. Upon construction, these are both set to -1, indicating that it is not known which channels are alpha or depth. Here is an example of setting up a 5-channel output that represents RGBAZ:

```
int channels = 5;
ImageSpec spec (width, length, channels, format);
spec.channelnames.push_back ("R");
spec.channelnames.push_back ("G");
spec.channelnames.push_back ("B");
```

```
spec.channelnames.push_back ("A");  
spec.channelnames.push_back ("Z");  
spec.alpha_channel = 3;  
spec.z_channel = 4;
```

There are two advantages to designating the alpha and depth channels in this manner:

- Some file formats may require that these channels be stored in a particular order, with a particular precision, or the `ImageOutput` may in some other way need to know about these special channels.
- Certain operations that make sense for colors should not apply to alpha or *z*. For example, if your call to `write` reduces precision (e.g., converts from `float` to integer pixels) it will typically add random *dither* to eliminate banding artifacts in the quantization. But for a variety of reasons, you want to add dither only to color channels and not to alpha. So setting `alpha_channel` will cause `write` to not dither that channel.

Linearity hints

We certainly hope that you are using only modern file formats that support high precision and extended range pixels (such as OpenEXR) and keeping all your images in a linear color space. But you may have to work with file formats that dictate the use of nonlinear color values. This is prevalent in formats that store pixels only as 8-bit values, since 256 values are not enough to linearly represent colors without banding artifacts in the dim values.

Since this can (and probably will) happen, the `ImageSpec` has fields that allow you to explain what color space your image pixels are in. Each individual `ImageOutput` should document how it uses this (or not).

The `ImageSpec` field `linearity` can take on any of the following values:

`ImageSpec::UnknownLinearity` the default, indicates that you have made no claim about the color space of your pixel data.

`ImageSpec::Linear` indicates that the pixel values you are passing represent linear values.

`ImageSpec::GammaCorrected` indicates that the color pixel values (but not alpha or *z*) that you are passing have already been gamma corrected (raised to the power $1/\gamma$), and that the gamma exponent may be found in the `gamma` field of the `ImageSpec`.

`ImageSpec::sRGB` indicates that the color pixel values that you are passing are already in sRGB color space.

`ImageSpec::AdobeRGB` indicates that the color pixel values that you are passing are already in Adobe RGB color space.

`ImageSpec::Rec709` indicates that the color pixel values that you are passing are already in Rec709 color space.

`ImageSpec::KodakLog` indicates that the color pixel values that you are passing are already in Kodak logarithmic color space.

Here is a simple example of setting up the ImageSpec when you know that the pixel values you are writing are linear:

```
ImageSpec spec (width, length, channels, format);
spec.linearity = ImageSpec::Linear;
...
```

If a particular ImageOutput implementation is required (by the rules of the file format it writes) to have pixels in a particular color space, then it will convert the color values of your image to the right color space if it is not already in that space. For example, JPEG images must be in sRGB space, so if you declare your pixels to be Linear, the JPEG ImageOutput will convert to sRGB.

If you leave the linearity set to the default of UnknownLinearity, the values will not be transformed, since the plugin can't be sure that it's not in the correct space to begin with.

The linearity only describes color channels. An ImageOutput plugin will assume that alpha or depth (*z*) channels (designated by the `alpha_channel` and `z_channel` fields, respectively) always represent linear values and should never be transformed.

Arbitrary metadata

For all other metadata that you wish to save in the file, you can attach the data to the ImageSpec using the `attribute()` methods. These come in polymorphic varieties that allow you to attach an attribute name and a value consisting of a single int, unsigned int, float, char*, or std::string, as shown in the following examples:

```
ImageSpec spec (...);
...

unsigned int u = 1;
spec.attribute ("Orientation", u);

float x = 72.0;
spec.attribute ("dotsize", f);

std::string s = "Fabulous image writer 1.0";
spec.attribute ("Software", s);
```

These are convenience routines for metadata that consist of a single value of one of these common types. For other data types, or more complex arrangements, you can use the more general form of `attribute()`, which takes arguments giving the name, type (as a TypeDesc), number of values (1 for a single value, > 1 for an array), and then a pointer to the data values. For example,

```
ImageSpec spec (...);

// Attach a 4x4 matrix to describe the camera coordinates
float mymatrix[16] = { ... };
spec.attribute ("worldtocamera", TypeDesc::TypeMatrix, &mymatrix);
```

```
// Attach an array of two floats giving the CIE neutral color
float neutral[2] = { ... };
spec.attribute ("adoptedNeutral", TypeDesc(TypeDesc::FLOAT, 2), &neutral);
```

In general, most image file formats (and therefore most ImageOutput implementations) are aware of only a small number of name/value pairs that they predefine and will recognize. Some file formats (OpenEXR, notably) do accept arbitrary user data and save it in the image file. If an ImageOutput does not recognize your metadata and does not support arbitrary metadata, that metadatum will be silently ignored and will not be saved with the file.

Each individual ImageOutput implementation should document the names, types, and meanings of all metadata attributes that they understand.

3.2.6 Controlling quantization

It is possible that your internal data format (that in which you compute pixel values that you pass to the write functions) is of greater precision or range than the native data format of the output file. This can occur either because you specified a lower-precision data format in the ImageSpec that you passed to `open()`, or else that the image file format dictates a particular data format that does not match your internal format. For example, you may compute float pixels and pass those to `write_image()`, but if you are writing a JPEG/JFIF file, the values must be stored in the file as 8-bit unsigned integers.

The conversion from floating-point formats to integer formats (or from higher to lower integer, which is done by first converting to float) is controlled by five fields within the ImageSpec: `quant_black`, `quant_white`, `quant_min`, `quant_max`, and `quant_dither`. Float 0.0 maps to the integer value given by `quant_black`, and float 1.0 maps to the integer value given by `quant_white`. Then, for color channels only (not alpha or depth), a random amount is added in the range $(-\text{quant_dither}.. \text{quant_dither})$, in order to reduce banding artifacts. The result is then clamped to lie within the range of `quant_min` and `quant_max`, inclusive. Finally, this result is truncated its integer value for final output. Here is the code that implements this transformation (T is the final output integer type):

```
float value = quant_black * (1 - input) + quant_white * input;
if (it's a color channel)
    value += quant_dither * (2 * random() - 1);
T output = (T) clamp ((int)(value + 0.5), quant_min, quant_max);
```

The values of the quantization parameters are set in one of three ways: (1) upon construction of the ImageSpec, they are set to the default quantization values for the given data format; (2) upon call to `ImageSpec::set_format()`, the quantization values are set to the defaults for the given data format; (3) or, after being first set up in this manner, you may manually change the quantization parameters in the ImageSpec, if you want something other than the default quantization.

Default quantization for each integer type is as follows:

Data Format	black	white	min	max	dither
UINT8	0	255	0	255	0.5
INT8	0	127	-128	127	0.5
UINT16	0	65535	0	65535	0.5
INT16	0	32767	-32768	32767	0.5
UINT	0	4294967295	0	4294967295	0.5
INT	0	2147483647	-2147483648	2147483647	0.5
FLOAT					
HALF	0	1	N/A	N/A	0
DOUBLE					

Note that the default is to use the entire positive range of each integer type to represent the floating-point (0..1) range. Floating-point types do not attempt to remap values, do not add dither, and do not clamp (except to their full floating-point range).

The default will almost always be what you want. But just as an example, here's how you would specify a quantization for a 16-bit file in which 1.0 maps to 16383 (14 bits of positive range) rather than filling the full 16 bit:

```
ImageSpec spec (width, length, channels, TypeDesc::UINT16);
spec.quant_black = 0;
spec.quant_white = 16383;
spec.quant_min = 0;
spec.quant_max = 16383;
spec.quant_dither = 0.5;
```

3.2.7 Random access and repeated transmission of pixels

All ImageOutput implementations that support scanlines and tiles should write pixels in strict order of increasing z slice, increasing y scanlines/rows within each slice, and increasing x column within each row. It is generally not safe to skip scanlines or tiles, or transmit them out of order, unless the plugin specifically advertises that it supports random access or rewrites, which may be queried using:

```
ImageOutput *out = ImageOutput::create (filename);
if (out->supports ("random_access"))
    ...
```

Similarly, you should assume the plugin will not correctly handle repeated transmissions of a scanline or tile that has already been sent, unless it advertises that it supports rewrites, which may be queried using:

```
if (out->supports ("rewrite"))
    ...
```

3.2.8 Multi-image files and MIP-maps

Some image file formats support multiple discrete subimages to be stored in one file, and/or multiple resolutions (MIP-map levels). Given a created ImageOutput, you can query whether multiple images may be stored in the file:

```

ImageOutput *out = ImageOutput::create (filename);
if (out->supports ("multiimage"))
    ...
if (out->supports ("mipmap"))
    ...

```

If you are working with an ImageOutput that supports multiple images or MIP-map levels, it is easy to write these images. All you have to do is, after writing all the pixels of one image but before calling `close()`, call `open()` again for the next subimage or MIP level and passing the appropriate value for the optional third *mode* argument. (See Section 3.3 for the full technical description of the arguments to `open()`.) The `close()` routine is called just once, after all subimages and MIP levels are completed.

Below is pseudocode for writing a MIP-map (a multi-resolution image used for texture mapping):

```

const char *filename = "foo.tif";
const int xres = 512, yres = 512;
const int channels = 3; // RGB
unsigned char *pixels = new unsigned char [xres*yres*channels];

// Create the ImageOutput
ImageOutput *out = ImageOutput::create (filename);

// Be sure we can support either mipmaps or subimages
if (! out->supports ("mipmap") && ! out->supports ("multiimage")) {
    std::cerr << "Cannot write a MIP-map\n";
    delete out;
    return;
}

// Set up spec for the highest resolution
ImageSpec spec (xres, yres, channels, TypeDesc::UINT8);

// Use Create mode for the first level.
ImageOutput::OpenMode appendmode = ImageOutput::Create;

// Write images, halving every time, until we're down to
// 1 pixel in either dimension
while (spec.width >= 1 && spec.height >= 1) {
    out->open (filename, spec, mode);
    out->write_image (TypeDesc::UINT8, pixels);
    // Assume halve() resamples the image to half resolution
    halve (pixels, spec.width, spec.height);
    // Don't forget to change spec for the next iteration
    spec.width /= 2;
    spec.height /= 2;
    // For subsequent levels, change the append mode argument to
    // open(). If the format doesn't support MIPmaps directly,
    // try to emulate it with subimages.
    if (out->supports("mipmap"))
        appendmode = ImageOutput::AppendMIPLevel;
}

```

```

        else
            appendmode = ImageOutput::AppendSubimage;
    }
    out->close ();
    delete out;

```

In this example, we have used `write_image()`, but of course `write_scanline()`, `write_tile()`, and `write_rectangle()` work as you would expect, on the current subimage.

3.2.9 Per-channel formats

Some image formats allow separate per-channel data formats (for example, half data for colors and float data for depth). When this is desired, the following steps are necessary:

1. Verify that the writer supports per-channel formats by checking `supports ("channelformats")`.
2. The `ImageSpec` passed to `open()` should have its `channelformats` vector filled with the types for each channel.
3. The call to `write_scanline`, `write_tile`, or `write_image` should pass a data pointer to the raw data, already in the native per-channel format of the file and contiguously packed, and specify that the data is of type `TypeDesc::UNKNOWN`.

For example, the following code fragment will write a 5-channel image to an OpenEXR file, consisting of R/G/B/A channels in half and a Z channel in float:

```

// Mixed data type for the pixel
struct Pixel { half r,g,b,a; float z; };
Pixel pixels[xres*yres];

ImageOutput *out = ImageOutput::create ("foo.exr");

// Double check that this format accepts per-channel formats
if (! out->supports("channelformats")) {
    delete out;
    return;
}

// Prepare an ImageSpec with per-channel formats
ImageSpec spec (xres, yres, 5, TypeDesc::FLOAT);
spec.channelformats.push_back (TypeDesc::HALF);
spec.channelformats.push_back (TypeDesc::HALF);
spec.channelformats.push_back (TypeDesc::HALF);
spec.channelformats.push_back (TypeDesc::HALF);
spec.channelformats.push_back (TypeDesc::HALF);
spec.channelformats.push_back (TypeDesc::FLOAT);
spec.channelnames.clear ();
spec.channelnames.push_back ("R");
spec.channelnames.push_back ("G");
spec.channelnames.push_back ("B");

```

```

spec.channelnames.push_back ("A");
spec.channelnames.push_back ("Z");

out->open (filename, spec);
out->write_image (TypeDesc::UNKNOWN, /* use channel formats */
                 pixels,             /* data buffer */
                 sizeof(Pixel));     /* pixel stride */

```

3.2.10 Copying an entire image

Suppose you want to copy an image, perhaps with alterations to the metadata but not to the pixels. You could open an `ImageInput` and perform a `read_image()`, and open another `ImageOutput` and call `write_image()` to output the pixels from the input image. However, for compressed images, this may be inefficient due to the unnecessary decompression and subsequent re-compression. In addition, if the compression is *lossy*, the output image may not contain pixel values identical to the original input.

A special `copy_image` method of `ImageOutput` is available that attempts to copy an image from an open `ImageInput` (of the same format) to the output as efficiently as possible without altering pixel values, if at all possible.

Not all format plugins will provide an implementation of `copy_image` (in fact, most will not), but the default implementation simply copies pixels one scanline or tile at a time (with decompression/recompression) so it's still safe to call. Furthermore, even a provided `copy_image` is expected to fall back on the default implementation if the input and output are not able to do an efficient copy. Nevertheless, this method is recommended for copying images so that maximal advantage will be taken in cases where savings can be had.

The following is an example use of `copy_image` to transfer pixels without alteration while modifying the image description metadata:

```

// Open the input file
const char *input = "input.jpg";
ImageInput *in = ImageInput::create (input);
ImageSpec in_spec;
in->open (input, in_spec);

// Make an output spec, identical to the input except for metadata
ImageSpec out_spec = in_spec;
out_spec.attribute ("ImageDescription", "My Title");

// Create the output file and copy the image
const char *output = "output.jpg";
ImageOutput *out = ImageOutput::create (output);
out->open (output, out_spec);
out->copy_image (in);

// Clean up
out->close ();
delete out;
in->close ();
delete in;

```

3.2.11 Custom search paths for plugins

When you call `ImageOutput::create()`, the `OpenImageIO` library will try to find a plugin that is able to write the format implied by your filename. These plugins are alternately known as DLL's on Windows (with the `.dll` extension), DSO's on Linux (with the `.so` extension), and dynamic libraries on Mac OS X (with the `.dylib` extension).

`OpenImageIO` will look for matching plugins according to *search paths*, which are strings giving a list of directories to search, with each directory separated by a colon (`:`). Within a search path, any substrings of the form `${FOO}` will be replaced by the value of environment variable `FOO`. For example, the searchpath `"${HOME}/plugins:/shared/plugins"` will first check the directory `"/home/tom/plugins"` (assuming the user's home directory is `/home/tom`), and if not found there, will then check the directory `"/shared/plugins"`.

The first search path it will check is that stored in the environment variable `IMAGEIO_LIBRARY_PATH`. It will check each directory in turn, in the order that they are listed in the variable. If no adequate plugin is found in any of the directories listed in this environment variable, then it will check the custom searchpath passed as the optional second argument to `ImageOutput::create()`, searching in the order that the directories are listed. Here is an example:

```
char *mysearch = "/usr/myapp/lib:${HOME}/plugins";
ImageOutput *out = ImageOutput::create (filename, mysearch);
...
```

3.2.12 Error checking

Nearly every `ImageOutput` API function returns a `bool` indicating whether the operation succeeded (`true`) or failed (`false`). In the case of a failure, the `ImageOutput` will have saved an error message describing in more detail what went wrong, and the latest error message is accessible using the `ImageOutput` method `geterror()`, which returns the message as a `std::string`.

The exception to this rule is `ImageOutput::create`, which returns `NULL` if it could not create an appropriate `ImageOutput`. And in this case, since no `ImageOutput` exists for which you can call its `geterror()` function, there exists a global `geterror()` function (in the `OpenImageIO` namespace) that retrieves the latest error message resulting from a call to `create`.

Here is another version of the simple image writing code from Section 3.1, but this time it is fully elaborated with error checking and reporting:

```
#include "imageio.h"
using namespace OpenImageIO;
...

const char *filename = "foo.jpg";
const int xres = 640, yres = 480;
const int channels = 3; // RGB
unsigned char pixels[xres*yres*channels];

ImageOutput *out = ImageOutput::create (filename);
if (! out) {
```

```

        std::cerr << "Could not create an ImageOutput for "
        << filename << ", error = "
        << OpenImageIO::geterror() << "\n";
    return;
}
ImageSpec spec (xres, yres, channels, TypeDesc::UINT8);

if (! out->open (filename, spec)) {
    std::cerr << "Could not open " << filename
    << ", error = " << out->geterror() << "\n";
    delete out;
    return;
}

if (! out->write_image (TypeDesc::UINT8, pixels)) {
    std::cerr << "Could not write pixels to " << filename
    << ", error = " << out->geterror() << "\n";
    delete out;
    return;
}

if (! out->close ()) {
    std::cerr << "Error closing " << filename
    << ", error = " << out->geterror() << "\n";
    delete out;
    return;
}

delete out;

```

3.3 ImageOutput Class Reference

```

static ImageOutput * create (const std::string &filename,
                           const std::string &plugin_searchpath="")

```

Create an ImageOutput that can be used to write an image file. The type of image file (and hence, the particular subclass of ImageOutput returned, and the plugin that contains its methods) is inferred from the extension of the file name. The `plugin_searchpath` parameter is a colon-separated list of directories to search for OpenImageIO plugin DSO/DLL's.

```

const char * format_name ()

```

Returns the canonical name of the format that this ImageOutput instance is capable of writing.

```

bool supports (const std::string &feature)

```


Given the name of a *feature*, tells if this ImageOutput instance supports that feature. The following features are recognized by this query:

- "tiles" Is this plugin able to write tiled images?
- "rectangles" Can this plugin accept arbitrary rectangular pixel regions (via `write_rectangle()`)? False indicates that pixels must be transmitted via `write_scanline()` (if scanline-oriented) or `write_tile()` (if tile-oriented, and only if `supports("tiles")` returns true).
- "random_access" May tiles or scanlines be written in any order? False indicates that they must be in successive order.
- "multiimage" Does this format support multiple subimages within a single file?
- "mipmap" Does this format support resolutions per image/subimage (MIP-map levels)?
- "volumes" Does this format support "3D" pixel arrays (a.k.a. volume images)?
- "channelformats" Does this format writer support per-channel data formats, respecting the ImageSpec's `channelformats` field? (If not, it only accepts a single data format for all channels and will ignore the `channelformats` field of the spec.)
- "rewrite" Does this plugin allow the same scanline or tile to be sent more than once? Generally this is true for plugins that implement some sort of interactive display, rather than a saved image file.
- "empty" Does this plugin support passing a NULL data pointer to the various `write` routines to indicate that the entire data block is composed of pixels with value zero. Plugins that support this achieve a speedup when passing blank scanlines or tiles (since no actual data needs to be transmitted or converted).

This list of queries may be extended in future releases. Since this can be done simply by recognizing new query strings, and does not require any new API entry points, addition of support for new queries does not break "link compatibility" with previously-compiled plugins.

```
bool open (const std::string &name, const ImageSpec &newspec,
          OpenMode mode=Create)
```

Open the file with given name, with resolution, and other format data as given in `newspec`. This function returns `true` for success, `false` for failure. Note that it is legal to call `open()` multiple times on the same file without a call to `close()`, if it supports multiimage and mode is `AppendSubimage`, or if it supports MIP-maps and mode is `AppendMIPlevel` – this is interpreted as appending a subimage, or a MIP level to the current subimage, respectively.

```
const ImageSpec & spec ()
```

Returns the spec internally associated with this currently open ImageOutput.

```
bool close ()
```

Closes the currently open file associated with this ImageOutput and frees any memory or resources associated with it.

```
bool write_scanline (int y, int z, TypeDesc format, const void *data,
                    stride_t xstride=AutoStride)
```

Write a full scanline that includes pixels $(*, y, z)$. For 2D non-volume images, z is ignored. The `xstride` value gives the distance between successive pixels (in bytes). Strides set to the special value `AutoStride` imply contiguous data, i.e.,

```
xstride = spec.nchannels*format.size()
```

This method automatically converts the data from the specified format to the actual output format of the file. If `format` is `TypeDesc::UNKNOWN`, the data is assumed to already be in the file's native format (including per-channel formats, as specified in the ImageSpec's `channelformats` field, if applicable). Return `true` for success, `false` for failure. It is a failure to call `write_scanline()` with an out-of-order scanline if this format driver does not support random access.

```
bool write_tile (int x, int y, int z, TypeDesc format, const void *data,
                stride_t xstride=AutoStride, stride_t ystride=AutoStride,
                stride_t zstride=AutoStride)
```

Write the tile with (x, y, z) as the upper left corner. For 2D non-volume images, z is ignored. The three stride values give the distance (in bytes) between successive pixels, scanlines, and volumetric slices, respectively. Strides set to the special value `AutoStride` imply contiguous data, i.e.,

```
xstride = spec.nchannels*format.size()
```

```
ystride = xstride*spec.tile_width
```

```
zstride = ystride*spec.tile_height
```

This method automatically converts the data from the specified format to the actual output format of the file. If `format` is `TypeDesc::UNKNOWN`, the data is assumed to already be in the file's native format (including per-channel formats, as specified in the ImageSpec's `channelformats` field, if applicable). Return `true` for success, `false` for failure. It is a failure to call `write_tile()` with an out-of-order tile if this format driver does not support random access.

```
bool write_rectangle (int xmin, int xmax, int ymin, int ymax, int zmin, int zmax,
                     TypeDesc format, const void *data,
                     stride_t xstride=AutoStride, stride_t ystride=AutoStride,
                     stride_t zstride=AutoStride)
```

Write pixels whose x coords range over $xmin...xmax$ (inclusive), y coords over $ymin...ymax$, and z coords over $zmin...zmax$. The three stride values give the distance (in bytes) between successive pixels, scanlines, and volumetric slices, respectively. Strides set to the special value `AutoStride` imply contiguous data, i.e.,

```
xstride = spec.nchannels*format.size()
```

```

    ystride = xstride*(xmax-xmin+1)
    zstride = ystride*(ymax-ymin+1)

```

This method automatically converts the data from the specified format to the actual output format of the file. If format is `TypeDesc::UNKNOWN`, the data is assumed to already be in the file's native format (including per-channel formats, as specified in the `ImageSpec`'s `channelformats` field, if applicable). Return `true` for success, `false` for failure. It is a failure to call `write_rectangle` for a format plugin that does not return `true` for `supports("rectangles")`.

```

bool write_image (TypeDesc format, const void *data,
    stride_t xstride=AutoStride, stride_t ystride=AutoStride,
    stride_t zstride=AutoStride,
    ProgressCallback progress_callback=NULL,
    void *progress_callback_data=NULL)

```

Write the entire image of `spec.width × spec.height × spec.depth` pixels, with the given strides and in the desired format. If format is `TypeDesc::UNKNOWN`, the data is assumed to already be in the file's native format (including per-channel formats, as specified in the `ImageSpec`'s `channelformats` field, if applicable). Strides set to the special value `AutoStride` imply contiguous data, i.e.,

```

    xstride = spec.nchannels * format.size()
    ystride = xstride * spec.width
    zstride = ystride * spec.height

```

The function will internally either call `write_scanline()` or `write_tile()`, depending on whether the file is scanline- or tile-oriented.

Because this may be an expensive operation, a progress callback may be passed. Periodically, it will be called as follows:

```

    progress_callback (progress_callback_data, float done)

```

where *done* gives the portion of the image (between 0.0 and 1.0) that has been written thus far.

```

bool copy_image (ImageInput *in)

```

Read the current subimage of `in`, and write it as the next subimage of `*this`, in a way that is efficient and does not alter pixel values, if at all possible. Both `in` and `this` must be a properly-opened `ImageInput` and `ImageOutput`, respectively, and their current images must match in size and number of channels. Return `true` if it works ok, `false` if for some reason the operation wasn't possible.

If a particular `ImageOutput` implementation does not supply a `copy_image` method, it will inherit the default implementation, which is to simply read scanlines or tiles from `in` and write them to `*this`. However, some format implementations may have a special technique for directly copying raw pixel data from the input to the output, when both input and output are the same file type and the same data format. This can be more efficient than `in->read_image` followed by `out->write_image`, and avoids any unintended pixel alterations, especially for formats that use lossy compression.

```
int send_to_output (const char *format, ...)
```

General message passing between client and image output server. This is currently undefined and is reserved for future use.

```
int send_to_client (const char *format, ...)
```

General message passing between client and image output server. This is currently undefined and is reserved for future use.

```
std::string geterror ()
```

Returns the current error string describing what went wrong if any of the public methods returned `false` indicating an error. (Hopefully the implementation plugin called `error()` with a helpful error message.)

4 Image I/O: Reading Images

4.1 Image Input Made Simple

Here is the simplest sequence required to open an image file, find out its resolution, and read the pixels (converting them into 8-bit values in memory, even if that's not the way they're stored in the file):

```
#include "imageio.h"
using namespace OpenImageIO;
...

const char *filename = "foo.jpg";
int xres, yres, channels;
unsigned char *pixels;

ImageInput *in = ImageInput::create (filename);
if (! in)
    return;
ImageSpec spec;
in->open (filename, spec);
xres = spec.width;
yres = spec.height;
channels = spec.nchannels;
pixels = new unsigned char [xres*yres*channels];
in->read_image (TypeDesc::UINT8, pixels);
in->close ();
delete in;
```

Here is a breakdown of what work this code is doing:

- Search for an ImageIO plugin that is capable of reading the file ("foo.jpg"), first by trying to deduce the correct plugin from the file extension, but if that fails, by opening every ImageIO plugin it can find until one will open the file without error. When it finds the right plugin, it creates a subclass instance of ImageInput that reads the right kind of file format.

```
ImageInput *in = ImageInput::create (filename);
```

- Open the file, read the header, and put all relevant metadata about the file in a specification structure.

```
ImageSpec spec;  
in->open (filename, spec);
```

- The specification contains vital information such as the dimensions of the image, number of color channels, and data type of the pixel values. This is enough to allow us to allocate enough space for the image.

```
xres = spec.width;  
yres = spec.height;  
channels = spec.nchannels;  
pixels = new unsigned char [xres*yres*channels];
```

Note that in this example, we don't care what data format is used for the pixel data in the file — we allocate enough space for unsigned 8-bit integer pixel values, and will rely on OpenImageIO's ability to convert to our requested format from the native data format of the file.

- Read the entire image, hiding all details of the encoding of image data in the file, whether the file is scanline- or tile-based, or what is the native format of the data in the file (in this case, we request that it be automatically converted to unsigned 8-bit integers).

```
in->read_image (TypeDesc::UINT8, pixels);
```

- Close the file, destroy and free the ImageInput we had created, and perform all other cleanup and release of any resources used by the plugin.

```
in->close ();  
delete in;
```

4.2 Advanced Image Input

Let's walk through some of the most common things you might want to do, but that are more complex than the simple example above.

4.2.1 Reading individual scanlines and tiles

The simple example of Section 4.1 read an entire image with one call. But sometimes you want to read a large image a little at a time and do not wish to retain the entire image in memory as you process it. OpenImageIO allows you to read images one scanline at a time or one tile at a time.

Reading individual scanlines

Individual scanlines may be read using the `read_scanline()` API call:

```
...
in->open (filename, spec);
unsigned char *scanline = new unsigned char [spec.width*spec.channels];
for (int y = 0; y < yres; ++y) {
    in->read_scanline (y, 0, TypeDesc::UINT8, scanline);
    ... process data in scanline[0..width*channels-1] ...
}
delete [] scanline;
in->close ();
...
```

The first two arguments to `read_scanline()` specify which scanline is being read by its vertical (*y*) scanline number (beginning with 0) and, for volume images, its slice (*z*) number (the slice number should be 0 for 2D non-volume images). This is followed by a `TypeDesc` describing the data type of the pixel buffer you are supplying, and a pointer to the pixel buffer itself. Additional optional arguments describe the data stride, which can be ignored for contiguous data (use of strides is explained in Section 4.2.3).

All `ImageInput` implementations will read scanlines in strict order (starting with scanline 0, then 1, up to `yres-1`, without skipping any). However, it's probably uncommon for an `ImageInput` to properly allow reading of out-of-order scanlines.

The full description of the `read_scanline()` function may be found in Section 4.3.

Reading individual tiles

Once you `open()` an image file, you can find out if it is a tiled image (and the tile size) by examining the `ImageSpec`'s `tile_width`, `tile_height`, and `tile_depth` fields. If they are zero, it's a scanline image and you should read pixels using `read_scanline()`, not `read_tile()`.

```
...
in->open (filename, spec);
if (spec.tile_width == 0) {
    ... read by scanline ...
} else {
    // Tiles
    int tilesize = spec.tile_width * spec.tile_height;
    unsigned char *tile = new unsigned char [tilesize * spec.channels];
    for (int y = 0; y < yres; y += spec.tile_height) {
        for (int x = 0; x < xres; x += spec.tile_width) {
            in->read_tile (x, y, 0, TypeDesc::UINT8, tile);
            ... process the pixels in tile[] ..
        }
    }
    delete [] tile;
}
in->close ();
...
```

The first three arguments to `read_tile()` specify which tile is being read by the pixel coordinates of any pixel contained in the tile: x (column), y (scanline), and z (slice, which should always be 0 for 2D non-volume images). This is followed by a `TypeDesc` describing the data format of the pixel buffer you are supplying, and a pointer to the pixel buffer. Pixel data will be written to your buffer in order of increasing slice, increasing scanline within each slice, and increasing column within each scanline. Additional optional arguments describe the data stride, which can be ignored for contiguous data (use of strides is explained in Section 4.2.3).

All `ImageInput` implementations are required to support reading tiles in arbitrary order (i.e., not in strict order of increasing y rows, and within each row, increasing x column, without missing any tiles).

The full description of the `read_tile()` function may be found in Section 4.3.

4.2.2 Converting formats

The code examples of the previous sections all assumed that your internal pixel data is stored as unsigned 8-bit integers (i.e., 0-255 range). But `OpenImageIO` is significantly more flexible.

You may request that the pixels be stored in any of several formats. This is done merely by passing the read function the data type of your pixel buffer, as one of the enumerated type `TypeDesc`.

It is not required that the pixel data buffer passed to `read_image()`, `read_scanline()`, or `read_tile()` actually be in the same data format as the data in the file being read. `OpenImageIO` will automatically convert from native data type of the file to the internal data format of your choice. For example, the following code will open a TIFF and read pixels into your internal buffer represented as `float` values. This will work regardless of whether the TIFF file itself is using 8-bit, 16-bit, or float values.

```
ImageInput *in = ImageInput::create ("myfile.tif");
ImageSpec spec;
in->open (filename, spec);
...
int numpixels = spec.width * spec.height;
float pixels = new float [numpixels * channels];
...
in->read_image (TypeDesc::FLOAT, pixels);
```

Note that `read_scanline()` and `read_tile()` have a parameter that works in a corresponding manner.

You can, of course, find out the native type of the file simply by examining `spec.format`. If you wish, you may then allocate a buffer big enough for an image of that type and request the native type when reading, therefore eliminating any translation among types and seeing the actual numerical values in the file.

4.2.3 Data Strides

In the preceding examples, we have assumed that the buffer passed to the read functions (i.e., the place where you want your pixels to be stored) is *contiguous*, that is:

- each pixel in memory consists of a number of data values equal to the number of channels in the file;
- successive column pixels within a row directly follow each other in memory, with the first channel of pixel x immediately following last channel of pixel $x - 1$ of the same row;
- for whole images or tiles, the data for each row immediately follows the previous one in memory (the first pixel of row y immediately follows the last column of row $y - 1$);
- for 3D volumetric images, the first pixel of slice z immediately follows the last pixel of slice $z - 1$.

Please note that this implies that `read_tile()` will write pixel data into your buffer so that it is contiguous in the shape of a single tile, not just an offset into a whole image worth of pixels.

The `read_scanline()` function takes an optional `xstride` argument, and the `read_image()` and `read_tile()` functions take optional `xstride`, `ystride`, and `zstride` values that describe the distance, in *bytes*, between successive pixel columns, rows, and slices, respectively, of your pixel buffer. For any of these values that are not supplied, or are given as the special constant `AutoStride`, contiguity will be assumed.

By passing different stride values, you can achieve some surprisingly flexible functionality. A few representative examples follow:

- Flip an image vertically upon reading, by using *negative y* stride:

```
unsigned char pixels[spec.width * spec.height * spec.nchannels];
int scanlinesize = spec.width * spec.nchannels * sizeof(pixels[0]);
...
in->read_image (TypeDesc::UINT8,
                (char *)pixels+(yres-1)*scanlinesize, // offset to last
                AutoStride,                          // default x stride
                -scanlinesize,                        // special y stride
                AutoStride);                          // default z stride
```

- Read a tile into its spot in a buffer whose layout matches a whole image of pixel data, rather than having a one-tile-only memory layout:

```
unsigned char pixels[spec.width * spec.height * spec.nchannels];
int pixelsize = spec.nchannels * sizeof(pixels[0]);
int scanlinesize = spec.width * pixelsize;
...
in->read_tile (x, y, 0, TypeDesc::UINT8,
              (char *)pixels + y*scanlinesize + x*pixelsize,
              pixelsize,
              scanlinesize);
```

Please consult Section 4.3 for detailed descriptions of the stride parameters to each read function.

4.2.4 Reading metadata

The `ImageSpec` that is filled in by `ImageInput::open()` specifies all the common properties that describe an image: data format, dimensions, number of channels, tiling. However, there may be a variety of additional *metadata* that are present in the image file and could be queried by your application.

The remainder of this section explains how to query additional metadata in the `ImageSpec`. It is up to the `ImageInput` to read these from the file, if indeed the file format is able to carry additional data. Individual `ImageInput` implementations should document which metadata they read.

Channel names

In addition to specifying the number of color channels, the `ImageSpec` also stores the names of those channels in its `channelnames` field, which is a `vector<std::string>`. Its length should always be equal to the number of channels (it's the responsibility of the `ImageInput` to ensure this).

Only a few file formats (and thus `ImageInput` implementations) have a way of specifying custom channel names, so most of the time you will see that the channel names follow the default convention of being named "R", "G", "B", and "A", for red, green, blue, and alpha, respectively.

Here is example code that prints the names of the channels in an image:

```
ImageInput *in = ImageInput::create (filename);
ImageSpec spec;
in->open (filename, spec);
for (int i = 0; i < spec.nchannels; ++i)
    std::cout << "Channel " << i << " is "
               << spec.channelnames[i] << "\n";
```

Specially-designated channels

The `ImageSpec` contains two fields, `alpha_channel` and `z_channel`, which designate which channel numbers represent alpha and *z* depth, if any. If either is set to `-1`, it indicates that it is not known which channel is used for that data.

If you are doing something special with alpha or depth, it is probably safer to respect the `alpha_channel` and `z_channel` designations (if not set to `-1`) rather than merely assuming that, for example, channel 3 is always the alpha channel.

Linearity hints

We certainly hope that you are using only modern file formats that support high precision and extended range pixels (such as OpenEXR) and keeping all your images in a linear color space. But you may have to work with file formats that dictate the use of nonlinear color values. This is prevalent in formats that store pixels only as 8-bit values, since 256 values are not enough to linearly represent colors without banding artifacts in the dim values.

The `ImageSpec` has a field that reveals what color space the image file is using. Each individual `ImageInput` is responsible for setting this properly.

The `ImageSpec` field `linearity` can take on any of the following values:

`ImageSpec::UnknownLinearity` indicates indicates it's unknown what color space the image file is using.

`ImageSpec::Linear` indicates that the color pixel values are known to be linear.

`ImageSpec::GammaCorrected` indicates that the color pixel values (but not alpha or z) in the file have already been gamma corrected (raised to the power $1/\gamma$), and that the gamma exponent may be found in the `gamma` field of the `ImageSpec`.

`ImageSpec::sRGB` indicates that the color pixel values in the file are in sRGB color space.

`ImageSpec::AdobeRGB` indicates that the color pixel values in the file are in Adobe RGB color space.

`ImageSpec::Rec709` indicates that the color pixel values in the file are in Rec709 color space.

`ImageSpec::Kodak` indicates that the color pixel values in the file are in Kodak logarithmic color space.

The `ImageInput` sets the `ImageSpec::linearity` field in a purely advisory capacity — the read will not convert pixel values among color spaces. Many image file formats only support nonlinear color spaces (for example, JPEG/JFIF dictates use of sRGB). So your application should intelligently deal with gamma-corrected and sRGB input.

The linearity only describes color channels. You should assume that alpha or depth (z) channels (designated by the `alpha_channel` and `z_channel` fields, respectively) always represent linear values and should never be transformed by your application.

Arbitrary metadata

All other metadata found in the file will be stored in the `ImageSpec`'s `extra_attribs` field, which is a `ParamValueList`, which is itself essentially a vector of `ParamValue` instances. Each `ParamValue` stores one meta-datum consisting of a name, type (specified by a `TypeDesc`), number of values, and data pointer.

If you know the name of a specific piece of metadata you want to use, you can find it using the `ImageSpec::find_attribute()` method, which returns a pointer to the matching `ParamValue`, or `NULL` if no match was found. An optional `TypeDesc` argument can narrow the search to only parameters that match the specified type as well as the name. Below is an example that looks for orientation information, expecting it to consist of a single integer:

```
ImageInput *in = ImageInput::create (filename);
ImageSpec spec;
in->open (filename, spec);
...
ParamValue *p = spec.find_attribute ("Orientation", TypeDesc::INT);
```

```

if (p) {
    int orientation = * (int *) p->data();
} else {
    std::cout << "No integer orientation in the file\n";
}

```

By convention, ImageInput plugins will save all integer metadata as 32-bit integers (TypeDesc::INT or TypeDesc::UINT), even if the file format dictates that a particular item is stored in the file as a 8- or 16-bit integer. This is just to keep client applications from having to deal with all the types. Since there is relatively little metadata compared to pixel data, there's no real memory waste of promoting all integer types to int32 metadata. Floating-point metadata and string metadata may also exist, of course.

It is also possible to step through all the metadata, item by item. This can be accomplished using the technique of the following example:

```

for (size_t i = 0; i < spec.extra_attribs.size(); ++i) {
    const ParamValue &p (spec.extra_attribs[i]);
    printf ("    %s: ", p.name.c_str());
    if (p.type() == TypeDesc::STRING)
        printf ("\"%s\"", *(const char **)p.data());
    else if (p.type() == TypeDesc::FLOAT)
        printf ("%g", *(const float *)p.data());
    else if (p.type() == TypeDesc::INT)
        printf ("%d", *(const int *)p.data());
    else if (p.type() == TypeDesc::UINT)
        printf ("%u", *(const unsigned int *)p.data());
    else
        printf ("<unknown data type>");
    printf ("\n");
}

```

Each individual ImageInput implementation should document the names, types, and meanings of all metadata attributes that they understand.

4.2.5 Multi-image files and MIP-maps

Some image file formats support multiple discrete subimages to be stored in one file, and/or multiple resolutions for each image to form a MIPmap. When you open() an ImageInput, it will by default point to the first (i.e., number 0) subimage in the file, and the highest resolution (level 0) MIP-map level. You can switch to viewing another subimage or MIP-map level using the seek_subimage() function:

```

ImageInput *in = ImageInput::create (filename);
ImageSpec spec;
in->open (filename, spec);
...
int subimage = 1;
int miplevel = 0;
if (in->seek_subimage (subimage, miplevel, spec)) {

```

```

    ...
} else {
    ... no such subimage/miplevel ...
}

```

The `seek_subimage()` function takes three arguments: the index of the subimage to switch to (starting with 0), the MIPmap level (starting with 0 for the highest-resolution level), and a reference to an `ImageSpec`, into which will be stored the spec of the new subimage/miplevel. The `seek_subimage()` function returns `true` upon success, and `false` if no such subimage or MIP level existed. It is legal to visit subimages and MIP levels out of order; the `ImageInput` is responsible for making it work properly. It is also possible to find out which subimage and MIP level is currently being viewed, using the `current_subimage()` and `current_miplevel()` functions, which return the index of the current subimage and MIP levels, respectively.

Below is pseudocode for reading all the levels of a MIP-map (a multi-resolution image used for texture mapping) that shows how to read multi-image files:

```

ImageInput *in = ImageInput::create (filename);
ImageSpec spec;
in->open (filename, spec);

int num_miplevels = 0;
while (in->seek_subimage (0, num_miplevels, spec)) {
    // Note: spec has the format of the current subimage/miplevel
    int npixels = spec.width * spec.height;
    int nchannels = spec.nchannels;
    unsigned char *pixels = new unsigned char [npixels * nchannels];
    in->read_image (TypeDesc::UINT8, pixels);

    ... do whatever you want with this level, in pixels ...

    delete [] pixels;
    ++num_miplevels;
}
// Note: we break out of the while loop when seek_subimage fails
// to find a next MIP level.

in->close ();
delete in;

```

In this example, we have used `read_image()`, but of course `read_scanline()` and `read_tile()` work as you would expect, on the current subimage and MIP level.

4.2.6 Per-channel formats

Some image formats allow separate per-channel data formats (for example, half data for colors and float data for depth). If you want to read the pixels in their true native per-channel formats, the following steps are necessary:

1. Check the `ImageSpec`'s `channelformats` vector. If non-empty, the channels in the file do not all have the same format.

2. When calling `read_scanline`, `read_tile`, or `read_image`, pass a format of `TypeDesc::UNKNOWN` to indicate that you would like the raw data in native per-channel format of the file written to your data buffer.

For example, the following code fragment will read a 5-channel image to an OpenEXR file, consisting of R/G/B/A channels in half and a Z channel in float:

```
ImageSpec spec;
ImageInput *in = ImageInput::create (filename);
in->open (filename, spec);

// Allocate enough space
unsigned char *pixels = new unsigned char [spec.image_bytes(true)];

in->read_image (TypeDesc::UNKNOWN, /* use native channel formats */
               pixels);           /* data buffer */

if (spec.channelformats.size() > 0) {
    ... the buffer contains packed data in the native
        per-channel formats ...
} else {
    ... the buffer contains all data per spec.format ...
}
```

4.2.7 Custom search paths for plugins

Please see Section 3.2.11 for discussion about search paths for finding plugins that implement `ImageOutput`.

In a similar fashion, calls to `ImageOutput::create()` will search for plugins in each directory listed in the environment variable `IMAGEIO_LIBRARY_PATH`, in the order that they are listed. If no adequate plugin is found, then it will check the custom searchpath passed as the optional second argument to `ImageInput::create()`. Here is an example:

```
char *mysearch = "/usr/myapp/lib:${HOME}/plugins";
ImageInput *in = ImageInput::create (filename, mysearch);
...
```

4.2.8 Error checking

Nearly every `ImageInput` API function returns a `bool` indicating whether the operation succeeded (`true`) or failed (`false`). In the case of a failure, the `ImageInput` will have saved an error message describing in more detail what went wrong, and the latest error message is accessible using the `ImageInput` method `geterror()`, which returns the message as a `std::string`.

The exception to this rule is `ImageInput::create`, which returns `NULL` if it could not create an appropriate `ImageInput`. And in this case, since no `ImageInput` exists for which you can call its `geterror()` function, there exists a global `geterror()` function (in the `OpenImageIO` namespace) that retrieves the latest error message resulting from a call to `create`.

Here is another version of the simple image reading code from Section 4.1, but this time it is fully elaborated with error checking and reporting:

```

#include "imageio.h"
using namespace OpenImageIO;
...

const char *filename = "foo.jpg";
int xres, yres, channels;
unsigned char *pixels;

ImageInput *in = ImageInput::create (filename);
if (! in) {
    std::cerr << "Could not create an ImageInput for "
                << filename << ", error = "
                << OpenImageIO::geterror() << "\n";
    return;
}

ImageSpec spec;
if (! in->open (filename, spec)) {
    std::cerr << "Could not open " << filename
                << ", error = " << in->geterror() << "\n";
    delete in;
    return;
}

xres = spec.width;
yres = spec.height;
channels = spec.nchannels;
pixels = new unsigned char [xres*yres*channels];

if (! in->read_image (TypeDesc::UINT8, pixels)) {
    std::cerr << "Could not read pixels from " << filename
                << ", error = " << in->geterror() << "\n";
    delete in;
    return;
}

if (! in->close ()) {
    std::cerr << "Error closing " << filename
                << ", error = " << in->geterror() << "\n";
    delete in;
    return;
}

delete in;

```

4.3 ImageInput Class Reference

```

ImageInput * create (const std::string &filename,
                    const std::string &plugin_searchpath="")

```

Create and return an ImageInput implementation that is able to read the given file.

The `plugin_searchpath` parameter is a colon-separated list of directories to search for OpenImageIO plugin DSO/DLL's (not a searchpath for the image itself!). This will actually just try every ImageIO plugin it can locate, until it finds one that's able to open the file without error. This just creates the `ImageInput`, it does not open the file.

```
const char * format_name (void) const
```

Return the name of the format implemented by this class.

```
bool open (const std::string &name, ImageSpec &newspec)
```

Opens the file with given name and seek to the first subimage in the file. Various file attributes are put in `newspec` and a copy is also saved internally to the `ImageInput` (retrievable via `spec()`). From examining `newspec` or `spec()`, you can discern the resolution, if it's tiled, number of channels, native data format, and other metadata about the image. Return `true` if the file was found and opened okay, otherwise `false`.

```
bool open (const std::string &name, ImageSpec &newspec,
           const ImageSpec &config)
```

Opens the file with given name, similarly to `open(name, newspec)`. However, in this version, any non-default fields of `config`, including metadata, will be taken to be configuration requests, preferences, or hints. The default implementation of `open(name, newspec, config)` will simply ignore `config` and calls the usual `open(name, newspec)`. But a plugin may choose to implement this version of `open` and respond in some way to the configuration requests. Supported configuration requests should be documented by each plugin.

```
const ImageSpec & spec (void) const
```

Returns a reference to the image format specification of the current subimage. Note that the contents of the `spec` are invalid before `open()` or after `close()`.

```
bool close ()
```

Closes an open image.

```
int current_subimage (void) const
```

Returns the index of the subimage that is currently being read. The first subimage (or the only subimage, if there is just one) is number 0.

```
bool seek_subimage (int subimage, int miplevel, ImageSpec &newspec)
```

Seek to the given subimage and MIP-map level within the open image file. The first subimage in the file has index 0, and for each subimage, the highest-resolution MIP level has index 0. Return `true` on success, `false` on failure (including that there is not a subimage or MIP level with those indices). The new subimage's vital statistics are put

in `newspec` (and also saved internally in a way that can be retrieved via `spec()`). The `ImageInput` is expected to give the appearance of random access to subimages and MIP levels — in other words, if it can't randomly seek to the given subimage or MIP level, it should transparently close, reopen, and sequentially read through prior subimages and levels.

```
bool read_scanline (int y, int z, TypeDesc format, void *data,
                   stride_t xstride=AutoStride)
```

Read the scanline that includes pixels $(*, y, z)$ into data ($z = 0$ for non-volume images), converting if necessary from the native data format of the file into the format specified. If format is `TypeDesc::UNKNOWN`, the data will be preserved in its native format (including per-channel formats, if applicable). The `xstride` value gives the data spacing of adjacent pixels (in bytes). Strides set to the special value `AutoStride` imply contiguous data, i.e.,

```
xstride = spec.nchannels * spec.pixel_size()
```

The `ImageInput` is expected to give the appearance of random access — in other words, if it can't randomly seek to the given scanline, it should transparently close, reopen, and sequentially read through prior scanlines. The base `ImageInput` class has a default implementation that calls `read_native_scanline()` and then does appropriate format conversion, so there's no reason for each format plugin to override this method.

```
bool read_scanline (int y, int z, float *data)
```

This simplified version of `read_scanline()` reads to contiguous float pixels.

```
bool read_tile (int x, int y, int z, TypeDesc format, void *data,
                stride_t xstride=AutoStride, stride_t ystride=AutoStride,
                stride_t zstride=AutoStride)
```

Read the tile that includes pixels $(*, y, z)$ into data ($z = 0$ for non-volume images), converting if necessary from the native data format of the file into the format specified. If format is `TypeDesc::UNKNOWN`, the data will be preserved in its native format (including per-channel formats, if applicable). The stride values give the data spacing of adjacent pixels, scanlines, and volumetric slices, respectively (measured in bytes). Strides set to the special value of `AutoStride` imply contiguous data, i.e.,

```
xstride = spec.nchannels * spec.pixel_size()
```

```
ystride = xstride * spec.tile_width
```

```
zstride = ystride * spec.tile_height
```

The `ImageInput` is expected to give the appearance of random access — in other words, if it can't randomly seek to the given tile, it should transparently close, reopen, and sequentially read through prior tiles. The base `ImageInput` class has a default implementation that calls `read_native_tile` and then does appropriate format conversion, so there's no reason for each format plugin to override this method.

```
bool read_tile (int x, int y, int z, float *data)
```

Simple version of `read_tile` that reads to contiguous float pixels.

```
bool read_image (TypeDesc format, void *data,
                stride_t xstride=AutoStride, stride_t ystride=AutoStride,
                stride_t zstride=AutoStride,
                ProgressCallback progress_callback=NULL,
                void *progress_callback_data=NULL)
```

Read the entire image of `spec.width * spec.height * spec.depth` pixels into `data` (which must already be sized large enough for the entire image) with the given strides, converting into the desired data format. If `format` is `TypeDesc::UNKNOWN`, the data will be preserved in its native format (including per-channel formats, if applicable). This function will automatically handle either tiles or scanlines in the file.

Strides set to the special value of `AutoStride` imply contiguous data, i.e.,

```
xstride = spec.nchannels * pixel_size()
ystride = xstride * spec.width
zstride = ystride * spec.height
```

The function will internally either call `read_scanline` or `read_tile`, depending on whether the file is scanline- or tile-oriented.

Because this may be an expensive operation, a progress callback may be passed. Periodically, it will be called as follows:

```
progress_callback (progress_callback_data, float done)
```

where *done* gives the portion of the image (between 0.0 and 1.0) that has been read thus far.

```
bool read_image (float *data)
```

Simple version of `read_image()` reads to contiguous float pixels.

```
bool read_native_scanline (int y, int z, void *data)
```

The `read_native_scanline()` function is just like `read_scanline()`, except that it keeps the data in the native format of the disk file and always reads into contiguous memory (no strides). It's up to the user to have enough space allocated and know what to do with the data. IT IS EXPECTED THAT EACH FORMAT PLUGIN WILL OVERRIDE THIS METHOD.

```
bool read_native_tile (int x, int y, int z, void *data)
```

The `read_native_tile()` function is just like `read_tile()`, except that it keeps the data in the native format of the disk file and always read into contiguous memory (no strides). It's up to the user to have enough space allocated and know what to do with the data. IT IS EXPECTED THAT EACH FORMAT PLUGIN WILL OVERRIDE THIS METHOD IF IT SUPPORTS TILED IMAGES.

```
int send_to_input (const char *format, ...)
```

General message passing between client and image input server. This is currently undefined and is reserved for future use.

```
int send_to_client (const char *format, ...)
```

General message passing between client and image input server. This is currently undefined and is reserved for future use.

```
std::string geterror () const
```

Returns the current error string describing what went wrong if any of the public methods returned `false` indicating an error. (Hopefully the implementation plugin called `error()` with a helpful error message.)

5 Writing ImageIO Plugins

5.1 Plugin Introduction

As explained in Chapters 4 and 3, the ImageIO library does not know how to read or write any particular image formats, but rather relies on plugins located and loaded dynamically at run-time. This set of plugins, and therefore the set of image file formats that OpenImageIO or its clients can read and write, is extensible without needing to modify OpenImageIO itself.

This chapter explains how to write your own OpenImageIO plugins. We will first explain separately how to write image file readers and writers, then tie up the loose ends of how to build the plugins themselves.

5.2 Image Readers

A plugin that reads a particular image file format must implement a *subclass* of `ImageInput` (described in Chapter 4). This is actually very straightforward and consists of the following steps, which we will illustrate with a real-world example of writing a JPEG/JFIF plug-in.

1. Read the base class definition from `imageio.h`. It may also be helpful to just use the `OpenImageIO` namespace, just to make your code a little less verbose.

```
#include "imageio.h"
using namespace OpenImageIO;
```

2. Declare three public items:
 - (a) An integer called `name_imageio_version` that identifies the version of the ImageIO protocol implemented by the plugin, defined in `imageio.h` as the constant `OPENIMAGEIO_PLUGIN_VERSION`. This allows the library to be sure it is not loading a plugin that was compiled against an incompatible version of OpenImageIO.
 - (b) A function named `name_input_imageio_create` that takes no arguments and returns a new instance of your `ImageInput` subclass. (Note that *name* is the name of your format, and must match the name of the plugin itself.)
 - (c) An array of `char *` called `name_input_extensions` that contains the list of file extensions that are likely to indicate a file of the right format. The list is terminated by a `NULL` pointer.

All of these items must be inside an ‘extern "C"’ block in order to avoid name mangling by the C++ compiler. Depending on your compiler, you may need to use special commands to dictate that the symbols will be exported in the DSO; we provide a special `DLLEXPORT` macro for this purpose, defined in `export.h`.

Putting this all together, we get the following for our JPEG example:

```
extern "C" {
    DLLEXPORT int jpeg_imageio_version = OPENIMAGEIO_PLUGIN_VERSION;
    DLLEXPORT JpgInput *jpeg_input_imageio_create () {
        return new JpgInput;
    }
    DLLEXPORT const char *jpeg_input_extensions[] = {
        "jpg", "jpe", "jpeg", NULL
    };
};
```

3. The definition and implementation of an `ImageInput` subclass for this file format. It must publicly inherit `ImageInput`, and must overload the following methods which are “pure virtual” in the `ImageInput` base class:

- (a) `format_name()` should return the name of the format, which ought to match the name of the plugin and by convention is strictly lower-case and contains no whitespace.
- (b) `open()` should open the file and return true, or should return false if unable to do so (including if the file was found but turned out not to be in the format that your plugin is trying to implement).
- (c) `close()` should close the file, if open.
- (d) `read_native_scanline` should read a single scanline from the file into the address provided, uncompressing it but keeping it in its native data format without any translation.
- (e) The virtual destructor, which should `close()` if the file is still open, addition to performing any other tear-down activities.

Additionally, your `ImageInput` subclass may optionally choose to overload any of the following methods, which are defined in the `ImageInput` base class and only need to be overloaded if the default behavior is not appropriate for your plugin:

- (f) `seek_subimage()`, only if your format supports reading multiple subimages within a single file.
- (g) `read_native_tile()`, only if your format supports reading tiled images.

Here is how the class definition looks for our JPEG example. Note that the JPEG/JFIF file format does not support multiple subimages or tiled images.

```

class JpgInput : public ImageInput {
public:
    JpgInput () { init(); }
    virtual ~JpgInput () { close(); }
    virtual const char * format_name (void) const { return "jpeg"; }
    virtual bool open (const std::string &name, ImageSpec &spec);
    virtual bool read_native_scanline (int y, int z, void *data);
    virtual bool close ();
private:
    FILE *m_fd;
    bool m_first_scanline;
    struct jpeg_decompress_struct m_cinfo;
    struct jpeg_error_mgr m_jerr;

    void init () { m_fd = NULL; }
};

```

Your subclass implementation of `open()`, `close()`, and `read_native_scanline()` are the heart of an `ImageInput` implementation. (Also `read_native_tile()` and `seek_subimage()`, for those image formats that support them.)

The remainder of this section simply lists the full implementation of our JPEG reader, which relies heavily on the open source `jpeg-6b` library to perform the actual JPEG decoding.

```

/*
Copyright 2008 Larry Gritz and the other authors and contributors.
All Rights Reserved.
Based on BSD-licensed software Copyright 2004 NVIDIA Corp.

Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions are
met:
* Redistributions of source code must retain the above copyright
  notice, this list of conditions and the following disclaimer.
* Redistributions in binary form must reproduce the above copyright
  notice, this list of conditions and the following disclaimer in the
  documentation and/or other materials provided with the distribution.
* Neither the name of the software's owners nor the names of its
  contributors may be used to endorse or promote products derived from
  this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS
"AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT
LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR
A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
OWNER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE

```

```

    OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

    (This is the Modified BSD License)
    */

#include <cassert>
#include <cstdio>

extern "C" {
#include "jpeglib.h"
}

#include "imageio.h"
#include "fmath.h"
#include "jpeg_pvt.h"

OIIO_PLUGIN_NAMESPACE_BEGIN

using namespace OpenImageIO;
using namespace Jpeg_imageio_pvt;

// See JPEG library documentation in /usr/share/doc/libjpeg-devel-6b

// N.B. The class definition for JpgInput is in jpeg_pvt.h.

// Export version number and create function symbols
OIIO_PLUGIN_EXPORTS_BEGIN

    DLLEXPORT int jpeg_imageio_version = OPENIMAGEIO_PLUGIN_VERSION;
    DLLEXPORT ImageInput *jpeg_input_imageio_create () {
        return new JpgInput;
    }
    DLLEXPORT const char *jpeg_input_extensions[] = {
        "jpg", "jpe", "jpeg", NULL
    };

OIIO_PLUGIN_EXPORTS_END

bool
JpgInput::open (const std::string &name, ImageSpec &newspec,
                const ImageSpec &config)
{
    const ImageIOParameter *p = config.find_attribute ("_jpeg:raw",
                                                        TypeDesc::TypeInt);

    m_raw = p && *(int *)p->data();
    return open (name, newspec);
}

```



```

bool
JpgInput::open (const std::string &name, ImageSpec &newspec)
{
    // Check that file exists and can be opened
    m_filename = name;
    m_fd = fopen (name.c_str(), "rb");
    if (m_fd == NULL) {
        error ("Could not open file \"%s\"", name.c_str());
        return false;
    }

    // Check magic number to assure this is a JPEG file
    int magic = 0;
    if (fread (&magic, 4, 1, m_fd) != 1) {
        error ("Empty file");
        return false;
    }

    rewind (m_fd);
    const int JPEG_MAGIC = 0xffd8ffe0, JPEG_MAGIC_OTHER_ENDIAN = 0xe0ffd8ff;
    const int JPEG_MAGIC2 = 0xffd8ffe1, JPEG_MAGIC2_OTHER_ENDIAN = 0xe1ffd8ff;
    if (magic != JPEG_MAGIC && magic != JPEG_MAGIC_OTHER_ENDIAN &&
        magic != JPEG_MAGIC2 && magic != JPEG_MAGIC2_OTHER_ENDIAN) {
        fclose (m_fd);
        m_fd = NULL;
        error ("\"%s\" is a JPEG file, magic number doesn't match", name.c_str());
        return false;
    }

    m_cinfo.err = jpeg_std_error (&m_jerr);
    jpeg_create_decompress (&m_cinfo);           // initialize decompressor
    jpeg_stdio_src (&m_cinfo, m_fd);             // specify the data source

    // Request saving of EXIF and other special tags for later spelunking
    for (int mark = 0; mark < 16; ++mark)
        jpeg_save_markers (&m_cinfo, JPEG_APP0+mark, 0xffff);
    jpeg_save_markers (&m_cinfo, JPEG_COM, 0xffff);    // comment marker

    jpeg_read_header (&m_cinfo, FALSE);             // read the file parameters
    if (m_raw)
        m_coeffs = jpeg_read_coefficients (&m_cinfo);
    else
        jpeg_start_decompress (&m_cinfo);           // start working
    m_next_scanline = 0;                            // next scanline we'll read

    m_spec = ImageSpec (m_cinfo.output_width, m_cinfo.output_height,
                        m_cinfo.output_components, TypeDesc::UINT8);

    for (jpeg_saved_marker_ptr m = m_cinfo.marker_list; m; m = m->next) {

```

```

        if (m->marker == (JPEG_APP0+1) &&
            ! strcmp ((const char *)m->data, "Exif"))
            decode_exif ((unsigned char *)m->data, m->data_length, m_spec);
        else if (m->marker == (JPEG_APP0+1) &&
            ! strcmp ((const char *)m->data, "http://ns.adobe.com/xap/1.0/")) {
#ifdef DEBUG
            std::cerr << "Found APP1 XMP! length " << m->data_length << "\n";
#endif
            std::string xml ((const char *)m->data, m->data_length);
            OpenImageIO::decode_xmp (xml, m_spec);
        }
        else if (m->marker == (JPEG_APP0+13) &&
            ! strcmp ((const char *)m->data, "Photoshop 3.0"))
            jpeg_decode_iptc ((unsigned char *)m->data);
        else if (m->marker == JPEG_COM) {
            if (! m_spec.find_attribute ("ImageDescription", TypeDesc::STRING))
                m_spec.attribute ("ImageDescription",
                                   std::string ((const char *)m->data));
        }
    }

    newspec = m_spec;
    return true;
}

```

```

bool
JpgInput::read_native_scanline (int y, int z, void *data)
{
    if (m_raw)
        return false;
    if (y < 0 || y >= (int)m_cinfo.output_height) // out of range scanline
        return false;
    if (m_next_scanline > y) {
        // User is trying to read an earlier scanline than the one we're
        // up to. Easy fix: close the file and re-open.
        ImageSpec dummyspec;
        int subimage = current_subimage();
        if (! close () ||
            ! open (m_filename, dummyspec) ||
            ! seek_subimage (subimage, dummyspec))
            return false; // Somehow, the re-open failed
        assert (m_next_scanline == 0 && current_subimage() == subimage);
    }
    while (m_next_scanline <= y) {
        // Keep reading until we're read the scanline we really need
        jpeg_read_scanlines (&m_cinfo, (JSAMPLE **)&data, 1); // read one scanline
        ++m_next_scanline;
    }
    return true;
}

```

```

}

bool
JpgInput::close ()
{
    if (m_fd != NULL) {
        // N.B. don't call finish_decompress if we never read anything
        if (m_next_scanline > 0) {
            // But if we've only read some scanlines, read the rest to avoid
            // errors
            std::vector<char> buf (spec().scanline_bytes());
            char *data = &buf[0];
            while (m_next_scanline < spec().height) {
                jpeg_read_scanlines (&m_cinfo, (JSAMPLE **)&data, 1);
                ++m_next_scanline;
            }
        }
        if (m_next_scanline > 0 || m_raw)
            jpeg_finish_decompress (&m_cinfo);
        jpeg_destroy_decompress (&m_cinfo);
        fclose (m_fd);
        m_fd = NULL;
    }
    init (); // Reset to initial state
    return true;
}

void
JpgInput::jpeg_decode_iptc (const unsigned char *buf)
{
    // APP13 blob doesn't have to be IPTC info. Look for the IPTC marker,
    // which is the string "Photoshop 3.0" followed by a null character.
    if (strcmp ((const char *)buf, "Photoshop 3.0"))
        return;
    buf += strlen("Photoshop 3.0") + 1;

    // Next are the 4 bytes "8BIM"
    if (strncmp ((const char *)buf, "8BIM", 4))
        return;
    buf += 4;

    // Next two bytes are the segment type, in big endian.
    // We expect 1028 to indicate IPTC data block.
    if (((buf[0] << 8) + buf[1]) != 1028)
        return;
    buf += 2;
}

```

```

// Next are 4 bytes of 0 padding, just skip it.
buf += 4;

// Next is 2 byte (big endian) giving the size of the segment
int segmentsize = (buf[0] << 8) + buf[1];
buf += 2;

OpenImageIO::decode_iptc_iim (buf, segmentsize, m_spec);
}

OIIO_PLUGIN_NAMESPACE_END

```

5.3 Image Writers

A plugin that writes a particular image file format must implement a *subclass* of `ImageOutput` (described in Chapter 3). This is actually very straightforward and consists of the following steps, which we will illustrate with a real-world example of writing a JPEG/JFIF plug-in.

1. Read the base class definition from `imageio.h`, just as with an image reader (see Section 5.2).
2. Declare three public items:
 - (a) An integer called `name_imageio_version` that identifies the version of the ImageIO protocol implemented by the plugin, defined in `imageio.h` as the constant `OPENIMAGEIO_PLUGIN_VERSION`. This allows the library to be sure it is not loading a plugin that was compiled against an incompatible version of `OpenImageIO`. Note that if your plugin has both a reader and writer and they are compiled as separate modules (C++ source files), you don't want to declare this in *both* modules; either one is fine.
 - (b) A function named `name_output_imageio_create` that takes no arguments and returns a new instance of your `ImageOutput` subclass. (Note that `name` is the name of your format, and must match the name of the plugin itself.)
 - (c) An array of `char *` called `name_output_extensions` that contains the list of file extensions that are likely to indicate a file of the right format. The list is terminated by a `NULL` pointer.

All of these items must be inside an `'extern "C"'` block in order to avoid name mangling by the C++ compiler. Depending on your compiler, you may need to use special commands to dictate that the symbols will be exported in the DSO; we provide a special `DLLEXPORT` macro for this purpose, defined in `export.h`.

Putting this all together, we get the following for our JPEG example:

```

extern "C" {
    DLLEXPORT int jpeg_imageio_version = OPENIMAGEIO_PLUGIN_VERSION;

```

```

    DLLEXPORT JpgOutput *jpeg_output_imageio_create () {
        return new JpgOutput;
    }
    DLLEXPORT const char *jpeg_input_extensions[] = {
        "jpg", "jpe", "jpeg", NULL
    };
};

```

3. The definition and implementation of an ImageOutput subclass for this file format. It must publicly inherit ImageOutput, and must overload the following methods which are “pure virtual” in the ImageOutput base class:

- (a) `format_name()` should return the name of the format, which ought to match the name of the plugin and by convention is strictly lower-case and contains no whitespace.
- (b) `supports()` should return `true` if its argument names a feature supported by your format plugin, `false` if it names a feature not supported by your plugin. See Section 3.3 for the list of feature names.
- (c) `open()` should open the file and return `true`, or should return `false` if unable to do so (including if the file was found but turned out not to be in the format that your plugin is trying to implement).
- (d) `close()` should close the file, if open.
- (e) `write_scanline` should write a single scanline to the file, translating from internal to native data format and handling strides properly.
- (f) The virtual destructor, which should `close()` if the file is still open, addition to performing any other tear-down activities.

Additionally, your ImageOutput subclass may optionally choose to overload any of the following methods, which are defined in the ImageOutput base class and only need to be overloaded if the default behavior is not appropriate for your plugin:

- (g) `write_tile()`, only if your format supports writing tiled images.
- (h) `write_rectangle()`, only if your format supports writing arbitrary rectangles.
- (i) `write_image()`, only if you have a more clever method of doing so than the default implementation that calls `write_scanline()` or `write_tile()` repeatedly.

Here is how the class definition looks for our JPEG example. Note that the JPEG/JFIF file format does not support multiple subimages or tiled images.

```

class JpgOutput : public ImageOutput {
public:
    JpgOutput () { init(); }
    virtual ~JpgOutput () { close(); }
    virtual const char * format_name (void) const { return "jpeg"; }
};

```

```

    virtual bool supports (const std::string &property) const { return false; }
    virtual bool open (const std::string &name, const ImageSpec &spec,
                      bool append=false);
    virtual bool write_scanline (int y, int z, TypeDesc format,
                                const void *data, stride_t xstride);

    bool close ();
private:
    FILE *m_fd;
    std::vector<unsigned char> m_scratch;
    struct jpeg_compress_struct m_cinfo;
    struct jpeg_error_mgr m_jerr;

    void init () { m_fd = NULL; }
};

```

Your subclass implementation of `open()`, `close()`, and `write_scanline()` are the heart of an `ImageOutput` implementation. (Also `write_tile()`, for those image formats that support tiled output.)

An `ImageOutput` implementation must properly handle all data formats and strides passed to `write_scanline()` or `write_tile()`, unlike an `ImageInput` implementation, which only needs to read scanlines or tiles in their native format and then have the super-class handle the translation. But don't worry, all the heavy lifting can be accomplished with the following helper functions provided as protected member functions of `ImageOutput`:

```

const void * to_native_scanline (TypeDesc format, const void *data,
                                stride_t xstride, std::vector<unsigned char> &scratch);

```

Convert a full scanline of pixels (pointed to by *data*) with the given *format* and strides into contiguous pixels in the native format (described by the `ImageSpec` returned by the `spec()` member function). The location of the newly converted data is returned, which may either be the original *data* itself if no data conversion was necessary and the requested layout was contiguous (thereby avoiding unnecessary memory copies), or may point into memory allocated within the *scratch* vector passed by the user. In either case, the caller doesn't need to worry about thread safety or freeing any allocated memory (other than eventually destroying the *scratch* vector).

```

const void * to_native_tile (TypeDesc format, const void *data,
                             stride_t xstride, stride_t ystride, stride_t zstride,
                             std::vector<unsigned char> &scratch);

```

Convert a full tile of pixels (pointed to by *data*) with the given *format* and strides into contiguous pixels in the native format (described by the `ImageSpec` returned by the `spec()` member function). The location of the newly converted data is returned, which may either be the original *data* itself if no data conversion was necessary and the requested layout was contiguous (thereby avoiding unnecessary memory copies), or may point into memory allocated within the *scratch* vector passed by the user. In either case, the caller

doesn't need to worry about thread safety or freeing any allocated memory (other than eventually destroying the scratch vector).

```
const void * to_native_rectangle (int xmin, int xmax, int ymin, int ymax,
                                int zmin, int zmax, TypeDesc format, const void *data,
                                stride_t xstride, stride_t ystride, stride_t zstride,
                                std::vector<unsigned char> &scratch);
```

Convert a rectangle of pixels (pointed to by *data*) with the given *format*, dimensions, and strides into contiguous pixels in the native format (described by the ImageSpec returned by the `spec()` member function). The location of the newly converted data is returned, which may either be the original *data* itself if no data conversion was necessary and the requested layout was contiguous (thereby avoiding unnecessary memory copies), or may point into memory allocated within the *scratch* vector passed by the user. In either case, the caller doesn't need to worry about thread safety or freeing any allocated memory (other than eventually destroying the scratch vector).

The remainder of this section simply lists the full implementation of our JPEG writer, which relies heavily on the open source jpeg-6b library to perform the actual JPEG encoding.

```
/*
Copyright 2008 Larry Gritz and the other authors and contributors.
All Rights Reserved.
Based on BSD-licensed software Copyright 2004 NVIDIA Corp.

Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions are
met:
* Redistributions of source code must retain the above copyright
  notice, this list of conditions and the following disclaimer.
* Redistributions in binary form must reproduce the above copyright
  notice, this list of conditions and the following disclaimer in the
  documentation and/or other materials provided with the distribution.
* Neither the name of the software's owners nor the names of its
  contributors may be used to endorse or promote products derived from
  this software without specific prior written permission.
THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS
"AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT
LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR
A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
OWNER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
```

```

    (This is the Modified BSD License)
    */

#include <cassert>
#include <cstdio>
#include <vector>

extern "C" {
#include "jpeglib.h"
}

#include "imageio.h"
#include "fmath.h"
#include "jpeg_pvt.h"

OIIO_PLUGIN_NAMESPACE_BEGIN

using namespace OpenImageIO;
using namespace Jpeg_imageio_pvt;

#define DBG if(0)

// See JPEG library documentation in /usr/share/doc/libjpeg-devel-6b

class JpgOutput : public ImageOutput {
public:
    JpgOutput () { init(); }
    virtual ~JpgOutput () { close(); }
    virtual const char * format_name (void) const { return "jpeg"; }
    virtual bool supports (const std::string &property) const { return false; }
    virtual bool open (const std::string &name, const ImageSpec &spec,
                      OpenMode mode=Create);
    virtual bool write_scanline (int y, int z, TypeDesc format,
                                const void *data, stride_t xstride);
    virtual bool close ();
    virtual bool copy_image (ImageInput *in);

private:
    FILE *m_fd;
    std::string m_filename;
    int m_next_scanline;           // Which scanline is the next to write?
    std::vector<unsigned char> m_scratch;
    struct jpeg_compress_struct m_cinfo;
    struct jpeg_error_mgr c_jerr;
    jvirt_barray_ptr *m_copy_coeffs;
    struct jpeg_decompress_struct *m_copy_decompressor;

```



```

    void init (void) {
        m_fd = NULL;
        m_copy_coeffs = NULL;
        m_copy_decompressor = NULL;
    }
};

OIIO_PLUGIN_EXPORTS_BEGIN

    DLLEXPORT ImageOutput *jpeg_output_imageio_create () {
        return new JpgOutput;
    }
    DLLEXPORT const char *jpeg_output_extensions[] = {
        "jpg", "jpe", "jpeg", NULL
    };

OIIO_PLUGIN_EXPORTS_END

bool
JpgOutput::open (const std::string &name, const ImageSpec &newspec,
                 OpenMode mode)
{
    if (mode != Create) {
        error ("%s does not support subimages or MIP levels", format_name());
        return false;
    }

    // Save name and spec for later use
    m_filename = name;
    m_spec = newspec;

    // Check for things this format doesn't support
    if (m_spec.width < 1 || m_spec.height < 1) {
        error ("Image resolution must be at least 1x1, you asked for %d x %d",
              m_spec.width, m_spec.height);
        return false;
    }
    if (m_spec.depth < 1)
        m_spec.depth = 1;
    if (m_spec.depth > 1) {
        error ("%s does not support volume images (depth > 1)", format_name());
        return false;
    }

    m_fd = fopen (name.c_str(), "wb");
    if (m_fd == NULL) {
        error ("Unable to open file \"%s\"", name.c_str());
    }
}

```

```

        return false;
    }

    int quality = 98;
    const ImageIOParameter *qual = newspec.find_attribute ("CompressionQuality",
                                                         TypeDesc::INT);

    if (qual)
        quality = * (const int *)qual->data();

    m_cinfo.err = jpeg_std_error (&c_jerr);           // set error handler
    jpeg_create_compress (&m_cinfo);                 // create compressor
    jpeg_stdio_dest (&m_cinfo, m_fd);                 // set output stream

    // Set image and compression parameters
    m_cinfo.image_width = m_spec.width;
    m_cinfo.image_height = m_spec.height;

    if (m_spec.nchannels == 3 || m_spec.nchannels == 4) {
        m_cinfo.input_components = 3;
        m_cinfo.in_color_space = JCS_RGB;
        m_spec.nchannels = 3; // Force RGBA -> RGB
        m_spec.alpha_channel = -1; // No alpha channel
    } else if (m_spec.nchannels == 1) {
        m_cinfo.input_components = 1;
        m_cinfo.in_color_space = JCS_GRAYSCALE;
    }

    m_cinfo.density_unit = 2; // RESUNIT_INCH;
    m_cinfo.X_density = 72;
    m_cinfo.Y_density = 72;
    m_cinfo.write_JFIF_header = true;

    if (m_copy_coeffs) {
        // Back door for copy()
        jpeg_copy_critical_parameters (m_copy_decompressor, &m_cinfo);
        DBG std::cout << "out open: copy_critical_parameters\n";
        jpeg_write_coefficients (&m_cinfo, m_copy_coeffs);
        DBG std::cout << "out open: write_coefficients\n";
    } else {
        // normal write of scanlines
        jpeg_set_defaults (&m_cinfo);                // default compression
        DBG std::cout << "out open: set_defaults\n";
        jpeg_set_quality (&m_cinfo, quality, TRUE);   // baseline values
        DBG std::cout << "out open: set_quality\n";
        jpeg_start_compress (&m_cinfo, TRUE);         // start working
        DBG std::cout << "out open: start_compress\n";
    }

    m_next_scanline = 0; // next scanline we'll write

    // Write JPEG comment, if sent an 'ImageDescription'
    ImageIOParameter *comment = m_spec.find_attribute ("ImageDescription",
                                                         TypeDesc::STRING);

```

```

    if (comment && comment->data()) {
        const char **c = (const char **) comment->data();
        jpeg_write_marker (&m_cinfo, JPEG_COM, (JOCTET*)*c, strlen(*c) + 1);
    }

    if (m_spec.linearity == ImageSpec::sRGB)
        m_spec.attribute ("Exif:ColorSpace", 1);

    // Write EXIF info, if we have anything
    std::vector<char> exif;
    encode_exif (m_spec, exif);
    if (exif.size())
        jpeg_write_marker (&m_cinfo, JPEG_APP0+1, (JOCTET*)&exif[0], exif.size());

    // Write IPTC IIM metadata tags, if we have anything
    std::vector<char> iptc;
    encode_iptc_iim (m_spec, iptc);
    if (iptc.size()) {
        static char photoshop[] = "Photoshop 3.0";
        std::vector<char> head (photoshop, photoshop+strlen(photoshop)+1);
        static char _8BIM[] = "8BIM";
        head.insert (head.end(), _8BIM, _8BIM+4);
        head.push_back (4);    // 0x0404
        head.push_back (4);
        head.push_back (0);    // four bytes of zeroes
        head.push_back (0);
        head.push_back (0);
        head.push_back (0);
        head.push_back ((char)(iptc.size() >> 8)); // size of block
        head.push_back ((char)(iptc.size() & 0xff));
        iptc.insert (iptc.begin(), head.begin(), head.end());
        jpeg_write_marker (&m_cinfo, JPEG_APP0+13, (JOCTET*)&iptc[0], iptc.size());
    }

    // Write XMP packet, if we have anything
    std::string xmp = OpenImageIO::encode_xmp (m_spec, true);
    if (! xmp.empty()) {
        static char prefix[] = "http://ns.adobe.com/xap/1.0/";
        std::vector<char> block (prefix, prefix+strlen(prefix)+1);
        block.insert (block.end(), xmp.c_str(), xmp.c_str()+xmp.length()+1);
        jpeg_write_marker (&m_cinfo, JPEG_APP0+1, (JOCTET*)&block[0], block.size());
    }

    m_spec.set_format (TypeDesc::UINT8); // JPG is only 8 bit

    return true;
}

```

bool

```

JpgOutput::write_scanline (int y, int z, TypeDesc format,
                          const void *data, stride_t xstride)
{
    y -= m_spec.y;
    if (y != m_next_scanline) {
        error ("Attempt to write scanlines out of order to %s",
              m_filename.c_str());
        return false;
    }
    if (y >= (int)m_cinfo.image_height) {
        error ("Attempt to write too many scanlines to %s", m_filename.c_str());
        return false;
    }
    assert (y == (int)m_cinfo.next_scanline);

    data = to_native_scanline (format, data, xstride, m_scratch);

    jpeg_write_scanlines (&m_cinfo, (JSAMPLE**)data, 1);
    ++m_next_scanline;

    return true;
}

bool
JpgOutput::close ()
{
    if (! m_fd)          // Already closed
        return true;

    if (m_next_scanline < spec().height && m_copy_coeffs == NULL) {
        // But if we've only written some scanlines, write the rest to avoid
        // errors
        std::vector<char> buf (spec().scanline_bytes(), 0);
        char *data = &buf[0];
        while (m_next_scanline < spec().height) {
            jpeg_write_scanlines (&m_cinfo, (JSAMPLE **)data, 1);
            // DBG std::cout << "out close: write_scanlines\n";
            ++m_next_scanline;
        }
    }

    if (m_next_scanline >= spec().height || m_copy_coeffs) {
        DBG std::cout << "out close: about to finish_compress\n";
        jpeg_finish_compress (&m_cinfo);
        DBG std::cout << "out close: finish_compress\n";
    } else {
        DBG std::cout << "out close: about to abort_compress\n";
        jpeg_abort_compress (&m_cinfo);
        DBG std::cout << "out close: abort_compress\n";
    }
}

```

```

    }
    DBG std::cout << "out close: about to destroy_compress\n";
    jpeg_destroy_compress (&m_cinfo);
    fclose (m_fd);
    m_fd = NULL;
    init();

    return true;
}

bool
JpgOutput::copy_image (ImageInput *in)
{
    if (in && !strcmp(in->format_name(), "jpeg")) {
        JpgInput *jpg_in = dynamic_cast<JpgInput *> (in);
        std::string in_name = jpg_in->filename ();
        DBG std::cout << "JPG copy_image from " << in_name << "\n";

        // Save the original input spec and close it
        ImageSpec orig_in_spec = in->spec();
        in->close ();
        DBG std::cout << "Closed old file\n";

        // Re-open the input spec, with special request that the JpgInput
        // will recognize as a request to merely open, but not start the
        // decompressor.
        ImageSpec in_spec;
        ImageSpec config_spec;
        config_spec.attribute ("_jpeg:raw", 1);
        in->open (in_name, in_spec, config_spec);

        // Re-open the output
        std::string out_name = m_filename;
        ImageSpec orig_out_spec = spec();
        close ();
        m_copy_coeffs = (jvirt_barray_ptr *)jpg_in->coeffs();
        m_copy_decompressor = &jpg_in->m_cinfo;
        open (out_name, orig_out_spec);

        // Strangeness -- the write_coefficients somehow sets things up
        // so that certain writes only happen in close(), which MUST
        // happen while the input file is still open. So we go ahead
        // and close() now, so that the caller of copy_image() doesn't
        // close the input file first and then wonder why they crashed.
        close ();

        return true;
    }
}

```

```
    return ImageOutput::copy_image (in);  
}  
  
OIIO_PLUGIN_NAMESPACE_END
```

5.4 Building ImageIO Plugins

FIXME – spell out how to compile and link plugins on each of the major platforms.

6 Bundled ImageIO Plugins

More information will be added soon.

- 6.1 BMP**
- 6.2 Cineon**
- 6.3 DDS**
- 6.4 DPX**
- 6.5 Field3D**
- 6.6 FITS**
- 6.7 HDR/RGBE**
- 6.8 ICO**
- 6.9 JPEG**
- 6.10 JPEG-2000**
- 6.11 OpenEXR**
- 6.12 PNG**
- 6.13 PNM**
- 6.14 Ptex**
- 6.15 SGI**
- 6.16 Softimage PIC**
- 6.17 Targa**
- 6.18 TIFF**
- 6.19 Zfile**

7 Cached Images

7.1 Image Cache Introduction and Theory of Operation

`ImageCache` is a utility class that allows an application to read pixels from a large number of image files while using a remarkably small amount of memory and other resources. Of course it is possible for an application to do this directly using `ImageInput` objects. But `ImageCache` offers the following advantages:

- `ImageCache` presents an even simpler user interface than `ImageInput`— the only supported operations are asking for an `ImageSpec` describing a subimage in the file, retrieving for a block of pixels, and locking/reading/releasing individual tiles. You refer to images by filename only; you don't need to keep track of individual file handles or `ImageInput` objects. You don't need to explicitly open or close files.
- The `ImageCache` is completely thread-safe; if multiple threads are accessing the same file, the `ImageCache` internals will handle all the locking and resource sharing.
- No matter how many image files you are accessing, the `ImageCache` will maintain a reasonable number of simultaneously-open files, automatically closing files that have not been needed recently.
- No matter how large the total pixels in all the image files you are dealing with are, the `ImageCache` will use only a small amount of memory. It does this by loading only the individual tiles requested, and as memory allotments are approached, automatically releasing the memory from tiles that have not been used recently.

In short, if you have an application that will need to read pixels from many large image files, you can rely on `ImageCache` to manage all the resources for you. It is reasonable to access thousands of image files totalling hundreds of GB of pixels, efficiently and using a memory footprint on the order of 50 MB.

Below are some simple code fragments that shows ImageCache in action:

```
#include "OpenImageIO/imagecache.h"

// Create an image cache and set some options
ImageCache *cache = ImageCache::create ();
cache->attribute ("max_memory_MB", 500.0);
cache->attribute ("autotile", 64);

// Get a block of pixels from a file.
// (for brevity of this example, let's assume that 'size' is the
// number of channels times the number of pixels in the requested region)
float pixels[size];
cache->get_pixels ("file1.jpg", 0, 0, xbegin, xend, ybegin, yend,
                 zbegin, zend, TypeDesc::FLOAT, pixels);

// Get information about a file
ImageSpec spec;
bool ok = cache->get_imagespec ("file2.exr", spec);
if (ok)
    std::cout << "resolution is " << spec.width << "x"
               << "spec.height << "\n";

// Request and hold a tile, do some work with its pixels, then release
ImageCache::Tile *tile;
tile = cache->get_tile ("file2.exr", 0, 0, x, y, z);
// The tile won't be freed until we release it, so this is safe:
TypeDesc format;
void *p = cache->tile_pixels (tile, format);
// Now p points to the raw pixels of the tile, whose data format
// is given by 'format'.
cache->release_tile (tile);
// Now cache is permitted to free the tile when needed

// Note that all files were referenced by name, we never had to open
// or close any files, and all the resource and memory management
// was automatic.

ImageCache::destroy (cache);
```

7.2 ImageCache API

7.2.1 Creating and destroying an image cache

ImageCache is an abstract API described as a pure virtual class. The actual internal implementation is not exposed through the external API of OpenImageIO. Because of this, you cannot construct or destroy the concrete implementation, so two static methods of ImageCache are provided:

```
static ImageCache *ImageCache::create (bool shared=true)
```

Creates a new ImageCache and returns a pointer to it. If `shared` is `true`, `create()` will return a pointer to a shared ImageCache (so that multiple parts of an application that request an ImageCache will all end up with the same one). If `shared` is `false`, a completely unique ImageCache will be created and returned.

```
static void ImageCache::destroy (ImageCache *x)
```

Destroys an allocated ImageCache, including freeing all system resources that it holds.

This is necessary to ensure that the memory is freed in a way that matches the way it was allocated within the library. Note that simply using `delete` on the pointer will not always work (at least, not on some platforms in which a DSO/DLL can end up using a different allocator than the main program).

It is safe to destroy even a shared ImageCache, as the implementation of `destroy()` will recognize a shared one and only truly release its resources if it has been requested to be destroyed as many times as shared ImageCache's were created.

7.2.2 Setting options and limits for the image cache

The following member functions of ImageCache allow you to set (and in some cases retrieve) options that control the overall behavior of the image cache:

```
bool attribute (const std::string &name, TypeDesc type, const void *val)
```

Sets an attribute (i.e., a property or option) of the ImageCache. The `name` designates the name of the attribute, `type` describes the type of data, and `val` is a pointer to memory containing the new value for the attribute.

If the ImageCache recognizes a valid attribute name that matches the type specified, the attribute will be set to the new value and `attribute()` will return `true`. If `name` is not recognized as a valid attribute name, or if the types do not match (e.g., `type` is `TypeDesc::FLOAT` but the named attribute is a string), the attribute will not be modified, and `attribute()` will return `false`.

Here are examples:

```
ImageCache *ts;  
...  
int maxfiles = 50;
```

```
ts->attribute ("max_open_files", TypeDesc::INT, &maxfiles);

const char *path = "/my/path";
ts->attribute ("searchpath", TypeDesc::STRING, &path);
```

Note that when passing a string, you need to pass a pointer to the `char*`, not a pointer to the first character. (Rationale: for an `int` attribute, you pass the address of the `int`. So for a string, which is a `char*`, you need to pass the address of the string, i.e., a `char**`).

The complete list of attributes can be found at the end of this section.

```
bool attribute (const std::string &name, int val)
bool attribute (const std::string &name, float val)
bool attribute (const std::string &name, double val)
bool attribute (const std::string &name, const char *val)
bool attribute (const std::string &name, const std::string & val)
```

Specialized versions of `attribute()` in which the data type is implied by the type of the argument.

For example, the following are equivalent to the example above for the general (pointer) form of `attribute()`:

```
ts->attribute ("max_open_files", 50);
ts->attribute ("searchpath", "/my/path");
```

```
bool getattribute (const std::string &name, TypeDesc type, void *val)
```

Gets the current value of an attribute of the `ImageCache`. The `name` designates the name of the attribute, `type` describes the type of data, and `val` is a pointer to memory where the user would like the value placed.

If the `ImageCache` recognizes a valid attribute name that matches the type specified, the attribute value will be stored at address `val` and `attribute()` will return `true`. If `name` is not recognized as a valid attribute name, or if the types do not match (e.g., `type` is `TypeDesc::FLOAT` but the named attribute is a string), no data will be written to `val`, and `attribute()` will return `false`.

Here are examples:

```
ImageCache *ts;
...
int maxfiles;
ts->getattribute ("max_open_files", TypeDesc::INT, &maxfiles);

const char *path;
ts->getattribute ("searchpath", TypeDesc::STRING, &path);
```

Note that when passing a string, you need to pass a pointer to the `char*`, not a pointer to the first character. Also, the `char*` will end up pointing to characters owned by the

ImageCache; the caller does not need to ever free the memory that contains the characters.

The complete list of attributes can be found at the end of this section.

```
bool getattribute (const std::string &name, int &val)
bool getattribute (const std::string &name, float &val)
bool getattribute (const std::string &name, double &val)
bool getattribute (const std::string &name, char **val)
bool getattribute (const std::string &name, std::string & val)
```

Specialized versions of `getattribute()` in which the data type is implied by the type of the argument.

For example, the following are equivalent to the example above for the general (pointer) form of `getattribute()`:

```
int maxfiles;
ts->getattribute ("max_open_files", &maxfiles);
const char *path;
ts->getattribute ("searchpath", &path);
```

Image cache attributes

Recognized attributes include the following:

`int max_open_files`

The maximum number of file handles that the image cache will hold open simultaneously. (Default = 100)

`float max_memory_MB`

The maximum amount of memory (measured in MB) that the image cache will use for its “tile cache.” (Default: 256.0 MB)

`string searchpath`

The search path for images: a colon-separated list of directories that will be searched in order for any image name that is not specified as an absolute path. (Default: no search path.)

`int autotile`

This attributes controls how the image cache deals with images that are not “tiled” (i.e., are stored as scanlines).

If `autotile` is set to 0 (the default), an untiled image will be treated as if it were a single tile of the resolution of the whole image. This is simple and fast, but can lead to poor cache behavior if you are simultaneously accessing many large untiled images.

If `autotile` is nonzero (e.g., 64 is a good recommended value), any untiled images will be read and cached as if they were constructed in tiles of size `autotile × autotile`. This leads to slightly more expensive disk access if you are using only a few images, but if you are using many untiled images, the caching will be much more efficient.

`int automip`

If `automip` is set to 0 (the default), an untiled single-subimage file will only be able to utilize that single subimage.

If `automip` is nonzero, any untiled, single-subimage (un-MIP-mapped) images will have lower-resolution MIP-map levels generated on-demand if pixels are requested from the lower-res subimages (that don't really exist). Essentially this makes the `ImageCache` pretend that the file is MIP-mapped even if it isn't.

`int forcefloat`

If set to nonzero, all image tiles will be converted to `float` type when stored in the image cache. This can be helpful especially for users of `ImageBuf` who want to simplify their image manipulations to only need to consider `float` data.

The default is zero, meaning that image pixels are not forced to be `float` when in cache.

`int accept_untiled`

When nonzero (the default), `ImageCache` accepts untiled images as usual. When set to zero, `ImageCache` will reject untiled images with an error condition, as if the file could not be properly read. This is sometimes helpful for applications that want to enforce use of tiled images only.

`int failure_retries`

When an `open()` or `read_tile()` calls fails, pause and try again, up to `failure_retries` times before truly returning a failure. This is meant to address spooky disk or network failures. The default is zero, meaning that failures of open or tile reading will immediately return as a failure.

7.2.3 Getting information about images

```
bool get_image_info (ustring filename, int subimage, int miplevel,
                    ustring dataname, TypeDesc datatype, void *data)
```

Retrieves information about the image named by `filename`. The `dataname` is a keyword indicating what information should be retrieved, `datatype` is the type of data expected, and `data` points to caller-owned memory where the results should be placed. It is up

to the caller to ensure that data contains enough space to hold an item of the requested datatype.

The return value is true if `get_image_info()` is able to find the requested dataname and it matched the requested datatype. If the requested data was not found, or was not of the right data type, `get_image_info()` will return false.

Supported dataname values include:

exists Return 1 if the file exists and is an image format that OpenImageIO knows how to read, otherwise return 0. The data pointer is not used.

subimages The number of subimages in the file, as an integer.

resolution The resolution of the image file, which is an array of 2 integers (described as `TypeDesc(INT,2)`).

miplevels The number of MIPmap levels for the specified subimage (an integer).

texturetype A string describing the type of texture of the given file, which describes how the texture may be used (also which texture API call is probably the right one for it). This currently may return one of: "unknown", "Plain Texture", "Volume Texture", "Shadow", or "Environment".

textureformat A string describing the format of the given file, which describes the kind of texture stored in the file. This currently may return one of: "unknown", "Plain Texture", "Volume Texture", "Shadow", "CubeFace Shadow", "Volume Shadow", "LatLong Environment", or "CubeFace Environment". Note that there are several kinds of shadows and environment maps, all accessible through the same API calls.

channels The number of color channels in the file (an integer).

format The native data format of the pixels in the file (an integer, giving the `TypeDesc::BASETYPE` of the data). Note that this is not necessarily the same as the data format stored in the image cache.

cachedformat The native data format of the pixels as stored in the image cache (an integer, giving the `TypeDesc::BASETYPE` of the data). Note that this is not necessarily the same as the native data format of the file.

viewingmatrix The viewing matrix, which is a 4×4 matrix (an `Imath::M44f`, described as `TypeDesc(FLOAT,MATRIX)`).

projectionmatrix The projection matrix, which is a 4×4 matrix (an `Imath::M44f`, described as `TypeDesc(FLOAT,MATRIX)`).

Anything else – For all other data names, the the metadata of the image file will be searched for an item that matches both the name and data type.

```
bool get_imagespec (ustring filename, ImageSpec &spec,
                  int subimage=0, int miplevel=0)
```

If the named image (and the specific subimage and MIP level) is found and able to be opened by an available image format plugin, and the designated subimage exists, this function copies its image specification for that subimage into `spec` and returns `true`. Otherwise, if the file is not found, could not be opened, is not of a format readable by any plugin that could be found, or the designated subimage did not exist in the file, the return value is `false` and `spec` will not be modified.

```
const ImageSpec * imagespec (ustring filename, int subimage=0, int miplevel=0)
```

If the named image is found and able to be opened by an available image format plugin, and the designated subimage exists, this function returns a pointer to an `ImageSpec` that describes it. Otherwise, if the file is not found, could not be opened, is not of a format readable by any plugin that could be found, or the designated subimage did not exist in the file, the return value is `NULL`.

This method is much more efficient than `get_imagespec()`, since it just returns a pointer to the spec held internally by the `ImageCache` (rather than copying the spec to the user's memory). However, the caller must beware that the pointer is only valid as long as nobody (even other threads) calls `invalidate()` on the file, or `invalidate_all()`, or destroys the `ImageCache`.

```
std::string resolve_filename (const std::string &filename)
```

Returns the true path to the given file name, with searchpath logic applied.

7.2.4 Getting pixels

```
bool get_pixels (ustring filename, int subimage, int miplevel,
                int xbegin, int xend, int ybegin, int yend, int zbegin, int zend,
                TypeDesc format, void *result)
```

Retrieve the rectangle of raw pixels of the designated subimage, storing the pixel values beginning at the address specified by `result`. The pixel values will be converted to the type specified by `format`. It is up to the caller to ensure that `result` points to an area of memory big enough to accommodate the requested rectangle (taking into consideration its dimensions, number of channels, and data format). The rectangular region to be retrieved includes `begin` but does not include `end` (much like STL `begin/end` usage). Requested pixels that are not part of the valid pixel data region of the image file will be filled with zero values.

7.2.5 Dealing with tiles

```
ImageCache::Tile get_tile (ustring filename, int subimage, int miplevel,
                          int x, int y, int z)
```


Find a tile given by an image filename, subimage, and pixel coordinates. An opaque pointer to the tile will be returned, or NULL if no such file (or tile within the file) exists or can be read. The tile will not be purged from the cache until after `release_tile()` is called on the tile pointer. This is thread-safe.

```
void release_tile (ImageCache::Tile *tile)
```

After finishing with a tile, `release_tile()` will allow it to once again be purged from the tile cache if required.

```
const void * tile_pixels (ImageCache::Tile *tile, TypeDesc &format)
```

For a tile retrieved by `get_tile()`, return a pointer to the pixel data itself, and also store in `format` the data type that the pixels are internally stored in (which may be different than the data type of the pixels in the disk file). This method should only be called on a tile that has been requested by `get_tile()` but has not yet been released with `release_tile()`.

```
void invalidate (ustring filename)
```

Invalidate any loaded tiles or open file handles associated with the filename, so that any subsequent queries will be forced to re-open the file or re-load any tiles (even those that were previously loaded and would ordinarily be reused). A client might do this if, for example, they are aware that an image being held in the cache has been updated on disk. This is safe to do even if other procedures are currently holding reference-counted tile pointers from the named image, but those procedures will not get updated pixels until they release the tiles they are holding.

```
void invalidate_all (bool force=false)
```

Invalidate all loaded tiles and open file handles, so that any subsequent queries will be forced to re-open the file or re-load any tiles (even those that were previously loaded and would ordinarily be reused). A client might do this if, for example, they are aware that an image being held in the cache has been updated on disk. This is safe to do even if other procedures are currently holding reference-counted tile pointers from the named image, but those procedures will not get updated pixels until they release the tiles they are holding. If `force` is true, everything will be invalidated, no matter how wasteful it is, but if `force` is false, in actuality files will only be invalidated if their modification times have been changed since they were first opened.

7.2.6 Errors and statistics

```
std::string geterror ()
```

If any other API routines return false, indicating that an error has occurred, this routine will retrieve the error and clear the error status. If no error has occurred since the last time `geterror()` was called, it will return an empty string.

```
std::string getstats (int level=1)
```

Returns a big string containing useful statistics about the ImageCache operations, suitable for saving to a file or outputting to the terminal. The `level` indicates the amount of detail in the statistics, with higher numbers (up to a maximum of 5) yielding more and more esoteric information.

8 Texture Access: TextureSystem

8.1 Texture System Introduction and Theory of Operation

Coming soon. FIXME

8.2 Helper Classes

8.2.1 Imath

The texture functionality of OpenImageIO uses the excellent open source Ilmbase package's Imath types when it requires 3D vectors and transformation matrixes. Specifically, we use `Imath::V3f` for 3D positions and directions, and `Imath::M44f` for 4×4 transformation matrices. To use these yourself, we recommend that you:

```
#include <OpenEXR/ImathVec.h>
#include <OpenEXR/ImathMatrix.h>
```

Please refer to the Ilmbase and OpenEXR documentation and header files for more complete information about use of these types in your own application. However, note that you are not strictly required to use these classes in your application — `Imath::V3f` has a memory layout identical to `float[3]` and `Imath::M44f` has a memory layout identical to `float[16]`, so as long as your own internal vectors and matrices have the same memory layout, it's ok to just cast pointers to them when passing as arguments to `TextureSystem` methods.

8.2.2 TextureOpt

`TextureOpt` is a structure that holds many options controlling single-point texture lookups. Because each texture lookup API call takes a reference to a `TextureOpt`, the call signatures remain uncluttered rather than having an ever-growing list of parameters, most of which will never vary from their defaults. Here is a brief description of the data members of a `TextureOpt` structure:

```
int nchannels
int firstchannel
```

The number of color channels to look up from the texture — for example, 1 (single channel), or 3 (for an RGB triple) — and the number of channels to look up. The defaults are `firstchannel = 0`, `nchannels = 1`.

Examples: To retrieve the first three channels (typically RGB), you should have `nchannels = 3`, `firstchannel = 0`. To retrieve just the blue channel, you should have `nchannels = 1`, `firstchannel = 2`.

`int subimage`

The subimage or face within the file to use for the texture lookup. The default is 0, and larger values only make sense for a texture file that supports subimages or separate images per face (such as Ptex). This will be ignored if the file does not have multiple subimages or separate per-face textures.

`Wrap swrap, twrap`

Specify the *wrap mode* for 2D texture lookups (and 3D volume texture lookups, using the additional `rwrap` field). These fields are ignored for shadow and environment lookups.

These specify what happens when texture coordinates are found to be outside the usual $[0, 1]$ range over which the texture is defined. Wrap is an enumerated type that may take on any of the following values:

`WrapBlack` The texture is black outside the $[0,1]$ range.

`WrapClamp` The texture coordinates will be clamped to $[0,1]$, i.e., the value outside $[0,1]$ will be the same as the color at the nearest point on the border.

`WrapPeriodic` The texture is periodic, i.e., wraps back to 0 after going past 1.

`WrapMirror` The texture presents a mirror image at the edges, i.e., the coordinates go from 0 to 1, then back down to 0, then back up to 1, etc.

`WrapDefault` Use whatever wrap might be specified in the texture file itself, or some other suitable default (caveat emptor).

The wrap mode does not need to be identical in the *s* and *t* directions.

`float swidth, twidth`

For each direction, gives a multiplier for the derivatives. Note that a width of 0 indicates a point sampled lookup (assuming that blur is also zero). The default width is 1, indicating that the derivatives should guide the amount of blur applied to the texture filtering (not counting any additional *blur* specified).

`float sblur, tblur`

For each direction, specifies an additional amount of pre-blur to apply to the texture (*after* derivatives are taken into account), expressed as a portion of the width of the texture. In other words, `blur = 0.1` means that the texture lookup should act as if the texture was pre-blurred with a filter kernel with a width 1/10 the size of the full image. The default blur amount is 0, indicating a sharp texture lookup.

float fill

Specifies the value that will be used for any color channels that are requested but not found in the file. For example, if you perform a 3-channel lookup on a 1-channel texture, the second two channels will get the fill value.

const float* missingcolor

If not NULL, indicates that a missing or broken texture should *not* be treated as an error, but rather will simply return the supplied color as the texture lookup color and `texture()` will return `true`. If the `missingcolor` field is left at its default (a NULL pointer), a missing or broken texture will be treated as an error and `texture()` will return `false`. Note: When not NULL, the data must point to *nchannels* contiguous floats.

float *dresultds, *dresultdt

If not NULL (the default), these specify locations in which to store the *derivatives* of the texture lookup, i.e., the change of the filtered texture per unit of *s* and *t*, respectively. Each must point to *nchannels* contiguous floats. If either is NULL, the derivative computations will not be performed.

float bias

For shadow map lookups only, this gives the “shadow bias” amount.

int samples

For shadow map lookups only, the number of samples to use for the lookup.

Wrap rwrap

float rblur, rwidth

float *dresultdr

Specifies wrap, blur, width, and derivative results for the third component of 3D volume texture lookups. These are not used for 2D texture lookups.

8.2.3 TextureOptions

`TextureOptions` is a structure that holds many options controlling batched texture lookups. Because each texture lookup API call takes a reference to a `TextureOptions`, the call signatures remain uncluttered rather than having an ever-growing list of parameters, most of which will never vary from their defaults. Here is a brief description of the data members of a `TextureOptions` structure:

int nchannels

int firstchannel

The number of color channels to look up from the texture — for example, 1 (single channel), or 3 (for an RGB triple) — and the number of channels to look up. The defaults are `firstchannel = 0`, `nchannels = 1`.

Examples: To retrieve the first three channels (typically RGB), you should have `nchannels = 3`, `firstchannel = 0`. To retrieve just the blue channel, you should have `nchannels = 1`, `firstchannel = 2`.

`int subimage`

The subimage or face within the file to use for the texture lookup. The default is 0, and larger values only make sense for a texture file that supports subimages or separate images per face (such as Ptex). This will be ignored if the file does not have multiple subimages or separate per-face textures.

Wrap `swrap`, `twrap`

Specify the *wrap mode* for 2D texture lookups (and 3D volume texture lookups, using the additional `rwrap` field). These fields are ignored for shadow and environment lookups.

These specify what happens when texture coordinates are found to be outside the usual $[0, 1]$ range over which the texture is defined. Wrap is an enumerated type that may take on any of the following values:

`WrapBlack` The texture is black outside the $[0, 1]$ range.

`WrapClamp` The texture coordinates will be clamped to $[0, 1]$, i.e., the value outside $[0, 1]$ will be the same as the color at the nearest point on the border.

`WrapPeriodic` The texture is periodic, i.e., wraps back to 0 after going past 1.

`WrapMirror` The texture presents a mirror image at the edges, i.e., the coordinates go from 0 to 1, then back down to 0, then back up to 1, etc.

`WrapDefault` Use whatever wrap might be specified in the texture file itself, or some other suitable default (caveat emptor).

The wrap mode does not need to be identical in the *s* and *t* directions.

`VaryingRef<float> swidth, twidth`

For each direction, gives a multiplier for the derivatives. Note that a width of 0 indicates a point sampled lookup (assuming that blur is also zero). The default width is 1, indicating that the derivatives should guide the amount of blur applied to the texture filtering (not counting any additional *blur* specified).

`VaryingRef<float> sblur, tblur`

For each direction, specifies an additional amount of pre-blur to apply to the texture (*after* derivatives are taken into account), expressed as a portion of the width of the texture. In other words, `blur = 0.1` means that the texture lookup should act as if the texture was pre-blurred with a filter kernel with a width 1/10 the size of the full image. The default blur amount is 0, indicating a sharp texture lookup.

`VaryingRef<float> fill`

Specifies the value that will be used for any color channels that are requested but not found in the file. For example, if you perform a 3-channel lookup on a 1-channel texture, the second two channels will get the fill value.

`VaryingRef<float> missingcolor`

If supplied, indicates that a missing or broken texture should *not* be treated as an error, but rather will simply return the supplied color as the texture lookup color and `texture()` will return `true`. If the `missingcolor` field is left at its default (a NULL pointer), a missing or broken texture will be treated as an error and `texture()` will return `false`.

Although this is a `VaryingRef<float>`, the data must point to *nchannels* contiguous floats, and if “varying,” the step size must be set to `nchannels*sizeof(float)`, not `sizeof(float)`.

`float *dresultds, *dresultdt`

If not NULL (the default), these specify locations in which to store the *derivatives* of the texture lookup, i.e., the change of the filtered texture per unit of *s* and *t*, respectively. Each must point to *nchannels* contiguous floats. If either is NULL, the derivative computations will not be performed.

`VaryingRef<float> bias`

For shadow map lookups only, this gives the “shadow bias” amount.

`VaryingRef<int> samples`

For shadow map lookups only, the number of samples to use for each lookup.

`Wrap rwrap`

`VaryingRef<float> rblur, rwidth`

`float *dresultdr`

Specifies wrap, blur, width, and derivative results for the third component of 3D volume texture lookups. These are not used for 2D texture lookups.

8.2.4 VaryingRef: encapsulate uniform and varying

Many texture access API routines are designed to look up texture efficiently at many points at once. Therefore, many of the parameters to the API routines, and many of the fields in `TextureOptions` need to accommodate both uniform and varying values. *Uniform* means that a single value may be used for each of the many simultaneous texture lookups, whereas *varying* means that a different value is provided for each of the positions where you are sampling the texture.

Please read the comments in "varyingref.h" for the full gory details, but here's all you really need to know about it to use the texture functionality. Let's suppose that we have a routine whose prototype looks like this:

```
void API (int n, VaryingRef<float> x);
```

This means that parameter x may either be a single value for use at each of the n texture lookups, or it may have n different values of x .

If you want to pass a uniform value, you may do any of the following:

```
float x;    // just one value
API (n, x); // automatically knows what to do!
API (n, &x); // Also ok to pass the pointer to x
API (n, VaryingRef<float>(x)); // Wordy but correct
API (n, Uniform(x)); // Shorthand
```

If you want to pass a varying value, i.e., an array of values,

```
float x[n]; // One value for each of n points
API (n, VaryingRef<float>(x), sizeof(x)); // Wordy but correct
API (n, Varying(x)); // Shorthand if stride is sizeof(x)
```

You can also initialize a `VaryingRef` directly:

```
float x;    // just one value
float y[n]; // array of values
VaryingRef<float> r;
r.init (&x); // Initialize to uniform
r.init (&x, 0); // Initialize to uniform the wordy way
r.init (&y, sizeof(float)); // Initialize to varying
...
API (n, r);
```

8.2.5 SIMD Run Flags

Many of the texture lookup API routines are written to accommodate queries about many points at once. Furthermore, only a subset of points may need to compute. This is all expressed using three parameters: `Runflag *runflags`, `int beginactive`, `int endactive`. There are also `VaryingRef` parameters such as `s` and `t` that act as if they are arrays.

The `beginactive` and `endactive` indices are the first (inclusive) and last (exclusive) points that should be computed, and for each point `runflags[i]` is nonzero if the point should be computed. To illustrate, here is how a routine might be written that would copy values in `arg` to `result` using `runflags`:


```
void copy (Runflag *runflags, int beginactive, int endactive,
          VaryingRef<float> arg, VaryingRef <float> result)
{
    for (int i = beginactive; i < endactive; ++i)
        if (runflags[i])
            result[i] = arg[i];
}
```

8.3 TextureSystem API

8.3.1 Creating and destroying texture systems

TextureSystem is an abstract API described as a pure virtual class. The actual internal implementation is not exposed through the external API of OpenImageIO. Because of this, you cannot construct or destroy the concrete implementation, so two static methods of TextureSystem are provided:

```
static TextureSystem *TextureSystem::create (bool share=true)
```

Creates a new TextureSystem and returns a pointer to it. If shared is true, the TextureSystem created will share its underlying ImageCache with any other TextureSystem's or ImageCache's that requested shared caches. If shared is false, a completely unique ImageCache will be created that is private to this particular TextureSystem.

```
static void TextureSystem::destroy (TextureSystem *x)
```

Destroys an allocated TextureSystem, including freeing all system resources that it holds.

This is necessary to ensure that the memory is freed in a way that matches the way it was allocated within the library. Note that simply using delete on the pointer will not always work (at least, not on some platforms in which a DSO/DLL can end up using a different allocator than the main program).

8.3.2 Setting options and limits for the texture system

The following member functions of TextureSystem allow you to set (and in some cases retrieve) options that control the overall behavior of the texture system:

```
bool attribute (const std::string &name, TypeDesc type, const void *val)
```

Sets an attribute (i.e., a property or option) of the TextureSystem. The name designates the name of the attribute, type describes the type of data, and val is a pointer to memory containing the new value for the attribute.

If the TextureSystem recognizes a valid attribute name that matches the type specified, the attribute will be set to the new value and attribute() will return true. If name is not recognized as a valid attribute name, or if the types do not match (e.g., type is TypeDesc::FLOAT but the named attribute is a string), the attribute will not be modified, and attribute() will return false.

Here are examples:

```
TextureSystem *ts;
...
int maxfiles = 50;
ts->attribute ("max_open_files", TypeDesc::INT, &maxfiles);

const char *path = "/my/path";
ts->attribute ("searchpath", TypeDesc::STRING, &path);
```

Note that when passing a string, you need to pass a pointer to the `char*`, not a pointer to the first character. (Rationale: for an `int` attribute, you pass the address of the `int`. So for a string, which is a `char*`, you need to pass the address of the string, i.e., a `char**`).

The complete list of attributes can be found at the end of this section.

```
bool attribute (const std::string &name, int val)
bool attribute (const std::string &name, float val)
bool attribute (const std::string &name, double val)
bool attribute (const std::string &name, const char *val)
bool attribute (const std::string &name, const std::string & val)
```

Specialized versions of `attribute()` in which the data type is implied by the type of the argument.

For example, the following are equivalent to the example above for the general (pointer) form of `attribute()`:

```
ts->attribute ("max_open_files", 50);
ts->attribute ("searchpath", "/my/path");
```

```
bool getattribute (const std::string &name, TypeDesc type, void *val)
```

Gets the current value of an attribute of the `TextureSystem`. The name designates the name of the attribute, type describes the type of data, and `val` is a pointer to memory where the user would like the value placed.

If the `TextureSystem` recognizes a valid attribute name that matches the type specified, the attribute value will be stored at address `val` and `attribute()` will return `true`. If name is not recognized as a valid attribute name, or if the types do not match (e.g., type is `TypeDesc::FLOAT` but the named attribute is a string), no data will be written to `val`, and `attribute()` will return `false`.

Here are examples:

```
TextureSystem *ts;
...
int maxfiles;
ts->getattribute ("max_open_files", TypeDesc::INT, &maxfiles);

const char *path;
ts->getattribute ("searchpath", TypeDesc::STRING, &path);
```

Note that when passing a string, you need to pass a pointer to the `char*`, not a pointer to the first character. Also, the `char*` will end up pointing to characters owned by the `TextureSystem`; the caller does not need to ever free the memory that contains the characters.

The complete list of attributes can be found at the end of this section.

```

bool getattribute (const std::string &name, int &val)
bool getattribute (const std::string &name, float &val)
bool getattribute (const std::string &name, double &val)
bool getattribute (const std::string &name, char **val)
bool getattribute (const std::string &name, std::string & val)

```

Specialized versions of `getattribute()` in which the data type is implied by the type of the argument.

For example, the following are equivalent to the example above for the general (pointer) form of `getattribute()`:

```

int maxfiles;
ts->getattribute ("max_open_files", &maxfiles);
const char *path;
ts->getattribute ("searchpath", &path);

```

Texture system attributes

Recognized attributes include the following:

```

int max_open_files
float max_memory_MB
string searchpath
int autotile
int automip
int accept_untiled
int failure_retries

```

These attributes are all passed along to the underlying `ImageCache` that is used internally by the `TextureSystem`. Please consult the `ImageCache` attribute list in Section 7.2.2 for explanations of these attributes.

`matrix worldtocommon`

The 4×4 matrix that provides the spatial transformation from “world” to a “common” coordinate system. This is used for shadow map lookups, in which the shadow map itself encodes the world coordinate system, but positions passed to `shadow()` are expressed in “common” coordinates.

`matrix commontoworld`

The 4×4 matrix that is the inverse of `worldtocommon` — that is, it transforms points from “common” to “world” coordinates.

You do not need to set `commontoworld` and `worldtocommon` separately; just setting either one will implicitly set the other, since each is the inverse of the other.

8.3.3 Texture Lookups

```
bool texture (ustring filename, TextureOpt &options,  
             float s, float t, float dsdx, float dtdx,  
             float dsdy, float dtdy, float *result)
```

Perform a filtered 2D texture lookup on a position centered at 2D coordinates (s, t) from the texture identified by filename, and using relevant texture options. The filtered results will be stored in result[].

We assume that this lookup will be part of an image that has pixel coordinates x and y. By knowing how s and t change from pixel to pixel in the final image, we can properly *filter* or antialias the texture lookups. This information is given via derivatives dsdx and dtdx that define the change in s and t per unit of x, and dsdy and dtdy that define the change in s and t per unit of y. If it is impossible to know the derivatives, you may pass 0 for them, but in that case you will not receive an antialiased texture lookup.

Fields within options that are honored for 2D texture lookups include the following:

```
int nchannels
```

The number of color channels to look up from the texture.

```
int firstchannel
```

The index of the first channel to look up from the texture.

```
int subimage
```

The subimage or face within the file. This will be ignored if the file does not have multiple subimages or separate per-face textures.

```
Wrap swrap, twrap
```

Specify the *wrap mode* for each direction, one of: WrapBlack, WrapClamp, WrapPeriodic, WrapMirror, or WrapDefault.

```
float swidth, twidth
```

For each direction, gives a multiplier for the derivatives.

```
float sblur, tblur
```

For each direction, specifies an additional amount of pre-blur to apply to the texture (*after* derivatives are taken into account), expressed as a portion of the width of the texture.

```
float fill
```

Specifies the value that will be used for any color channels that are requested but not found in the file.

```
const float *missingcolor
```

If not NULL, specifies the color that will be returned for missing or broken textures (rather than being an error).

```
float *dresultds, *dresultdt
```

If not NULL, specifies locations in which to store the s and t derivatives of the filtered texture lookup.

This function returns true upon success, or false if the file was not found or could not be opened by any available ImageIO plugin.

```
bool texture (ustring filename, TextureOptions &options,  
             Runflag *runflags, int beginactive, int endactive,  
             VaryingRef<float> s, VaryingRef<float> t,  
             VaryingRef<float> dsdx, VaryingRef<float> dtdx,  
             VaryingRef<float> dsdy, VaryingRef<float> dtdy,  
             float *result)
```

Perform filtered 2D texture lookups on a collection of positions all at once, which may be much more efficient than repeatedly calling the single-point version of `texture()`. The parameters s , t , $dsdx$, $dtdx$, and $dsdy$, $dtdy$ are now `VaryingRef`'s that may refer to either a single or an array of values, as are many of the fields in the options.

Texture will be computed at indices `beginactive` through `endactive` (exclusive of the end), but only at indices where `runflags[i]` is nonzero. Results will be stored at corresponding positions of `result`, that is, `result[i*n ... (i+1)*n-1]` where n is the number of channels requested by `options.nchannels`.

This function returns true upon success, or false if the file was not found or could not be opened by any available ImageIO plugin.

8.3.4 Volume Texture Lookups

```
bool texture3d (ustring filename, TextureOpt &options,
               const Imath::V3f &P,
               const Imath::V3f &dPdx,
               const Imath::V3f &dPdy,
               const Imath::V3f &dPdz,
               float *result)
```

Perform a filtered 3D volumetric texture lookup on a position centered at 3D position *P* from the texture identified by *filename*, and using relevant texture options. The filtered results will be stored in *result* [].

We assume that this lookup will be part of an image that has pixel coordinates *x* and *y* and depth *z*. By knowing how *P* changes from pixel to pixel in the final image, and as we step in *z* depth, we can properly *filter* or antialias the texture lookups. This information is given via derivatives *dPdx*, *dPdy*, and *dPdz* that define the changes in *P* per unit of *x*, *y*, and *z*, respectively. If it is impossible to know the derivatives, you may pass 0 for them, but in that case you will not receive an antialiased texture lookup.

Fields within *options* that are honored for 3D texture lookups include the following:

`int nchannels`

The number of color channels to look up from the texture.

`int firstchannel`

The index of the first channel to look up from the texture.

`Wrap swrap, twrap, rwrap`

Specify the wrap modes for each direction, one of: `WrapBlack`, `WrapClamp`, `WrapPeriodic`, `WrapMirror`, or `WrapDefault`.

`float swidth, twidth, rwidth`

For each direction, gives a multiplier for the derivatives.

`float sblur, tblur, rblur`

For each direction, specifies an additional amount of pre-blur to apply to the texture (*after* derivatives are taken into account), expressed as a portion of the width of the texture.

`float fill`

Specifies the value that will be used for any color channels that are requested but not found in the file.

`const float *missingcolor`

If not `NULL`, specifies the color that will be returned for missing or broken textures (rather than being an error).

```
float *dresultds, *dresultdt, *dresultdr
```

If not NULL, specifies locations in which to store the s , t , and r derivatives of the filtered texture lookup.

This function returns true upon success, or false if the file was not found or could not be opened by any available ImageIO plugin.

```
bool texture3d (ustring filename, TextureOptions &options,
               Runflag *runflags, int beginactive, int endactive,
               VaryingRef<Imath::V3f> P,
               VaryingRef<Imath::V3f> dPdx,
               VaryingRef<Imath::V3f> dPdy,
               VaryingRef<Imath::V3f> dPdz,
               float *result)
```

Perform filtered 3D volumetric texture lookups on a collection of positions all at once, which may be much more efficient than repeatedly calling the single-point version of `texture()`. The parameters `P`, `dPdx`, `dPdy`, and `dPdz` are now `VaryingRef`'s that may refer to either a single or an array of values, as are all the fields in the options.

Texture will be computed at indices `beginactive` through `endactive` (exclusive of the end), but only at indices where `runflags[i]` is nonzero. Results will be stored at corresponding positions of `result`, that is, `result[i*n ... (i+1)*n-1]` where n is the number of channels requested by `options.nchannels`.

This function returns true upon success, or false if the file was not found or could not be opened by any available ImageIO plugin.

8.3.5 Shadow Lookups

```
bool shadow (ustring filename, TextureOptions &opt,
             const Imath::V3f &P,
             const Imath::V3f &dPdx,
             const Imath::V3f &dPdy,
             float *result)
```

Perform a shadow map lookup on a position centered at 3D coordinate P (in a designated “common” space) from the shadow map identified by `filename`, and using relevant texture options. The filtered results will be stored in `result[]`.

We assume that this lookup will be part of an image that has pixel coordinates `x` and `y`. By knowing how P changes from pixel to pixel in the final image, we can properly *filter* or antialias the texture lookups. This information is given via derivatives `dPdx` and `dPdy` that define the changes in P per unit of `x` and `y`, respectively. If it is impossible to know the derivatives, you may pass 0 for them, but in that case you will not receive an antialiased texture lookup.

Fields within options that are honored for 2D texture lookups include the following:

```
float swidth, twidth
```

For each direction, gives a multiplier for the derivatives.

```
float sblur, tblur
```

For each direction, specifies an additional amount of pre-blur to apply to the texture (*after* derivatives are taken into account), expressed as a portion of the width of the texture.

```
float bias
```

Specifies the amount of *shadow bias* to use — this effectively ignores shadow occlusion that is closer than the bias amount to the surface, helping to eliminate self-shadowing artifacts.

```
int samples
```

Specifies the number of samples to use when evaluating the shadow map. More samples will give a smoother, less noisy, appearance to the shadows, but may also take longer to compute.

This function returns `true` upon success, or `false` if the file was not found or could not be opened by any available ImageIO plugin.

```
bool shadow (ustring filename, TextureOptions &options,
             Runflag *runflags, int beginactive, int endactive,
             VaryingRef<Imath::V3f> P,
             VaryingRef<Imath::V3f> dPdx,
             VaryingRef<Imath::V3f> dPdy,
             float *result)
```

Perform filtered shadow map lookups on a collection of positions all at once, which may be much more efficient than repeatedly calling the single-point version of `shadow()`. The parameters `P`, `dPdx`, and `dPdy` are now `VaryingRef`'s that may refer to either a single or an array of values, as are many the fields in the `options`.

Shadow lookups will be computed at indices `beginactive` through `endactive` (exclusive of the end), but only at indices where `runflags[i]` is nonzero. Results will be stored at corresponding positions of `result`, that is, `result[i*n ... (i+1)*n-1]` where n is the number of channels requested by `options.nchannels`.

This function returns `true` upon success, or `false` if the file was not found or could not be opened by any available `ImageIO` plugin.

8.3.6 Environment Lookups

```
bool environment (usttring filename, TextureOpt &options,
                 const Imath::V3f &R,
                 const Imath::V3f &dRdx,
                 const Imath::V3f &dRdy,
                 float *result)
```

Perform a filtered directional environment map lookup in the direction of vector *R*, from the texture identified by *filename*, and using relevant texture options. The filtered results will be stored in *result* [].

We assume that this lookup will be part of an image that has pixel coordinates *x* and *y*. By knowing how *R* changes from pixel to pixel in the final image, we can properly *filter* or antialias the texture lookups. This information is given via derivatives *dRdx* and *dRdy* that define the changes in *R* per unit of *x* and *y*, respectively. If it is impossible to know the derivatives, you may pass 0 for them, but in that case you will not receive an antialiased texture lookup.

Fields within options that are honored for 3D texture lookups include the following:

`int nchannels`

The number of color channels to look up from the texture.

`int firstchannel`

The index of the first channel to look up from the texture.

`float swidth, twidth`

For each direction, gives a multiplier for the derivatives.

`float sblur, tblur`

For each direction, specifies an additional amount of pre-blur to apply to the texture (*after* derivatives are taken into account), expressed as a portion of the width of the texture.

`float fill`

Specifies the value that will be used for any color channels that are requested but not found in the file.

This function returns `true` upon success, or `false` if the file was not found or could not be opened by any available ImageIO plugin.

```
bool environment (usttring filename, TextureOptions &options,
                 Runflag *runflags, int beginactive, int endactive,
                 VaryingRef<Imath::V3f> R,
                 VaryingRef<Imath::V3f> dRdx,
                 VaryingRef<Imath::V3f> dRdy,
                 float *result)
```

Perform filtered directional environment map lookups on a collection of directions all at once, which may be much more efficient than repeatedly calling the single-point version of `environment()`. The parameters `R`, `dRdx`, and `dRdy` are now `VaryingRef`'s that may refer to either a single or an array of values, as are many the fields in the `options`.

Results will be computed at indices `beginactive` through `endactive` (exclusive of the end), but only at indices where `runflags[i]` is nonzero. Results will be stored at corresponding positions of `result`, that is, `result[i*n ... (i+1)*n-1]` where n is the number of channels requested by `options.nchannels`.

This function returns `true` upon success, or `false` if the file was not found or could not be opened by any available `ImageIO` plugin.

8.3.7 Texture Metadata and Raw Texels

```
bool get_texture_info (ustring filename, int subimage,
                      ustring dataname, TypeDesc datatype, void *data)
```

Retrieves information about the texture named by filename. The dataname is a keyword indicating what information should be retrieved, datatype is the type of data expected, and data points to caller-owned memory where the results should be placed. It is up to the caller to ensure that data contains enough space to hold an item of the requested datatype.

The return value is true if get_texture_info() is able to find the requested dataname and it matched the requested datatype. If the requested data was not found, or was not of the right data type, get_texture_info() will return false.

Supported dataname values include:

exists Return 1 if the file exists and is an image format that OpenImageIO knows how to read, otherwise return 0. The data pointer is not used.

subimages The number of subimages/faces in the file, as an integer.

resolution The resolution of the texture file, which is an array of 2 integers (described as TypeDesc(INT,2)).

resolution (int[3]) The 3D resolution of the texture file, which is an array of 3 integers (described as TypeDesc(INT,3)) The third value will be 1 unless it's a volumetric (3D) image.

miplevels The number of MIPmap levels for the specified subimage (an integer).

texturetype A string describing the type of texture of the given file, which describes how the texture may be used (also which texture API call is probably the right one for it). This currently may return one of: "unknown", "Plain Texture", "Volume Texture", "Shadow", or "Environment".

textureformat A string describing the format of the given file, which describes the kind of texture stored in the file. This currently may return one of: "unknown", "Plain Texture", "Volume Texture", "Shadow", "CubeFace Shadow", "Volume Shadow", "LatLong Environment", or "CubeFace Environment". Note that there are several kinds of shadows and environment maps, all accessible through the same API calls.

channels The number of color channels in the file (an integer).

viewingmatrix The viewing matrix, which is a 4×4 matrix (an Imath::M44f, described as TypeDesc(FLOAT,MATRIX)).

projectionmatrix The projection matrix, which is a 4×4 matrix (an Imath::M44f, described as TypeDesc(FLOAT,MATRIX)).

Anything else – For all other data names, the the metadata of the image file will be searched for an item that matches both the name and data type.

```
bool get_imagespec (ustring filename, int subimage, ImageSpec &spec)
```

If the named image is found and able to be opened by an available image format plugin, this function copies its image specification into `spec` and returns `true`. Otherwise, if the file is not found, could not be opened, or is not of a format readable by any plugin that could be found, the return value is `false`.

```
const ImageSpec * imagespec (ustring filename, int subimage)
```

If the named image is found and able to be opened by an available image format plugin, and the designated subimage exists, this function returns a pointer to an `ImageSpec` that describes it. Otherwise, if the file is not found, could not be opened, is not of a format readable by any plugin that could be find, or the designated subimage did not exist in the file, the return value is `NULL`.

This method is much more efficient than `get_imagespec()`, since it just returns a pointer to the spec held internally by the underlying `ImageCache` (rather than copying the spec to the user's memory). However, the caller must beware that the pointer is only valid as long as nobody (even other threads) calls `invalidate()` on the file, or `invalidate_all()`, or destroys the `TextureSystem`.

```
bool get_texels (ustring filename, TextureOptions &options, int level,
                int xbegin, int xend, int ybegin, int yend,
                int zbegin, int zend, TypeDesc format, void *result)
```

Retrieve a rectangle of raw unfiltered texels at the named MIP-map level, storing the texel values beginning at the address specified by `result`. Note that the face/subimage is communicated through `options.face`. The texel values will be converted to the type specified by `format`. It is up to the caller to ensure that `result` points to an area of memory big enough to accommodate the requested rectangle (taking into consideration its dimensions, number of channels, and data format). The rectangular region to be retrieved includes `begin` but does not include `end` (much like STL `begin/end` usage). Requested pixels that are not part of the valid pixel data region of the image file will be filled with zero values.

Fields within `options` that are honored for raw texel retrieval include the following:

```
int nchannels
```

The number of color channels to look up from the texture.

```
int firstchannel
```

The index of the first channel to look up from the texture.

```
VaryingRef<float> fill
```

Specifies the value that will be used for any color channels that are requested but not found in the file.

Return `true` if the file is found and could be opened by an available `ImageIO` plugin, otherwise return `false`.

```
std::string resolve_filename (const std::string &filename)
```

Returns the true path to the given file name, with searchpath logic applied.

8.3.8 Miscellaneous – Statistics, errors, flushing the cache

```
std::string geterror ()
```

If any other API routines return false, indicating that an error has occurred, this routine will retrieve the error and clear the error status. If no error has occurred since the last time `geterror()` was called, it will return an empty string.

```
std::string getstats (int level=1, bool icstats=true)
```

Returns a big string containing useful statistics about the ImageCache operations, suitable for saving to a file or outputting to the terminal. The `level` indicates the amount of detail in the statistics, with higher numbers (up to a maximum of 5) yielding more and more esoteric information. If `icstats` is true, the returned string will also contain all the statistics of the underlying ImageCache, but if false will only contain texture-specific statistics.

```
void invalidate (ustring filename)
```

Invalidate any loaded tiles or open file handles associated with the filename, so that any subsequent queries will be forced to re-open the file or re-load any tiles (even those that were previously loaded and would ordinarily be reused). A client might do this if, for example, they are aware that an image being held in the cache has been updated on disk. This is safe to do even if other procedures are currently holding reference-counted tile pointers from the named image, but those procedures will not get updated pixels until they release the tiles they are holding.

```
void invalidate_all (bool force=false)
```

Invalidate all loaded tiles and open file handles, so that any subsequent queries will be forced to re-open the file or re-load any tiles (even those that were previously loaded and would ordinarily be reused). A client might do this if, for example, they are aware that an image being held in the cache has been updated on disk. This is safe to do even if other procedures are currently holding reference-counted tile pointers from the named image, but those procedures will not get updated pixels until they release the tiles they are holding. If `force` is true, everything will be invalidated, no matter how wasteful it is, but if `force` is false, in actuality files will only be invalidated if their modification times have been changed since they were first opened.

9 Python Bindings

10 Image Buffer

Part II

Image Utilities

11 The `iv` Image Viewer

The `iv` program is a great interactive image viewer. Because `iv` is built on top on OpenImageIO, it can display images of any formats readable by ImageInput plugins on hand.

More documentation on this later.

12 Getting Image information

With `iinfo`

The `iinfo` program will print either basic information (name, resolution, format) or detailed information (including all metadata) found in images. Because `iinfo` is built on top on OpenImageIO, it will print information about images of any formats readable by ImageInput plugins on hand.

12.1 Using `iinfo`

The `iinfo` utility is invoked as follows:

```
iinfo [options] filename ...
```

Where *filename* (and any following strings) names the image file(s) whose information should be printed. The image files may be of any format recognized by OpenImageIO (i.e., for which ImageInput plugins are available).

In its most basic usage, it simply prints the resolution, number of channels, pixel data type, and file format type of each of the files listed:

```
$ iinfo img_6019m.jpg grid.tif lenna.png

img_6019m.jpg : 1024 x 683, 3 channel, uint8 jpeg
grid.tif      : 512 x 512, 3 channel, uint8 tiff
lenna.png     : 120 x 120, 4 channel, uint8 png
```

The `-s` flag also prints the uncompressed sizes of each image file, plus a sum for all of the images:

```
$ iinfo -s img_6019m.jpg grid.tif lenna.png

img_6019m.jpg : 1024 x 683, 3 channel, uint8 jpeg (2.00 MB)
grid.tif      : 512 x 512, 3 channel, uint8 tiff (0.75 MB)
lenna.png     : 120 x 120, 4 channel, uint8 png (0.05 MB)
Total size: 2.81 MB
```

The `-v` option turns on *verbose mode*, which exhaustively prints all metadata about each image:

```
$ iinfo -v img_6019m.jpg

img_6019m.jpg : 1024 x 683, 3 channel, uint8 jpeg
  channel list: R, G, B
  Color space: sRGB
  ImageDescription: "Family photo"
  Make: "Canon"
  Model: "Canon EOS DIGITAL REBEL XT"
  Orientation: 1 (normal)
  XResolution: 72
  YResolution: 72
  ResolutionUnit: 2 (inches)
  DateTime: "2008:05:04 19:51:19"
  Exif:YCbCrPositioning: 2
  ExposureTime: 0.004
  FNumber: 11
  Exif:ExposureProgram: 2 (normal program)
  Exif:ISOSpeedRatings: 400
  Exif:DateTimeOriginal: "2008:05:04 19:51:19"
  Exif:DateTimeDigitized: "2008:05:04 19:51:19"
  Exif:ShutterSpeedValue: 7.96579 (1/250 s)
  Exif:ApertureValue: 6.91887 (f/11)
  Exif:ExposureBiasValue: 0
  Exif:MeteringMode: 5 (pattern)
  Exif:Flash: 16 (no flash, flash supression)
  Exif:FocalLength: 27 (27 mm)
  Exif:ColorSpace: 1
  Exif:PixelXDimension: 2496
  Exif:PixelYDimension: 1664
  Exif:FocalPlaneXResolution: 2855.84
  Exif:FocalPlaneYResolution: 2859.11
  Exif:FocalPlaneResolutionUnit: 2 (inches)
  Exif:CustomRendered: 0 (no)
  Exif:ExposureMode: 0 (auto)
  Exif:WhiteBalance: 0 (auto)
  Exif:SceneCaptureType: 0 (standard)
  Keywords: "Carly; Jack"
```

If the input file has multiple subimages, extra information summarizing the subimages will be printed:

```
$ iinfo img_6019m.tx

img_6019m.tx : 1024 x 1024, 3 channel, uint8 tiff (11 subimages)

$ iinfo -v img_6019m.tx

img_6019m.tx : 1024 x 1024, 3 channel, uint8 tiff
  11 subimages: 1024x1024 512x512 256x256 128x128 64x64 32x32 16x16 8x8 4x4 2x2 1x1
  channel list: R, G, B
  tile size: 64 x 64
```

...

Furthermore, the `-a` option will print information about all individual subimages:

```
$ iinfo -a ../sample-images/img_6019m.tx

img_6019m.tx : 1024 x 1024, 3 channel, uint8 tiff (11 subimages)
subimage 0: 1024 x 1024, 3 channel, uint8 tiff
subimage 1: 512 x 512, 3 channel, uint8 tiff
subimage 2: 256 x 256, 3 channel, uint8 tiff
subimage 3: 128 x 128, 3 channel, uint8 tiff
subimage 4: 64 x 64, 3 channel, uint8 tiff
subimage 5: 32 x 32, 3 channel, uint8 tiff
subimage 6: 16 x 16, 3 channel, uint8 tiff
subimage 7: 8 x 8, 3 channel, uint8 tiff
subimage 8: 4 x 4, 3 channel, uint8 tiff
subimage 9: 2 x 2, 3 channel, uint8 tiff
subimage 10: 1 x 1, 3 channel, uint8 tiff

$ iinfo -v -a img_6019m.tx
img_6019m.tx : 1024 x 1024, 3 channel, uint8 tiff
  11 subimages: 1024x1024 512x512 256x256 128x128 64x64 32x32 16x16 8x8 4x4 2x2 1x1
subimage 0: 1024 x 1024, 3 channel, uint8 tiff
  channel list: R, G, B
  tile size: 64 x 64
...
subimage 1: 512 x 512, 3 channel, uint8 tiff
  channel list: R, G, B
...
...
```

12.2 iinfo command-line options

`--help`

Prints usage information to the terminal.

`-v`

Verbose output — prints all metadata of the image files.

`-a`

Print information about all subimages in the file(s).

-f

Print the filename as a prefix to every line. For example,

```
$ iinfo -v -f img_6019m.jpg

img_6019m.jpg : 1024 x 683, 3 channel, uint8 jpeg
img_6019m.jpg : channel list: R, G, B
img_6019m.jpg : Color space: sRGB
img_6019m.jpg : ImageDescription: "Family photo"
img_6019m.jpg : Make: "Canon"
...
```

-m *pattern*

Match the *pattern* (specified as an extended regular expression) against data metadata field names and print only data fields whose names match. The default is to print all data fields found in the file (if **-v** is given).

For example,

```
$ iinfo -v -f -m ImageDescription test*.jpg

test3.jpg :      ImageDescription: "Birthday party"
test4.jpg :      ImageDescription: "Hawaii vacation"
test5.jpg :      ImageDescription: "Bob's graduation"
test6.jpg :      ImageDescription: <unknown>
```

Note: the **-m** option is probably not very useful without also using the **-v** and **-f** options.

--hash

Displays a SHA-1 hash of the pixel data of the image (and of each subimage if combined with the **-a** flag).

-s

Show the image sizes, including a sum of all the listed images.

13 Converting Image Formats With `iconvert`

13.1 Overview

The `iconvert` program will read an image (from any file format for which an `ImageInput` plugin can be found) and then write the image to a new file (in any format for which an `ImageOutput` plugin can be found). In the process, `iconvert` can optionally change the file format or data format (for example, converting floating-point data to 8-bit integers), apply gamma correction, switch between tiled and scanline orientation, or alter or add certain metadata to the image.

The `iconvert` utility is invoked as follows:

```
iconvert [options] input output
```

Where *input* and *output* name the input image and desired output filename. The image files may be of any format recognized by `OpenImageIO` (i.e., for which `ImageInput` plugins are available). The file format of the output image will be inferred from the file extension of the output filename (e.g., "`foo.tif`" will write a TIFF file).

Alternately, any number of files may be specified as follows:

```
iconvert --inplace [options] file1 file2 ..
```

When the `--inplace` option is used, any number of file names ≥ 1 may be specified, and the image conversion commands are applied to each file in turn, with the output being saved under the original file name. This is useful for applying the same conversion to many files, or simply if you want to replace the input with the output rather than create a new file with a different name.

13.2 `iconvert` Recipes

This section will give quick examples of common uses of `iconvert`.

Converting between file formats

It's a snap to converting among image formats supported by `OpenImageIO` (i.e., for which `ImageInput` and `ImageOutput` plugins can be found). The `iconvert` utility will simply infer the file format from the file extension. The following example converts a PNG image to JPEG:

```
iconvert lena.png lena.jpg
```

Changing the data format or bit depth

Just use the `-d` option to specify a pixel data format. For example, assuming that `in.tif` uses 16-bit unsigned integer pixels, the following will convert it to an 8-bit unsigned pixels:

```
iconvert -d uint8 in.tif out.tif
```

Changing the compression

The following command converts writes a TIFF file, specifically using LZW compression:

```
iconvert --compression lzw in.tif out.tif
```

The following command writes its results as a JPEG file at a compression quality of 50 (pretty severe compression):

```
iconvert --quality 50 big.jpg small.jpg
```

Gamma-correcting an image

The following gamma-corrects the pixels, raising all pixel values to $x^{1/2.2}$ upon writing:

```
iconvert -g 2.2 in.tif out.tif
```

Converting between scanline and tiled images

Convert a scanline file to a tiled file with 16×16 tiles:

```
iconvert --tile 16 16 s.tif t.tif
```

Convert a tiled file to scanline:

```
iconvert --scanline t.tif s.tif
```

Converting images in place

You can use the `--inplace` flag to cause the output to *replace* the input file, rather than create a new file with a different name. For example, this will re-compress all of your TIFF files to use ZIP compression (rather than whatever they currently are using):

```
iconvert --inplace --compression zip *.tif
```

Change the file modification time to the image capture time

Many image formats (including JPEGs from digital cameras) contain an internal time stamp indicating when the image was captured. But the time stamp on the file itself (what you'd see in a directory listing from your OS) most likely shows when the file was last copied, not when it was created or captured. You can use the following command to re-stamp your files so that the file system modification time matches the time that the digital image was originally captured:

```
iconvert --inplace --adjust-time *.jpg
```

Add captions, keywords, IPTC tags

For formats that support it, you can add a caption/image description, keywords, or arbitrary string metadata:

```
iconvert --inplace --adjust-time --caption "Hawaii vacation" *.jpg
```

```
iconvert --inplace --adjust-time --keyword "John" img18.jpg img21.jpg
```

```
iconvert --inplace --adjust-time --attrib IPTC:State "HI" \  
--attrib IPTC:City "Honolulu" *.jpg
```

13.3 iconvert command-line options

--help

Prints usage information to the terminal.

-v

Verbose status messages.

--inplace

Causes the output to *replace* the input file, rather than create a new file with a different name.

Without this flag, `iconvert` expects two file names, which will be used to specify the input and output files, respectively.

But when `--inplace` option is used, any number of file names ≥ 1 may be specified, and the image conversion commands are applied to each file in turn, with the output being saved under the original file name. This is useful for applying the same conversion to many files.

For example, the following example will add the caption “Hawaii vacation” to all JPEG files in the current directory:

```
iconvert --inplace --adjust-time --caption "Hawaii vacation" *.jpg
```

`-d datatype`

Attempt to sets the output pixel data type to one of: `uint8`, `sint8`, `uint16`, `sint16`, `half`, `float`, `double`.

If the `-d` option is not supplied, the output data type will be the same as the data format of the input file.

In either case, the output file format itself (implied by the file extension of the output filename) may trump the request if the file format simply does not support the requested data type.

`-g gamma`

Applies a gamma correction of $1/\text{gamma}$ to the pixels as they are output.

`--sRGB`

Explicitly tags the image as being in sRGB color space. Note that this does not alter pixel values, it only marks which color space those values refer to (and only works for file formats that understand such things). An example use of this command is if you have an image that is not explicitly marked as being in any particular color space, but you know that the values are sRGB.

`--tile x y`

Requests that the output file be tiled, with the given $x \times y$ tile size, if tiled images are supported by the output format. By default, the output file will take on the tiledness and tile size of the input file.

`--scanline`

Requests that the output file be scanline-oriented (even if the input file was tile-oriented), if scanline orientation is supported by the output file format. By default, the output file will be scanline if the input is scanline, or tiled if the input is tiled.

`--separate`

`--contig`

Forces either “separate” (e.g., RRR...GGG...BBB) or “contiguous” (e.g., RGBRGBRGB...) packing of channels in the file. If neither of these options are present, the output file will have the same kind of channel packing as the input file. Of course, this is ignored if the output file format does not support a choice or does not support the particular choice requested.

--compression *method*

Sets the compression method for the output image. Each ImageOutput plugin will have its own set of methods that it supports.

By default, the output image will use the same compression technique as the input image (assuming it is supported by the output format, otherwise it will use the default compression method of the output plugin).

--quality *q*

Sets the compression quality, on a 1–100 floating-point scale. This only has an effect if the particular compression method supports a quality metric (as JPEG does).

--no-copy-image

Ordinarily, `iconvert` will attempt to use `ImageOutput::copy_image` underneath to avoid de/recompression or alteration of pixel values, unless other settings clearly contradict this (such as any settings that must alter pixel values). The use of `--no-copy-image` will force all pixels to be decompressed, read, and compressed/written, rather than copied in compressed form. We're not exactly sure when you would need to do this, but we put it in just in case.

--adjust-time

When this flag is present, after writing the output, the resulting file's modification time will be adjusted to match any "DateTime" metadata in the image. After doing this, a directory listing will show file times that match when the original image was created or captured, rather than simply when `iconvert` was run. This has no effect on image files that don't contain any "DateTime" metadata.

--caption *text*

Sets the image metadata "ImageDescription". This has no effect if the output image format does not support some kind of title, caption, or description metadata field. Be careful to enclose *text* in quotes if you want your caption to include spaces or certain punctuation!

--keyword *text*

Adds a keyword to the image metadata "Keywords". Any existing keywords will be preserved, not replaced, and the new keyword will not be added if it is an exact duplicate of existing keywords. This has no effect if the output image format does not support some kind of keyword field.

Be careful to enclose *text* in quotes if you want your keyword to include spaces or certain punctuation. For image formats that have only a single field for keywords, OpenImageIO will concatenate the keywords, separated by semicolon (';'), so don't use semicolons within your keywords.

`--clear-keywords`

Clears all existing keywords in the image.

`--attrib name text`

Sets the named image metadata attribute to a string given by *text*. For example, you could explicitly set the IPTC location metadata fields with:

```
iconvert --attrib "IPTC:City" "Berkeley" in.jpg out.jpg
```

`--orientation orient`

Explicitly sets the image's "Orientation" metadata to a numeric value (see Section B.2 for the numeric codes). This only changes the metadata field that specifies how the image should be displayed, it does NOT alter the pixels themselves, and so has no effect for image formats that don't support some kind of orientation metadata.

`--rotcw`

`--rotccw`

`--rot180`

Adjusts the image's "Orientation" metadata by rotating it 90° clockwise, 90° degrees counter-clockwise, or 180°, respectively, compared to its current setting. This only changes the metadata field that specifies how the image should be displayed, it does NOT alter the pixels themselves, and so has no effect for image formats that don't support some kind of orientation metadata.

14 Searching Image Metadata With `igrep`

The `igrep` program search one or more image files for metadata that match a string or regular expression.

14.1 Using `igrep`

The `igrep` utility is invoked as follows:

```
igrep [options] pattern filename ...
```

Where *pattern* is a POSIX.2 regular expression (just like the Unix/Linux `grep(1)` command), and *filename* (and any following names) specify images or directories that should be searched. An image file will “match” if any of its metadata contains values contain substring that are recognized regular expression. The image files may be of any format recognized by OpenImageIO (i.e., for which ImageInput plugins are available).

Example:

```
$ igrep Jack *.jpg
bar.jpg: Keywords = Carly; Jack
foo.jpg: Keywords = Jack
test7.jpg: ImageDescription = Jack on vacation
```

14.2 `igrep` command-line options

`--help`

Prints usage information to the terminal.

`-d`

Print directory names as it recurses. This only happens if the `-r` option is also used.

`-E`

Interpret the pattern as an extended regular expression (just like `egrep` or `grep -E`).

-f

Match the expression against the filename, as well as the metadata within the file.

-i

Ignore upper/lower case distinctions. Without this flag, the expression matching will be case-sensitive.

-l

Simply list the matching files by name, surpressing the normal output that would include the metadata name and values that matched. For example:

```
$ igrep Jack *.jpg
bar.jpg: Keywords = Carly; Jack
foo.jpg: Keywords = Jack
test7.jpg: ImageDescription = Jack on vacation

$ igrep -l Jack *.jpg
bar.jpg
foo.jpg
test7.jpg
```

-r

Recurse into directories. If this flag is present, any files specified that are directories will have any image file contained therein to be searched for a match (an so on, recursively).

-v

Invert the sense of matching, to select image files that *do not* match the expression.

15 Comparing Images With `idiff`

15.1 Overview

The `idiff` program compares two images, printing a report about how different they are and optionally producing a third image that records the pixel-by-pixel differences between them. There are a variety of options and ways to compare (absolute pixel difference, various thresholds for warnings and errors, and also an optional perceptual difference metric).

Because `idiff` is built on top on `OpenImageIO`, it can compare two images of any formats readable by `ImageInput` plugins on hand. They may have any (or different) file formats, data formats, etc.

15.2 Using `idiff`

The `idiff` utility is invoked as follows:

```
idiff [options] image1 image2
```

Where *input1* and *input2* are the names of two image files that should be compared. They may be of any format recognized by `OpenImageIO` (i.e., for which image-reading plugins are available).

If the two input images are not the same resolutions, or do not have the same number of channels, the comparison will return `FAILURE` immediately and will not attempt to compare the pixels of the two images. If they are the same dimensions, the pixels of the two images will be compared, and a report will be printed including the mean and maximum error, how many pixels were above the warning and failure thresholds, and whether the result is `PASS`, `WARNING`, or `FAILURE`. For example:

```
$ idiff a.jpg b.jpg

Comparing "a.jpg" and "b.jpg"
Mean error = 0.00450079
RMS error = 0.00764215
Peak SNR = 42.3357
Max error = 0.254902 @ (700, 222, B)
574062 pixels (82.1%) over 1e-06
574062 pixels (82.1%) over 1e-06
FAILURE
```

The “mean error” is the average difference (per channel, per pixel). The “max error” is the largest difference in any pixel channel, and will point out on which pixel and channel it was found. It will also give a count of how many pixels were above the warning and failure thresholds.

The metadata of the two images (e.g., the comments) are not currently compared; only differences in pixel values are taken into consideration.

Raising the thresholds

By default, if any pixels differ between the images, the comparison will fail. You can allow *some* differences to still pass by raising the failure thresholds. The following example will allow images to pass the comparison test, as long as no more than 10% of the pixels differ by 0.004 (just above a 1/255 threshold):

```
idiff -fail 0.004 -failpercent 10 a.jpg b.jpg
```

But what happens if a just a few pixels are very different? Maybe you want that to fail, also. The following adjustment will fail if at least 10% of pixels differ by 0.004, or if *any* pixel differs by more than 0.25:

```
idiff -fail 0.004 -failpercent 10 -hardfail 0.25 a.jpg b.jpg
```

If none of the failure criteria are met, and yet some pixels are still different, it will still give a WARNING. But you can also raise the warning threshold in a similar way:

```
idiff -fail 0.004 -failpercent 10 -hardfail 0.25 \  
-warn 0.004 -warnpercent 3 a.jpg b.jpg
```

The above example will PASS as long as fewer than 3% of pixels differ by more than 0.004. If it does, it will be a WARNING as long as no more than 10% of pixels differ by 0.004 and no pixel differs by more than 0.25, otherwise it is a FAILURE.

Output a difference image

Ordinary text output will tell you how many pixels failed or were warnings, and which pixel had the biggest difference. But sometimes you need to see visually where the images differ. You can get `idiff` to save an image of the differences between the two input images:

```
idiff -o diff.tif -abs a.jpg b.jpg
```

The `-abs` flag saves the absolute value of the differences (i.e., all positive values or zero). If you omit the `-abs`, pixels in which `a.jpg` have smaller values than `b.jpg` will be negative in the difference image (be careful in this case of using a file format that doesn’t support negative values).

You can also scale the difference image with the `-scale`, making them easier to see. And the `-od` flag can be used to output a difference image only if the comparison fails, but not if the images pass within the designated threshold (thus saving you the trouble and space of saving a black image).

15.3 idiff Reference

The various command-line options are discussed below:

General options

`--help`

Prints usage information to the terminal.

`-v`

Verbose output — more detail about what it finds when comparing images. (Currently, there is no extra info to print.)

`-a`

Compare all subimages. Without this flag, only the first subimage of each file will be compared.

Thresholds and comparison options

`-fail A`

`-failpercent B`

`-hardfail C`

Sets the threshold for FAILURE: if more than $B\%$ of pixels (on a 0-100 floating point scale) are greater than A different, or if *any* pixels are more than C different. The defaults are to fail if more than 0% (any) pixels differ by more than 0.00001 (1e-6), and C is infinite.

`-warn A`

`-warnpercent B`

`-hardwarn C`

Sets the threshold for WARNING: if more than $B\%$ of pixels (on a 0-100 floating point scale) are greater than A different, or if *any* pixels are more than C different. The defaults are to warn if more than 0% (any) pixels differ by more than 0.00001 (1e-6), and C is infinite.

`-p`

Does an additional test on the images to attempt to see if they are *perceptually* different (whether you are likely to discern a difference visually), using Hector Yee's metric. If this option is enabled, the statistics will additionally show a report on how many pixels failed the perceptual test, and the test overall will fail if more than the "fail percentage" failed the perceptual test.

Difference image output

`-o outfile`

Outputs a *difference image* to the designated file. This difference image pixels consist are each of the value of the corresponding pixel from *image1* minus the value of the pixel *image2*.

The file extension of the output file is used to determine the file format to write (e.g., "out.tif" will write a TIFF file, "out.jpg" will write a JPEG, etc.). The data format of the output file will be format of whichever of the two input images has higher precision (or the maximum precision that the designated output format is capable of, if that is less than either of the input images).

Note that pixels whose value is lower in *image1* than in *image2*, this will result in negative pixels (which may be clamped to zero if the image format does not support negative values)), unless the `-abs` option is also used.

`-abs`

Will cause the output image to consist of the *absolute value* of the difference between the two input images (so all values in the difference image ≥ 0).

`-scale factor`

Scales the values in the difference image by the given (floating point) factor. The main use for this is to make small actual differences more visible in the resulting difference image by giving a large scale factor.

`-od`

Causes a difference image to be produce *only* if the image comparison fails. That is, even if the `-o` option is used, images that are within the comparison threshold will not write out a useless black (or nearly black) difference image.

Process return codes

The `idiff` program will return a code that can be used by scripts to indicate the results:

- 0 OK: the images match within the warning and error thresholds.
- 1 Warning: the errors differ a little, but within error thresholds.
- 2 Failure: the errors differ a lot, outside error thresholds.
- 3 The images weren't the same size and couldn't be compared.
- 4 File error: could not find or open input files, etc.

16 Making Tiled MIP-Map Texture Files With `maketx`

16.1 Overview

The `maketx` program will read an image (from any file format for which an `ImageInput` plugin can be found) and then write it in a form in which it will have high performance when used by `TextureSystem` (Chapter 8). This involves converting it to tiled (versus scanline) orientation, writing multiple subimages at different resolutions (MIP-map), and setting a variety of header or metadata fields appropriately for texture maps.

The `maketx` utility is invoked as follows:

```
maketx [options] input... -o output
```

Where *input* and *output* name the input image and desired output filename. The input files may be of any image format recognized by `OpenImageIO` (i.e., for which `ImageInput` plugins are available). The file format of the output image will be inferred from the file extension of the output filename (e.g., "`foo.tif`" will write a TIFF file).

16.2 `maketx` command-line options

`--help`

Prints usage information to the terminal.

`-v`

Verbose status messages, including runtime statistics and timing.

`-o outputname`

Sets the name of the output texture.

`--format formatname`

Specifies the image format of the output file (e.g., "`tiff`", "`OpenEXR`", etc.). If `--format` is not used, `maketx` will guess based on the file extension of the output filename; if it is not a recognized format extension, TIFF will be used by default.

-d *datatype*

Attempt to sets the output pixel data type to one of: `uint8`, `sint8`, `uint16`, `sint16`, `half`, `float`, `double`.

If the `-d` option is not supplied, the output data type will be the same as the data format of the input file.

In either case, the output file format itself (implied by the file extension of the output filename) may trump the request if the file format simply does not support the requested data type.

--tile *x y*

Specifies the tile size of the output texture. If not specified, `maketx` will make 64×64 tiles.

--separate

Forces “separate” (e.g., RRR...GGG...BBB) packing of channels in the output file. Without this option specified, “contiguous” (e.g., RGBRGBRGB...) packing of channels will be used for those file formats that support it.

--update

Ordinarily, textures are created unconditionally (which could take several seconds for large input files if read over a network) and will be stamped with the current time.

The `--update` option enables *update mode* if the output file already exists and has the same time stamp as the input file, the texture will not be recreated. If the output file does not exist or has a different time than the input file, then the texture will be created be given the time stamp of the input file.

--wrap *wrapmode***--swrap *wrapmode* --twrap *wrapmode***

Sets the default *wrap mode* for the texture, which determines the behavior when the texture is sampled outside the $[0, 1]$ range. Valid wrap modes are: `black`, `clamp`, `periodic`, `mirror`. The default, if none is set, is `black`. The `--wrap` option sets the wrap mode in both directions simultaneously, while the `--swrap` and `--twrap` may be used to set them individually in the *s* (horizontal) and *t* (vertical) directions.

Although this sets the default wrap mode for a texture, note that the wrap mode may have an override specified in the texture lookup at runtime.

--resize

Causes the highest-resolution level of the MIP-map to be a power-of-two resolution in each dimension (by rounding up the resolution of the input image). There is no good reason to do this for the sake of OIIO’s texture system, but some users may require it in order to create MIP-map images that are compatible with both OIIO and other texturing systems that require power-of-2 textures.

`--nomipmap`

Causes the output to *not* be MIP-mapped, i.e., only will have the highest-resolution level.

`--checknan`

Checks every pixel of the input image to ensure that no NaN or Inf values are present. If such non-finite pixel values are found, an error message will be printed and `maketx` will terminate without writing the output image (returning an error code).

`--ingamma value`

`--outgamma value`

Not currently implemented

`--Mcamera ...16 floats...`

`--Mscreen ...16 floats...`

Sets the camera and screen matrices (sometimes called N1 and NP, respectively, by some renderers) in the texture file, overriding any such matrices that may be in the input image (and would ordinarily be copied to the output texture).

`--hash`

Computes a SHA-1 hash on the input file's pixels and embeds this hash in the "ImageDescription" metadata of the output texture. This is useful in helping the `TextureSystem` identify duplicate textures at runtime.

Part III

Appendices

A Building OpenImageIO

B Metadata conventions

The ImageSpec class, described thoroughly in Section 2.2, provides the basic description of an image that are essential across all formats — resolution, number of channels, pixel data format, etc. Individual images may have additional data, stored as name/value pairs in the `extra_attrs` field. Though literally *anything* can be stored in `extra_attrs` — it’s specifically designed for format- and user-extensibility — this chapter establishes some guidelines and lays out all of the field names that OpenImageIO understands.

B.1 Description of the image

"ImageDescription" : string

The image description, title, caption, or comments.

"Keywords" : string

Semicolon-separated keywords describing the contents of the image. (Semicolons are used rather than commas because of the common case of a comma being part of a keyword itself, e.g., “Kurt Vonnegut, Jr.” or “Washington, DC.”)

"Artist" : string

The artist, creator, or owner of the image.

"Copyright" : string

Any copyright notice or owner of the image.

"DateTime" : string

The creation date of the image, in the following format: YYYY:MM:DD HH:MM:SS (exactly 19 characters long, not including a terminating NULL). For example, 7:30am on Dec 31, 2008 is encoded as "2008:12:31 07:30:00".

"DocumentName" : string

The name of an overall document that this image is a part of.

"Software" : string

The software that was used to create the image.

"HostComputer" : string

The name or identity of the computer that created the image.

B.2 Display hints

"Orientation" : int

By default, image pixels are ordered from the top of the display to the bottom, and within each scanline, from left to right (i.e., the same ordering as English text and scan progression on a CRT). But the "Orientation" field can suggest that it should be displayed with a different orientation, according to the TIFF/EXIF conventions:

- 1 normal (top to bottom, left to right)
- 2 flipped horizontally (top to bottom, right to left)
- 3 rotate 180° (bottom to top, right to left)
- 4 flipped vertically (bottom to top, left to right)
- 5 transposed (left to right, top to bottom)
- 6 rotated 90° clockwise (right to left, top to bottom)
- 7 transverse (right to left, bottom to top)
- 8 rotated 90° counter-clockwise (left to right, bottom to top)

"PixelAspectRatio" : float

The aspect ratio (x/y) of the individual pixels, with square pixels being 1.0 (the default).

"XResolution" : float

"YResolution" : float

"ResolutionUnit" : string

The number of horizontal (x) and vertical (y) pixels per resolution unit. This ties the image to a physical size (where applicable, such as with a scanned image, or an image that will eventually be printed).

Different file formats may dictate different resolution units. For example, the TIFF ImageIO plugin supports "none", "in", and "cm".

B.3 Disk file format info/hints

"BitsPerSample" : int

Number of bits per sample *in the file*.

Note that this may not match the reported `ImageSpec::format`, if the plugin is translating from an unsupported format. For example, if a file stores 4 bit grayscale per channel, the "BitsPerSample" may be 4 but the format field may be `TypeDesc::UINT8` (because the `OpenImageIO` APIs do not support fewer than 8 bits per sample).

"planarconfig" : string

"contig" indicates that the file has contiguous pixels (RGB RGB RGB...), whereas "separate" indicate that the file stores each channel separately (RRR...GGG...BBB...).

Note that only contiguous pixels are transmitted through the `OpenImageIO` APIs, but this metadata indicates how it is (or should be) stored in the file, if possible.

"compression" : string

Indicates the type of compression the file uses. Supported compression modes will vary from `ImageInput` plugin to plugin, and each plugin should document the modes it supports. If `ImageInput::open` is called with an `ImageSpec` that specifies an compression mode not supported by that `ImageInput`, it will choose a reasonable default. As an example, the TIFF `ImageInput` plugin supports "none", "lzw", "ccittrle", "zip" (the default), "packbits".

"CompressionQuality" : int

Indicates the quality of compression to use (0–100), for those plugins and compression methods that allow a variable amount of compression, with higher numbers indicating higher image fidelity.

B.4 Photographs or scanned images

The following metadata items are specific to photos or captured images.

"Make" : string

For captured or scanned image, the make of the camera or scanner.

"Model" : string

For captured or scanned image, the model of the camera or scanner.

"ExposureTime" : float

The exposure time (in seconds) of the captured image.

"FNumber" : float

The f/stop of the camera when it captured the image.

B.5 Texture Information

Several standard metadata are very helpful for images that are intended to be used as textures (especially for OpenImageIO's TextureSystem).

"textureformat" : string

The kind of texture that this image is intended to be. We suggest the following names:

"Plain Texture"	Ordinary 2D texture
"Volume Texture"	3D volumetric texture
"Shadow"	Ordinary z -depth shadow map
"CubeFace Shadow"	Cube-face shadow map
"Volume Shadow"	Volumetric ("deep") shadow map
"LatLong Environment"	Latitude-longitude (rectangular) environment map
"CubeFace Environment"	Cube-face environment map

"wrapmodes" : string

Give the intended texture *wrap mode* indicating what happens with texture coordinates outside the $[0...1]$ range. We suggest the following names: "black", "periodic", "clamp", "mirror". If the wrap mode is different in each direction, they should simply be separated by a comma. For example, "black" means black wrap in both directions, whereas "clamp,periodic" means to clamp in u and be periodic in v .

"fovcot" : float

The cotangent (x/y) of the field of view of the original image (which may not be the same as the aspect ratio of the pixels of the texture, which may have been resized).

"worldtocamera" : matrix44

For shadow maps or rendered images this item (of type TypeDesc::PT_MATRIX) is the world-to-camera matrix describing the camera position.

"worldtoscreen" : matrix44

For shadow maps or rendered images this item (of type TypeDesc::PT_MATRIX) is the world-to-screen matrix describing the full projection of the 3D view onto a $[-1...1] \times [-1...1]$ 2D domain.

"updirection" : string

For environment maps, indicates which direction is "up" (valid values are "y" or "z"), to disambiguate conventions for environment map orientation.

B.6 Exif metadata

The following Exif metadata tags correspond to items in the “standard” set of metadata.

Exif tag	OpenImageIO metadata convention
ColorSpace	(stored in ImageSpec::Linearity)
ExposureTime	"ExposureTime"
FNumber	"FNumber"

The other remaining Exif metadata tags all include the “Exif:” prefix to keep it from clashing with other names that may be used for other purposes.

"Exif:ExposureProgram" : int

The exposure program used to set exposure when the picture was taken:

- 0 unknown
- 1 manual
- 2 normal program
- 3 aperture priority
- 4 shutter priority
- 5 Creative program (biased toward depth of field)
- 6 Action program (biased toward fast shutter speed)
- 7 Portrait mode (closeup photo with background out of focus)
- 8 Landscape mode (background in focus)

"Exif:SpectralSensitivity" : string

The camera’s spectral sensitivity, using the ASTM conventions.

"Exif:ISOSpeedRatings" : int

The ISO speed and ISO latitude of the camera as specified in ISO 12232.

"Exif:DateTimeOriginal" : string

Date and time that the original image data was generated (in "YYYY:MM:DD HH:MM:SS" format).

"Exif:DateTimeDigitized" : string

Date and time that the image was stored as digital data (in "YYYY:MM:DD HH:MM:SS" format).

"Exif:CompressedBitsPerPixel" : float

The compression mode used, measured in compressed bits per pixel.

"Exif:ShutterSpeedValue" : float

Shutter speed, in APEX units: $-\log_2(\textit{exposuretime})$

"Exif:ApertureValue" : float

Aperture, in APEX units: $2\log_2(\textit{fnumber})$

"Exif:BrightnessValue" : float

Brightness value, assumed to be in the range of $-99.99 - 99.99$.

"Exif:ExposureBiasValue" : float

Exposure bias, assumed to be in the range of $-99.99 - 99.99$.

"Exif:MaxApertureValue" : float

Smallest F number of the lens, in APEX units: $2\log_2(\textit{fnumber})$

"Exif:SubjectDistance" : float

Distance to the subject, in meters.

"Exif:MeteringMode" : int

The metering mode:

0	unknown
1	average
2	center-weighted average
3	spot
4	multi-spot
5	pattern
6	partial
255	other

"Exif:LightSource" : int

The kind of light source:

0	unknown
1	daylight
2	tungsten (incandescent light)
4	flash
9	fine weather
10	cloudy weather
11	shade
12	daylight fluorescent (D 5700-7100K)
13	day white fluorescent (N 4600-5400K)
14	cool white fluorescent (W 3900 - 4500K)
15	white fluorescent (WW 3200 - 3700K)
17	standard light A
18	standard light B
19	standard light C
20	D55
21	D65
22	D75
23	D50
24	ISO studio tungsten
255	other light source

"Exif:Flash" int

A sum of:

1	if the flash fired
0	no strobe return detection function
4	strobe return light was not detected
6	strobe return light was detected
8	compulsary flash firing
16	compulsary flash supression
24	auto-flash mode
32	no flash function (0 if flash function present)
64	red-eye reduction supported (0 if no red-eye reduction mode)

"Exif:FocalLength" : float

Actual focal length of the lens, in mm.

"Exif:SubsecTime" : string

Fractions of a second to augment the "DateTime" (expressed as text of the digits to the right of the decimal).

"Exif:SubsecTimeOriginal" : string

Fractions of a second to augment the "Exif:DateTimeOriginal" (expressed as text of the digits to the right of the decimal).

"Exif:SubsecTimeDigitized" : string

Fractions of a second to augment the "Exif:DateTimeDigital" (expressed as text of the digits to the right of the decimal).

"Exif:PixelXDimension" : int

"Exif:PixelYDimension" : int

The *x* and *y* dimensions of the valid pixel area. FIXME – better explanation?

"Exif:FlashEnergy" : float

Strobe energy when the image was captures, measured in Beam Candle Power Seconds (BCPS).

"Exif:FocalPlaneXResolution" : float

"Exif:FocalPlaneYResolution" : float

"Exif:FocalPlaneResolutionUnit" : int

The number of pixels in the *x* and *y* dimension, per resolution unit. The code for resolution units is: 2 for inches.

"Exif:ExposureIndex" : float

The exposure index selected on the camera.

"Exif:SensingMethod" : int

The image sensor type on the camra:

- | | |
|---|------------------------------|
| 1 | undefined |
| 2 | one-chip color area sensor |
| 3 | two-chip color area sensor |
| 4 | three-chip color area sensor |
| 5 | color sequential area sensor |
| 7 | trilinear sensor |
| 8 | color trilinear sensor |

"Exif:FileSource" : int

Set to 3, if captured by a digital camera, otherwise it should not be present.

"Exif:SceneType" : int

Set to 1, if a directly-photographed image, otherwise it should not be present.

"Exif:CustomRendered" : int

Set to 0 for a normal process, 1 if some custom processing has been performed on the image data.

"Exif:ExposureMode" : int

The exposure mode:

- 0 auto
- 1 manual
- 2 auto-bracket

"Exif:WhiteBalance" : int

Set to 0 for auto white balance, 1 for manual white balance.

"Exif:DigitalZoomRatio" : float

The digital zoom ratio used when the image was shot.

"Exif:FocalLengthIn35mmFilm" : int

The equivalent focal length of a 35mm camera, in mm.

"Exif:SceneCaptureType" : int

The type of scene that was shot:

- 0 standard
- 1 landscape
- 2 portrait
- 3 night scene

"Exif:GainControl" : float

The degree of overall gain adjustment:

- 0 none
- 1 low gain up
- 2 high gain up
- 3 low gain down
- 4 high gain down

"Exif:Contrast" : int

The direction of contrast processing applied by the camera:

- 0 normal
- 1 soft
- 2 hard

"Exif:Saturation" : int

The direction of saturation processing applied by the camera:

- 0 normal
- 1 low saturation
- 2 high saturation

"Exif:Sharpness" : int

The direction of sharpness processing applied by the camera:

- 0 normal
- 1 soft
- 2 hard

"Exif:SubjectDistanceRange" : int

The distance to the subject:

- 0 unknown
- 1 macro
- 2 close
- 3 distant

"Exif:ImageUniqueID" : string

A unique identifier for the image, as 16 ASCII hexadecimal digits representing a 128-bit number.

B.7 GPS Exif metadata

The following GPS-related Exif metadata tags correspond to items in the “standard” set of metadata.

"GPS:LatitudeRef" : string

Whether the "GPS:Latitude" tag refers to north or south: "N" or "S".

"GPS:Latitude" : float[3]

The degrees, minutes, and seconds of latitude (see also "GPS:LatitudeRef").

"GPS:LongitudeRef" : string

Whether the "GPS:Longitude" tag refers to east or west: "E" or "W".

"GPS:Longitude" : float[3]

The degrees, minutes, and seconds of longitude (see also "GPS:LongitudeRef").

"GPS:AltitudeRef" : string

A value of 0 indicates that the altitude is above sea level, 1 indicates below sea level.

"GPS:Altitude" : float

Absolute value of the altitude, in meters, relative to sea level (see "GPS:AltitudeRef" for whether it's above or below sea level).

"GPS:TimeStamp" : float[3]

Gives the hours, minutes, and seconds, in UTC.

"GPS:Satellites" : string

Information about what satellites were visible.

"GPS:Status" : string

"A" indicates a measurement in progress, "V" indicates measurement interoperability.

"GPS:MeasureMode" : string

"2" indicates a 2D measurement, "3" indicates a 3D measurement.

"GPS:DOP" : float

Data degree of precision.

"GPS:SpeedRef" : string

Indicates the units of the related "GPS:Speed" tag: "K" for km/h, "M" for miles/h, "N" for knots.

"GPS:Speed" : float

Speed of the GPS receiver (see "GPS:SpeedRef" for the units).

"GPS:TrackRef" : string

Describes the meaning of the "GPS:Track" field: "T" for true direction, "M" for magnetic direction.

"GPS:Track" : float

Direction of the GPS receiver movement (from 0–359.99). The related "GPS:TrackRef" indicate whether it's true or magnetic.

"GPS:ImgDirectionRef" : string

Describes the meaning of the "GPS:ImgDirection" field: "T" for true direction, "M" for magnetic direction.

"GPS:ImgDirection" : float

Direction of the image when captured (from 0–359.99). The related "GPS:ImgDirectionRef" indicate whether it's true or magnetic.

"GPS:MapDatum" : string

The geodetic survey data used by the GPS receiver.

"GPS:DestLatitudeRef" : string

Whether the "GPS:DestLatitude" tag refers to north or south: "N" or "S".

"GPS:DestLatitude" : float[3]

The degrees, minutes, and seconds of latitude of the destination (see also "GPS:DestLatitudeRef").

"GPS:DestLongitudeRef" : string

Whether the "GPS:DestLongitude" tag refers to east or west: "E" or "W".

"GPS:DestLongitude" : float[3]

The degrees, minutes, and seconds of longitude of the destination (see also "GPS:DestLongitudeRef").

"GPS:DestBearingRef" : string

Describes the meaning of the "GPS:DestBearing" field: "T" for true direction, "M" for magnetic direction.

"GPS:DestBearing" : float

Bearing to the destination point (from 0–359.99). The related "GPS:DestBearingRef" indicate whether it's true or magnetic.

"GPS:DestDistanceRef" : string

Indicates the units of the related "GPS:DestDistance" tag: "K" for km, "M" for miles, "N" for knots.

"GPS:DestDistance" : float

Distance to the destination (see "GPS:DestDistanceRef" for the units).

"GPS:ProcessingMethod" : string

Processing method information.

"GPS:AreaInformation" : string

Name of the GPS area.

"GPS:DateStamp" : string

Date according to the GPS device, in format "YYYY:MM:DD".

"GPS:Differential" : int

If 1, indicates that differential correction was applied.

B.8 IPTC metadata

The IPTC (International Press Telecommunications Council) publishes conventions for storing image metadata, and this standard is growing in popularity and is commonly used in photo-browsing programs to record captions and keywords.

The following IPTC metadata items correspond exactly to metadata in the OpenImageIO conventions, so it is recommended that you use the standards and that plugins supporting IPTC metadata respond likewise:

IPTC tag	OpenImageIO metadata convention
Caption	"ImageDescription"
Keyword	IPTC keywords should be concatenated, separated by semicolons (;), and stored as the "Keywords" attribute.
ExposureTime	"ExposureTime"
CopyrightNotice	"Copyright"
Creator	"Artist"

The remainder of IPTC metadata fields should use the following names, prefixed with "IPTC:" to avoid conflicts with other plugins or standards.

"IPTC:ObjectName" : string

The name of the object in the picture.

"IPTC:Instructions" : string

Special instructions for handling the image.

"IPTC:AuthorsPosition" : string

The job title or position of the creator of the image.

"IPTC:City" : string

"IPTC:State" : string

"IPTC:Country" : string

The city, state, and country of the location of the image.

"IPTC:Headline" : string

Any headline that is meant to accompany the image.

"IPTC:Provider" : string

The provider of the image, or credit line.

"IPTC:Source" : string

The source of the image.

"IPTC:Contact" : string

The contact information for the image (possibly including name, address, email, etc.).

"IPTC:CaptionWriter" : string

The name of the person who wrote the caption or description of the image.

B.9 Extension conventions

To avoid conflicts with other plugins, or with any additional standard metadata names that may be added in future versions of `OpenImageIO`, it is strongly advised that writers of new plugins should prefix their metadata with the name of the format, much like the "Exif:" and "IPTC:" metadata.

C Glossary

Channel One of several data values present in each pixel. Examples include red, green, blue, alpha, etc. The data in one channel of a pixel may be represented by a single number, whereas the pixel as a whole requires one number for each channel.

Client A client (as in “client application”) is a program or library that uses `OpenImageIO` or any of its constituent libraries.

Data format The type of numerical representation used to store a piece of data. Examples include 8-bit unsigned integers, 32-bit floating-point numbers, etc.

Image File Format The specification and data layout of an image on disk. For example, TIFF, JPEG/JFIF, OpenEXR, etc.

Image Format Plugin A DSO/DLL that implements the `ImageInput` and `ImageOutput` classes for a particular image file format.

Format Plugin See *image format plugin*.

Metadata Data about data. As used in `OpenImageIO`, this means Information about an image, beyond describing the values of the pixels themselves. Examples include the name of the artist that created the image, the date that an image was scanned, the camera settings used when a photograph was taken, etc.

Native data format The *data format* used in the disk file representing an image. Note that with `OpenImageIO`, this may be different than the data format used by an application to store the image in the computer’s RAM.

Pixel One pixel element of an image, consisting of one number describing each *channel* of data at a particular location in an image.

Plugin See *image format plugin*.

Scanline A single horizontal row of pixels of an image. See also *tile*.

Scanline Image An image whose data layout on disk is organized by breaking the image up into horizontal scanlines, typically with the ability to read or write an entire scanline at once. See also *tiled image*.

Tile A rectangular region of pixels of an image. A rectangular tile is more spatially coherent than a scanline that stretches across the entire image — that is, a pixel’s neighbors are most likely in the same tile, whereas a pixel in a scanline image will typically have most of its immediate neighbors on different scanlines (requiring additional scanline reads in order to access them).

Tiled Image An image whose data layout on disk is organized by breaking the image up into rectangular regions of pixels called *tiles*. All the pixels in a tile can be read or written at once, and individual tiles may be read or written separately from other tiles.

Volume Image A 3-D set of pixels that has not only horizontal and vertical dimensions, but also a “depth” dimension.

Index

attribute, 77, 92

crop windows, 23

data formats, 7

getAttribute, 78, 93

iconvert, 119

idiff, 127

igrep, 125

iinfo, 115

Image Cache, 75–84

Image I/O API, 7–15, 17–53

ImageOutput, 17

ImageSpec, 9

iv, 113

maketx, 131

Orientation, 140

overscan, 23

Texture System, 85–105

TextureOpt, 85